

Rakshitha Prakash

Senior Data Engineer Consultant

📍 Dallas, Texas | 📞 (469) 607-9466 | @ rakshithaaprakash54@gmail.com |
🌐 <https://www.linkedin.com/in/rakshithaaprakash54>

Summary

Seasoned data engineer focused on architecting and delivering Microsoft-centric data solutions. I design scalable lakehouse and warehouse architectures using Azure Data Factory, ADLS Gen2, Synapse, and Azure Databricks, build metadata-driven ETL/ELT pipelines, and author dimensional (star) models and advanced SQL transformations to enable analytics and reliable reporting. I partner closely with business stakeholders to discover requirements, lead design workshops, document solutions, and mentor teams—bringing strong communication, ownership, and a continuous-growth mindset aligned with JourneyTeam values.

Experience

First Republic Bank Corporation **Sept 2024 - Present**
Azure Data Engineer Dallas, Texas

- Designed and implemented large-scale customer data pipelines using Azure Data Factory, Apache NiFi, and Informatica to ingest, transform, and consolidate customer profiles, transaction data, and engagement metrics from CRM, web, mobile apps, e-commerce, and third-party platforms into centralized Azure data platforms supporting email and SMS marketing initiatives
- Architected and managed Azure Data Lake Storage Gen2 and Azure Synapse Analytics environments to process high-volume customer databases supporting segmentation, personalization, and targeted marketing campaigns while ensuring GDPR and CCPA compliance with data governance, privacy, and consent management standards
- Developed real-time customer data processing pipelines using Azure Event Hubs, Azure Stream Analytics, Apache Kafka, and Apache Flink to capture digital interactions, behavioral signals, and customer engagement events for marketing automation and campaign targeting
- Engineered scalable data models and transformations using Azure Databricks (PySpark, Scala, SQL) to support Customer 360 views, audience segmentation, recommendation engines, and personalized customer communications for email and SMS marketing programs
- Integrated structured and unstructured customer data across Azure Cosmos DB, PostgreSQL, MongoDB, MySQL, and Cassandra, enabling low-latency access for email and SMS marketing platforms, customer engagement systems, and campaign execution
- Centralized customer analytics using Snowflake on Azure, delivering a single source of truth for marketing operations, CRM, and digital engagement platforms, while improving performance through partitioning and clustering strategies for large marketing datasets
- Implemented workflow orchestration and automation using Apache Airflow, enabling fault-tolerant and auditable ETL processes with data quality validation, lineage tracking, and compliance monitoring via Talend and custom Python frameworks
- Deployed containerized data services and APIs using Docker, Kubernetes, and Spring Boot, supporting secure and scalable data access for marketing technology platforms, customer communication systems, and campaign management tools
- Automated CI/CD pipelines using Azure DevOps, Jenkins, Git, and Maven, ensuring reliable, secure, and compliant deployment of customer data pipelines and marketing data solutions across environments
- Monitored data quality, troubleshoot issues, and resolved discrepancies affecting marketing campaigns, ensuring data accuracy, integrity, and availability for campaign targeting and personalization at scale

Tenet Healthcare Dallas, Texas, USA

AWS Data Engineer Dec 2023 – Aug 2024

- Architected and implemented cloud-native customer data pipelines using AWS Glue, AWS Lambda, Amazon S3, and Amazon Redshift to ingest, transform, and store large-scale customer, operational, and engagement data from multiple sources supporting marketing and analytics initiatives
- Designed near-real-time data processing solutions leveraging Amazon Kinesis, Apache Kafka, and AWS Athena to monitor customer interactions, service utilization, and engagement metrics for marketing and analytics teams
- Developed and optimized enterprise customer databases within Amazon Redshift and Amazon RDS (PostgreSQL/MySQL) to support audience segmentation, campaign targeting, and customer analytics dashboards
- Built and maintained scalable big data processing frameworks using Apache Spark, Hive, and Hadoop on Amazon EMR, enabling analytics on historical customer behavior, engagement patterns, and campaign performance trends
- Implemented robust ETL and ELT workflows using AWS Glue, Apache NiFi, Talend, and dbt, consolidating customer data from CRM systems, web platforms, mobile apps, and third-party vendors into centralized data lakes
- Delivered executive-level dashboards and self-service analytics using Amazon QuickSight, Power BI, Snowflake, and Redshift, providing visibility into customer engagement KPIs, campaign performance, and audience insights
- Established secure and automated CI/CD pipelines utilizing AWS CodePipeline, Jenkins, GitLab CI/CD, Docker, and Kubernetes, supporting reliable deployment of customer data engineering components and analytics services
- Developed high-performance SQL solutions across SQL Server, Oracle, and Redshift, including complex views, stored procedures, and query optimization strategies to meet data privacy and compliance requirements
- Supported audience insights, reporting, and experimentation through reliable data access, ensuring data accuracy and scalability for marketing programs

Tenet Healthcare **Dec 2023 - Aug 2024**
AWS Data Engineer Dallas, Texas

- Architected and implemented cloud-native customer data pipelines using AWS Glue, AWS Lambda, Amazon S3, and Amazon Redshift to ingest, transform, and store large-scale customer, operational, and engagement data from multiple sources supporting marketing and analytics initiatives
- Designed near-real-time data processing solutions leveraging Amazon Kinesis, Apache Kafka, and AWS Athena to monitor customer interactions, service utilization, and engagement metrics for marketing and analytics teams
- Developed and optimized enterprise customer databases within Amazon Redshift and Amazon RDS (PostgreSQL/MySQL) to support audience segmentation, campaign targeting, and customer analytics dashboards
- Built and maintained scalable big data processing frameworks using Apache Spark, Hive, and Hadoop on Amazon EMR, enabling analytics on historical customer behavior, engagement patterns, and campaign performance trends
- Implemented robust ETL and ELT workflows using AWS Glue, Apache NiFi, Talend, and dbt, consolidating customer data from CRM systems, web platforms, mobile apps, and third-party vendors into centralized data lakes
- Delivered executive-level dashboards and self-service analytics using Amazon QuickSight, Power BI, Snowflake, and Redshift, providing visibility into customer engagement KPIs, campaign performance, and audience insights
- Established secure and automated CI/CD pipelines utilizing AWS CodePipeline, Jenkins, GitLab CI/CD, Docker, and Kubernetes, supporting reliable deployment of customer data engineering components and analytics services
- Developed high-performance SQL solutions across SQL Server, Oracle, and Redshift, including complex views, stored procedures, and query optimization strategies to meet data privacy and compliance requirements
- Supported audience insights, reporting, and experimentation through reliable data access, ensuring data accuracy and scalability for marketing programs

Accenture **Nov 2021 - Aug 2023**
Data Engineer Chennai, India

- Designed and maintained scalable batch and real-time data pipelines to ingest, transform, and process large volumes of customer, transactional, and operational datasets from enterprise applications and external platforms
- Developed large-scale data processing workflows using Apache Spark, Python, and SQL, enabling reliable transformation and aggregation of high-volume datasets for analytics and reporting
 - Implemented enterprise data integration solutions connecting multiple data sources including CRM systems, web platforms, transactional databases, and third-party data providers into centralized data platforms
 - Built and optimized ETL and ELT pipelines using distributed data processing frameworks and data orchestration tools to support scalable analytics and business intelligence workloads
 - Designed and managed data warehouse environments enabling high-performance analytical queries, customer analytics, and campaign reporting across large datasets
 - Developed robust data models, transformations, and aggregation layers to support business intelligence tools and reporting platforms used by analytics and marketing teams
 - Automated data pipeline deployment and infrastructure management using Terraform, CI/CD pipelines, and Git-based workflows, improving reliability and release efficiency
 - Optimized data processing performance and query efficiency across large datasets, ensuring reliable and timely data availability for enterprise analytics and decision-making

Tata Consultancy Services **Jul 2019 - Nov 2021**
Data Engineer Chennai, India

- Designed, built, and maintained highly scalable batch and streaming data pipelines using Apache Spark, Hadoop, and enterprise ETL frameworks, supporting large volumes of customer, operational, and transactional datasets
- Developed and optimized cloud-based and on-premise data integration solutions leveraging Azure Synapse Analytics, SQL Server, and PostgreSQL, enabling reliable customer analytics, reporting, and data consumption
- Implemented real-time data ingestion and processing architectures utilizing Apache Kafka and Azure Stream Analytics, ensuring low-latency processing of customer engagement and transactional data
- Engineered robust ETL workflows and orchestration frameworks using Python, PySpark, Scala, Apache Airflow, and Azure Data Factory, integrating structured and unstructured customer data from MongoDB, Cassandra, and relational sources into centralized data platforms
- Built and supported enterprise analytics and reporting ecosystems using Databricks, Apache Hive, Power BI, and Tableau, while enabling secure deployments through CI/CD pipelines (Jenkins, Git, Maven) and containerized platforms using Docker and Kubernetes
- Collaborated with cross-functional teams to translate business requirements into data solutions, ensuring data quality and performance

Education

Texas Tech University **Aug 2023 - May 2025**
Computer and Information Science Master's

Skills

Programming & Query Languages

SQL (Advanced), PySpark, Python, Scala, Java, Bash

Database Management (SQL & NoSQL)

Azure SQL Database, PostgreSQL, MySQL, SQL Server, Oracle, Cosmos DB, MongoDB, Cassandra, DynamoDB, Relational & NoSQL Databases

Data Engineering & ETL/ELT

Azure Data Factory, Azure Databricks, Azure Synapse Analytics, Metadata-driven pipelines, Apache NiFi, Informatica, Talend, dbt, SSIS, AWS Glue, Apache Airflow

Big Data & Distributed Processing

Apache Spark, Azure Databricks, Hadoop, Hive, HDFS, Presto, Amazon EMR

Streaming & Real-Time Processing

Azure Event Hubs, Azure Stream Analytics, Apache Kafka, Spark Streaming, Apache Flink, AWS Kinesis

Data Warehousing & Analytics

Azure Synapse Analytics, Snowflake, Amazon Redshift, Teradata, Vertica

Cloud Platforms

Azure (Synapse, Data Factory, Databricks, ADLS/Storage, Azure SQL Database, Cosmos DB, Logic Apps, Azure Functions, Event Hubs), AWS (S3, Redshift, Glue, Lambda, EMR, Athena, Kinesis, DynamoDB), Cloud-native data platforms

Workflow Orchestration & Automation

Azure Data Factory, Apache Airflow, Azure Logic Apps, Terraform, CloudFormation, Ansible, AWS Step Functions

DevOps, CI/CD & Containers

Docker, Kubernetes, Jenkins, GitHub Actions, GitLab CI/CD, Azure DevOps, AWS CodePipeline

BI & Data Visualization

Power BI, Tableau, Looker, Amazon QuickSight

Data Governance & Compliance

Data Governance, Data Quality Management, GDPR, CCPA, Consent Management, Data privacy

Version Control & Collaboration

Git, GitHub, GitLab, Bitbucket