

PROFESSIONAL SUMMARY

Highly skilled **Data Engineer** with 4+ years of hands-on experience creating, optimizing, and maintaining scalable data pipelines and cloud-based solutions. I thrive on turning complex datasets into actionable insights using modern technologies like **AWS, Azure, Spark, and Python**. I'm passionate about streamlining workflows, automating **ETL** processes, and ensuring data reliability across distributed systems. I enjoy collaborating with teams to leverage data for innovation, predictive analytics, and smarter business decisions in fast-paced environments.

TECHNICAL SKILLS

- **Programming & Scripting:** Python, SQL, Scala, Shell Script
 - **Big Data Frameworks:** Apache Spark, Hadoop, Kafka, Hive
 - **ETL & Orchestration:** Airflow, AWS Glue, Azure Data Factory, Talend
 - **Cloud Platforms:** AWS (S3, Redshift, Lambda, EMR, IAM), Azure (Synapse, Databricks, Data Lake)
 - **Databases & Warehousing:** Snowflake, PostgreSQL, MySQL, SQL Server, MongoDB
 - **Data Visualization:** Power BI, Tableau
 - **Version Control & CI/CD:** Git, GitHub Actions, Jenkins, AWS CodePipeline
 - **Monitoring & Logging:** CloudWatch, Grafana, Prometheus
 - **Other Tools:** Docker, Kubernetes, Confluence, Jira, Agile
-

PROFESSIONAL EXPERIENCE

IDS | Atlanta, GA

Data Engineer | Oct 2023 - Present

- Designed and deployed data ingestion pipelines using **Azure Data Factory, Azure Functions, and Python**, reducing data latency by 35%.
 - Optimized ETL workflows with **PySpark and Azure Data Factory**, decreasing data processing time from 6 hours to 2.5 hours.
 - Engineered an **Azure Synapse Analytics** data warehouse with dynamic partitioning, improving query performance by 40%.
 - Developed robust CI/CD pipelines for data workflows using **Azure DevOps**, increasing deployment reliability.
 - Implemented data validation frameworks with **Great Expectations**, achieving 99.8% data accuracy.
 - Automated schema evolution and metadata management in **Azure Data Lake Storage (ADLS)** using Azure Data Factory and Python.
 - Collaborated with data scientists to integrate **ML feature stores** into training pipelines, accelerating model development.
 - Migrated legacy on-prem SQL Server systems to **Azure Synapse or Snowflake on Azure**, reducing infrastructure costs by 25%.
 - Enhanced data governance and lineage tracking using **Azure Purview**, improving audit readiness.
 - Partnered with BI teams to create **Power BI** dashboards, enabling faster data-driven decisions by 20%.
 - Implemented **role-based access control (RBAC)** across Azure environments to ensure SOC 2 compliance.
 - Monitored and optimized pipeline performance using **Azure Monitor, Grafana, and Prometheus**, improving system uptime by 15%.
-

Purevisitx | Austin, TX

Data Engineer | Jan 2023 - Sep 2023

- Built and maintained **data pipelines** using **Apache Spark, Kafka, and Airflow** to process large-scale datasets efficiently.
 - Developed **ETL frameworks in Python** for real-time data ingestion from **Kafka streams** into **Azure Synapse Analytics**.
 - Improved **data model efficiency** by redesigning schemas in **Snowflake**, reducing storage usage by 30%.
 - Integrated **Azure Data Factory** pipelines with **Databricks notebooks**, enabling scalable batch processing.
 - Implemented **Delta Lake architecture** to ensure data consistency and support incremental data loads.
 - Automated **data quality checks** using **PyDeequ**, improving trust in downstream analytics and reporting.
 - Created **monitoring dashboards** using **Grafana and Prometheus** to track pipeline performance metrics.
 - Assisted in migrating **ETL jobs from legacy systems to PySpark**, enhancing maintainability and runtime efficiency.
 - Supported BI teams in building semantic models and dashboards using **Power BI and Tableau** for actionable insights.
 - Documented data engineering processes in **Confluence**, improving knowledge sharing and onboarding for new team members.
 - Collaborated with **DevOps and QA teams** to ensure robust, fault-tolerant data pipelines and infrastructure.
 - Participated in **Agile sprints** using **Jira**, contributing to timely delivery of data engineering tasks.
-

DBS Bank | Mumbai, India

Junior Data Engineer | Sep 2019 – Sep 2020

- Hands-on experience with **Python, SQL, Scala, and Shell scripting** for automating financial data pipelines and workflow optimization.
- Knowledge of big data frameworks like **Apache Spark, Hadoop, Kafka, and Hive** for processing large-scale **transactional and risk-related datasets**.

- Experience building **ETL pipelines** using **Apache Airflow**, **AWS Glue**, and **Talend** to integrate data from multiple banking systems.
- Skilled in **AWS ecosystem** — **S3** for secure data storage, **Redshift** for warehousing, **Lambda** for serverless compute, and **EMR** for high-volume data processing.
- Exposure to **data modeling, normalization**, and **query optimization** across **Snowflake**, **PostgreSQL**, **MySQL**, **SQL Server**, and **MongoDB** for financial analytics.
- Strong understanding of **data security, IAM policies, and encryption** practices within **AWS** to meet banking compliance standards.
- Proficient in **data visualization** using **Power BI** and **Tableau** for building dashboards that track KPIs, credit risk, and financial performance.
- Familiar with **CI/CD pipelines** using **Git**, **GitHub Actions**, **Jenkins**, and **AWS CodePipeline** for automated deployments.
- Knowledge of **monitoring and logging** with **CloudWatch**, **Grafana**, and **Prometheus** for maintaining high data reliability and traceability.
- Strong collaboration and agile mindset using **Confluence**, **Jira**, and **Agile methodologies** for delivering data projects in cross-functional banking teams.

INTERNSHIP

Massachusetts General Hospital | Cambridge, MA

Data Science Intern | Jan 2022 - Sep 2022

- Completed comprehensive **HIPAA compliance** and **PII data handling** training, ensuring strict adherence to healthcare privacy and security regulations.
- Analyzed **Alzheimer’s disease research databases** containing sensitive patient data and collaborated with **clinical research teams** to build a secure **Drupal** website for data visualization.
- Designed and deployed a **multi-purpose Drupal site** with **role-based access controls (RBAC)** to allow authorized hospital users and researchers to view de-identified patient data.
- Implemented **security protocols, audit trails**, and **data encryption** to ensure **PHI (Protected Health Information)** remained fully compliant with **HIPAA** requirements.
- Developed **Python scripts** using **Regular Expressions (Regex)** to categorize medical diagnoses while maintaining strict data anonymization standards.

CERTIFICATIONS

- | | |
|--|---------------------|
| • Microsoft Certified Azure Data Engineer Associate | Jul 2024 - Jul 2026 |
| • Databricks Certified Data Engineer Associate | Feb 2024 - Feb 2026 |
| • Microsoft Certified Fabric Analytics Engineer Associate | Dec 2024 - Dec 2025 |
| • Microsoft Certified Fabric Data Engineer Associate | Apr 2025 - Apr 2026 |

EDUCATION

Masters in Data Analytics Engineering - Northeastern University