

Donthineni Jaswanth

Data Engineer

USA (Open to Relocate) | +1 (945) 268-5288 | jaswanthdonthi123@gmail.com | [LinkedIn](#)

SUMMARY

Senior Data Engineer with 5+ years of experience building and scaling cloud-native data platforms on AWS and Azure. Strong expertise in PySpark, Spark SQL, Glue, Databricks, Snowflake, and Redshift, processing terabyte-scale datasets in production environments. Proven track record in designing robust data models, optimizing ETL/ELT pipelines, and implementing monitoring, CI/CD, and data quality frameworks. Experienced in supporting analytics and machine learning workloads with high reliability, performance, and governance in enterprise-grade systems.

EXPERIENCE

Data Engineer | CVS Health, USA

June 2025 – Present

- Architected and owned AWS-based data platforms (Glue, S3, Redshift, Snowflake) processing 5+ TB/day, improving ingestion scalability by 30% and reducing end-to-end latency from 6 to 4 hours.
- Developed PySpark, Spark SQL, and ELT frameworks supporting structured, semi-structured, and unstructured data, enabling 40% higher throughput for analytics and ML pipelines.
- Designed enterprise-grade dimensional and analytics data models in Redshift & Snowflake, reducing BI and ad-hoc query time by 60%.
- Implemented data quality, observability, and monitoring frameworks using CloudWatch, alerts, and SLA checks, reducing pipeline failures by 35%.
- Partnered with Data Science & Analytics teams to deliver curated datasets powering predictive models, contributing to 15% increase in marketing ROI.
- Automated Apache Airflow orchestration and CI/CD pipelines (Git, IaC), cutting manual deployment effort by 25% and enabling zero-downtime releases twice per week.
- Ensured HIPAA-compliant data handling, access control, and audit readiness across cloud data assets.

Data Engineer | Accenture, India

June 2020 – July 2023

- Designed and delivered cloud-scale ETL/ELT pipelines using Azure Data Factory, Databricks, Synapse, increasing enterprise data availability by 30%.
- Built high-performance Spark and SQL transformations handling terabyte-scale datasets, reducing query execution times by 25%.
- Implemented data warehousing and semantic models supporting 10+ business stakeholders, accelerating decision-making across finance, supply chain, and operations.
- Led data governance, lineage, and metadata management using Azure Purview, reducing data quality incidents by 20% and improving audit compliance.
- Collaborated with data scientists and analysts to ensure reliable, analytics-ready datasets for forecasting and reporting initiatives.
- Spearheaded legacy on-prem to Azure cloud migration, implementing CI/CD and cost-optimized architecture, reducing infrastructure and maintenance costs by 15%.

Jr. Data Engineer | ICICI Bank, India

Jan 2019 – May 2020

- Built and supported enterprise ETL pipelines using SQL, Python, and Informatica, integrating data from ERP and CRM systems, improving availability by 30%.
- Performed data cleansing, validation, and standardization, reducing reporting errors by 25% and ensuring consistent analytics outputs.
- Assisted in data modeling and schema optimization, improving BI query performance and reducing data retrieval times by 20%.
- Supported cloud migration initiatives (AWS & Azure), improving system scalability and cutting report generation time by 35%.
- Monitored and resolved pipeline issues, maintaining 99% data accuracy and compliance with banking governance standards.

SKILLS

- ETL/ELT Tools & Frameworks:** AWS Glue, Azure Data Factory, Databricks, Informatica, Apache Spark, Airflow, Jenkins, GitLab, Azure DevOps
- Programming Languages:** Python, SQL, PySpark
- Cloud Platforms:** AWS (Redshift, S3, Glue, CloudWatch), Azure (Synapse, Data Lake, Purview, HDInsight)
- Data Warehousing & Database:** Redshift, Snowflake, SQL Server, MySQL, PostgreSQL
- Big Data Technologies:** Apache Hadoop, Databricks, Kafka, Spark Streaming
- Data Modeling & Architecture:** Star Schema, Snowflake Schema, Dimensional Modeling, Schema Optimization
- Data Integration & Transformation:** Structured, Semi-structured, Unstructured Data, API Integration, Data Wrangling, Data Cleansing, Data Validation
- Reporting & Analytics:** Power BI, Tableau, BI Dashboards, Predictive Analytics, Business KPIs
- CI/CD & Automation:** Airflow, Pipeline Orchestration, CI/CD Automation (Jenkins, GitLab), Version Control (Git, GitHub)
- Data Governance & Compliance:** Azure Purview, Metadata Management, Data Lineage
- Collaboration & Methodologies:** Cross-functional Collaboration, Agile, JIRA, Stakeholder Communication

EDUCATION

Master in Advanced Data Analytics

University of North Texas, Denton, Texas, USA

Aug 2023 – May 2025

Bachelor of Technology (B Tech) in Computer Science Engineering

Chebrolu Engineering College - JNTU - Kakinada, India

Jun 2015 – May 2019

PROJECTS

Enterprise Cloud Data Platform Modernization (AWS)

- **Overview:** Designed and implemented a scalable, cloud-native data platform to ingest, transform, and serve analytics-ready data for enterprise BI and ML use cases.
- Architected end-to-end AWS data pipelines (S3, Glue, PySpark, Redshift, Airflow) processing 5+ TB/day, reducing data latency by 30% and improving pipeline scalability and reliability.
- Implemented data quality checks, monitoring, and automated CI/CD deployments, reducing pipeline failures by 35% and enabling faster, zero-downtime releases for analytics teams.

Analytics & Reporting Optimization Using Snowflake and Spark

- **Overview:** Optimized analytics performance and reporting efficiency by redesigning data models and improving transformation workflows for large-scale datasets.
- Designed dimensional and analytics data models in Snowflake and Redshift, accelerating BI dashboards and ad-hoc queries, reducing report generation time by 60%.
- Built and optimized Spark SQL and PySpark ELT workflows for structured and semi-structured data, improving query performance by 25% and enabling faster decision-making for business stakeholders.