

Steven Mitchell

Spanish Fork, UT 84660 | (385) 437-6742 | stevenml9883@gmail.com | www.linkedin.com/in/steven-mitchell-487a01376

Summary

Accomplished Data Engineer with 9 years of experience, proven track record and strong expertise in Azure Data Engineering PaaS offerings, namely Azure Data Factory, Databricks, Azure Storage, Azure SQL DB, Azure Synapse Analytics and Power BI. Expertise in Enterprise Data Warehouse (EDW), ETL, Reporting, OLAP, OLTP processes with extensive exposure in Azure Data Factory and Databricks.

Education

M.S. IN COMPUTER SCIENCE | APR 2015 | UTAH VALLEY UNIVERSITY

B.S. IN COMPUTER SCIENCE | MAY 2013 | UTAH VALLEY UNIVERSITY

Skills

- **Cloud Platforms & Services:** Azure Data Factory, Azure Databricks, Azure SQL Database, Azure Synapse Analytics, Azure DevOps, Event Hubs
- **Databricks:** Delta Lake, Medallion Architecture, Administration, Delta Live Tables, Databricks Structured Streaming
- **Programming & Scripting:** Python, SQL, PySpark, Scala
- **ETL & Data Engineering:** ETL Pipelines, Data Mesh Architecture, Airflow, Data Quality Frameworks, REST API Integration, Real-Time Streaming, Power BI.
- **Security & Governance:** Unity Catalog, Fine-Grained Access Control, SOX Compliance
- **DevOps & Automation:** CI/CD Pipeline

Experience

SENIOR DATA ENGINEER | VERCEL | JULY 2023 – PRESENT

- Engineered enterprise-scale ETL pipelines using Python, Scala, and Azure Databricks Runtime 13.3 LTS ML with Delta Live Tables, achieving 90% infrastructure cost savings and cross-region Unity Catalog governance for 5TB+ daily data processing.
- Developed production-grade Apache Spark applications on Azure Databricks using PySpark and Scala, implementing 200+ data pipelines with Common Data Model standardization for master data management, enabling data mesh architecture across multiple business domains with automated Azure DevOps CI/CD pipelines
- Implemented real-time streaming architecture using Event Hubs Premium and Databricks Structured Streaming, processing 1M+ retail transaction events daily with Change Data Feed integration for delta lake tables, achieving sub-second latency and 99.9% data accuracy across US and EMEA regions
- Built robust Python and Scala monitoring framework on Azure Databricks for Delta Lake transaction logs using Unity Catalog's fine-grained access control, automating operational metrics collection through REST APIs and reducing incident resolution time by 70% while maintaining SOX compliance

- Optimized Azure Databricks cluster configurations (Standard_D16s_v5) and Apache Spark jobs using advanced Scala performance tuning and Photon engine, resulting in 40% cost reduction through automated job clustering and pipeline scheduling managed via Azure DevOps across development, QA, and production environments
- Architected end-to-end data quality framework using PySpark and Unity Catalog's metadata API, integrating 100+ automated quality checks within Azure Data Factory pipelines and Azure DevOps release workflows, maintaining 15-minute SLAs for critical retail business tables while enabling self-service data discovery
- Engineered Azure Data Factory/Databricks ETL transformation, slashing 50TB daily processing by 30% across 20+ systems. Architected zero-downtime pipelines powering instant analytics with guaranteed fault tolerance.
- Reengineered Azure Synapse queries, slashing finance reporting 40% and cutting 3+ hours weekly. Delivered instant financial analytics through optimized data models.

DATA ENGINEER | STRIPE | JAN 2018 – JUN 2023

- Optra Data Accelerator (ODA), an Azure Databricks Lakehouse solution, achieved 90% infrastructure cost savings and 70% developer effort reduction through metadata-driven Big Data ingestion and refinement while overseeing product architecture, design, and adoption strategy
- Engineered metadata-driven transformation capabilities leveraging Spark SQL and PySpark, implementing a custom data quality library and comprehensive Delta Lake operational metrics store that enabled real-time monitoring, alerting, and end-to-end pipeline observability
- Developed **Terraform modules for Azure Synapse Analytics**, automating workspace provisioning and access control, reducing manual setup efforts by 60%. Integrated **Terraform with Azure Key Vault** to securely manage secrets, API keys, and database credentials for data engineering pipelines.
- Developed enterprise integrations, including Dynamics365 CRM connector customization using Microsoft's open-source Spark Common Data Model and Event Hubs ingestion with Spark Structured Streaming, enhancing data processing capabilities across diverse sources
- Architected Unity Catalog integration for advanced data governance while optimizing 6000 plus lines of PySpark codebase
- Worked with 5 Azure Databricks developers, delivering 100+ production-grade notebooks across 45+ sprints with zero post-production defects while maintaining strict quality standards
- Provided technical support during product adoption for 2 enterprise products, including advanced troubleshooting, debugging, and root cause analysis for complex data engineering workloads, ensuring seamless implementation and optimal performance across 8 major modules, implementing robust data privacy handling in Azure Databricks environment
- Engineered Azure Data Factory ETL architecture, transforming 100TB data processing with 30% faster pipeline execution. Automated real-time ingestion from 25+ SAP modules while maintaining zero data loss.
- Optimized complex Azure SQL queries and indexing strategies, accelerating business intelligence reporting by 25%. Reduced monthly Azure costs \$50K through strategic data warehouse migration.

DATA ANALYST | ALTAIR | MAY 2015 – DEC 2017

- Optra Data Accelerator (ODA), an Azure Databricks Lakehouse solution, achieved 90% infrastructure cost savings and 70% developer effort reduction through metadata-driven Big Data ingestion and refinement while overseeing product architecture, design, and adoption strategy
- Engineered metadata-driven transformation capabilities leveraging Spark SQL and PySpark, implementing a custom data quality library and comprehensive Delta Lake operational metrics store that enabled real-time monitoring, alerting, and end-to-end pipeline observability

- Developed **Terraform modules for Azure Synapse Analytics**, automating workspace provisioning and access control, reducing manual setup efforts by 60%. Integrated **Terraform with Azure Key Vault** to securely manage secrets, API keys, and database credentials for data engineering pipelines.
- Developed enterprise integrations, including Dynamics365 CRM connector customization using Microsoft's open-source Spark Common Data Model and Event Hubs ingestion with Spark Structured Streaming, enhancing data processing capabilities across diverse sources
- Architected Unity Catalog integration for advanced data governance while optimizing 6000 plus lines of PySpark codebase
- Worked with 5 Azure Databricks developers, delivering 100+ production-grade notebooks across 45+ sprints with zero post-production defects while maintaining strict quality standards
- Provided technical support during product adoption for 2 enterprise products, including advanced troubleshooting, debugging, and root cause analysis for complex data engineering workloads, ensuring seamless implementation and optimal performance across 8 major modules, implementing robust data privacy handling in Azure Databricks environment
- Engineered Azure Data Factory ETL architecture, transforming 100TB data processing with 30% faster pipeline execution. Automated real-time ingestion from 25+ SAP modules while maintaining zero data loss.
- Optimized complex Azure SQL queries and indexing strategies, accelerating business intelligence reporting by 25%. Reduced monthly Azure costs \$50K through strategic data warehouse migration.

Certifications & Trainings

- Microsoft Certified: Azure Fundamentals (AZ-900)
- Microsoft Certified: Azure Data Engineer Associate (DP-203)
- Databricks Certified Data Engineer Associate