

Daniel Rowell

Orange City, FL 32763 | (347) 860-9591 | drowellde@outlook.com

PROFESSIONAL SUMMARY

Senior Data Engineer with 10+ years of experience designing scalable cloud data platforms and high-volume data pipelines supporting analytics, machine learning, and real-time data products. Expert in distributed data processing, modern lakehouse architectures, and enterprise ETL/ELT frameworks. Strong experience with Spark, Kafka, Airflow, Databricks, Snowflake, and cloud-native architectures across AWS, Azure, and GCP. Proven track record delivering reliable data platforms, optimizing large-scale data pipelines, and collaborating with analytics and ML teams to enable data-driven decision making.

TECHNICAL SKILLS

- **Languages:** Python, SQL, Scala, Java, Bash
- **Big Data:** Apache Spark, PySpark, Hadoop, Kafka, Spark Structured Streaming, Apache Flink
- **Data Platforms:** Databricks, Snowflake, Delta Lake, Apache Iceberg, Redshift, BigQuery
- **Workflow Orchestration:** Apache Airflow, Prefect, Dagster
- **Data Modeling:** Dimensional Modeling, Data Vault, Star/Snowflake Schema, SCD Type-2
- **Cloud Platforms:** AWS (S3, EMR, Glue, Lambda, Redshift, Kinesis), Azure (ADLS, Synapse, Data Factory), GCP (BigQuery, Dataflow, Cloud Storage)
- **DevOps & Infrastructure:** Docker, Kubernetes, Terraform, CI/CD, GitHub Actions, Jenkins
- **Data Quality & Governance:** Great Expectations, Data Lineage, Metadata Management, Data Catalogs
- **BI & Analytics:** Tableau, Power BI, Looker

PROFESSIONAL EXPERIENCE

Slalom

Mar 2022 - Present

Senior Data Engineer

Miami, FL

- Architected enterprise data platforms for Fortune 500 retail, healthcare, and financial services clients using Data Vault modeling, CI/CD pipelines, and Redshift, improving data accessibility and governance across the organization
- Designed modern lakehouse architectures with Databricks, Delta Lake, and cloud object storage, enabling faster analytics queries and unified data access
- Constructed large-scale Spark and PySpark pipelines that process billions of records daily, reducing batch processing time and supporting near-real-time reporting
- Deployed real-time streaming pipelines with Kafka and Spark Structured Streaming, achieving sub-second data latency for critical dashboards
- Led migration of legacy ETL pipelines to cloud-native ELT frameworks using Dagster and Hadoop, cutting maintenance effort and increasing pipeline reliability
- Engineer scalable ingestion frameworks with Dagster pipelines, CI/CD automation, and Redshift/Hadoop storage, delivering API, SaaS, relational, and event data to the data lake and enabling downstream analytics teams to access fresh data within minutes
- Orchestrate data workflows in Apache Airflow using event-driven patterns, reducing pipeline latency and improving reliability for nightly data loads
- Establish automated data-quality validation frameworks with Great Expectations, catching data anomalies before they reach production and maintaining data trust
- Collaborate with data scientists to design feature pipelines in Spark and Data Vault, delivering ready-to-use features for machine-learning models and shortening model-training cycles
- Optimize Spark workloads by applying partitioning, caching, and query-optimization techniques, cutting job runtimes and lowering compute costs
- Model enterprise data structures using Data Vault methodology and Data Catalog tools to enable scalable BI and analytics workloads, improving data accessibility for business users
- Implement CI/CD pipelines for data platform deployment using Terraform and GitHub Actions, reducing deployment time and minimizing manual errors
- Deploy monitoring and observability frameworks with Dagster and Prometheus for large-scale data pipelines, providing real-time alerts that improve issue-resolution speed
- Mentor junior engineers and conduct code reviews via GitHub pull requests, raising code quality and accelerating onboarding of new team members
- Lead architecture discussions with stakeholders and define enterprise data standards using Data Vault and Data Catalogs, aligning data practices across the organization
- Enhance pipeline reliability and reduce execution time by 40% through targeted performance optimization.

NielsenIQ

Mar 2018 - Feb 2022

Senior Data Engineer

Chicago, IL

- Engineered scalable ETL pipelines in BigQuery and Azure Data Factory that processed large consumer analytics datasets, reducing data load time and supporting daily reporting

- Designed distributed data processing frameworks with Spark, Python, and SQL on Azure Synapse, enabling analysts to run complex queries faster
- Architected ingestion pipelines using Azure ADLS and Delta Lake to capture both structured and semi-structured data, improving data availability for downstream models
- Orchestrated Kafka streaming pipelines enabling near real-time analytics and insights.
- Modeled dimensional data structures, building star schemas and SCD Type 2 tables in BigQuery, which supported enterprise reporting and reduced report generation errors
- Coordinated complex workflows using Apache Airflow, ensuring reliable and automated pipeline execution.
- Tuned data warehouse queries to enhance dashboard performance and accelerate data retrieval.
- Integrated datasets from APIs, flat files, relational databases, and cloud storage systems using Python scripts and SQL, enabling unified data access for downstream analytics
- Constructed large-scale batch ingestion pipelines with AWS Glue and Spark to handle multi-terabyte datasets.
- Leveraged metadata-driven ETL frameworks to reduce development time and improve team efficiency.
- Partnered with analytics and BI teams to deliver enterprise reporting dashboards.
- Applied partitioning and indexing strategies to optimize query performance in Snowflake and Redshift.
- Designed SQL transformations to support Snowflake and Redshift data warehouses for reporting.
- Established automated testing frameworks to ensure pipeline reliability and data accuracy.

Segment

Jun 2015 - Feb 2018

Data Engineer

Austin, TX

- Engineered scalable event-driven pipelines using Dagster and Azure Data Factory to ingest and process high-volume customer interaction data from web, mobile, and backend systems, enabling near-real-time analytics for product teams
- Created Python ingestion services with GitHub Actions CI/CD that integrated REST APIs, SDK event streams, and third-party SaaS sources, delivering reliable data feeds to the data lake
- Designed distributed Spark and PySpark pipelines on Azure Synapse to transform and enrich behavioral data, providing clean datasets for downstream analytics platforms
- Deployed Kafka-based streaming pipelines orchestrated with Apache Airflow to process events in real time, improving customer activity tracking accuracy
- Constructed ETL workflows that routed event data into Amazon Redshift and Snowflake, consolidating data for reporting and reducing manual data preparation
- Architected scalable data schemas in Snowflake using SQL to enable customer analytics, attribution modeling, and product-usage insights, improving data accessibility for analysts
- Orchestrated automated workflows with Apache Airflow and GitHub Actions to manage ingestion, transformation, and loading pipelines, reducing manual effort and ensuring timely data delivery
- Established data-validation and monitoring frameworks in Dagster and Azure Data Factory to detect anomalies and maintain data integrity, enabling early detection of data-quality issues
- Collaborated with product engineers to implement event tracking across web and mobile applications using Azure ADLS and Snowflake, providing comprehensive usage data for product decisions
- Optimized large-scale Spark workloads by applying partitioning, caching, and Parquet file formats, lowering processing time and resource consumption
- Authored reusable Python libraries for ingestion and transformation pipelines, adopted across multiple teams.
- Enabled analytics teams with clean, structured datasets powering BI dashboards and product metrics.
- Supported scaling of the cloud-based data infrastructure to accommodate rapid startup growth.
- Partnered with analytics engineers and data scientists to create datasets facilitating experimentation and product analytics.

EDUCATION

University of Florida

Bachelor of Science, Computer Science

Aug 2011 - May 2015

Gainesville, FL