

Adrian M.

Senior Data Engineer | Data Platform Architect • Big Data & AI Systems • Real-Time Data Pipelines • Multi-Cloud Platforms

✉ adrian.datacore1@gmail.com

☎ (806) 279-1513

📍 Austin, TX 78701, USA

🌐 [linkedin.com/in/adriandataworks/](https://www.linkedin.com/in/adriandataworks/)

SUMMARY

I'm a Senior Lead Data Engineer with 9+ years of experience designing and operating cloud-native, multi-cloud data platforms across healthcare, AI, and enterprise environments. I build high-throughput, fault-tolerant batch and streaming ETL/ELT pipelines that power analytics, real-time decisioning, and AI-ready workloads, leveraging Kafka, Spark, Flink, Delta Lake, dbt, Databricks, Snowflake, Airflow, Fivetran, Prefect, and modern lakehouse solutions (Iceberg, Hudi). Skilled across AWS, Azure, and GCP, I deliver scalable, automated, cost-optimized platforms supporting BI, ML pipelines, GenAI applications, RAG/vector search, and self-service analytics, while ensuring HIPAA, SOC2, GDPR compliance and integrating healthcare interoperability standards (FHIR, HL7). My work spans cloud migrations, platform modernization, lakehouse adoption, and AI/ML orchestration. I lead through hands-on technical ownership, guiding teams via architecture design, system trade-offs, and best practices, driving adoption of data mesh, DevSecOps, and MLOps principles. My platforms emphasize performance, observability, reliability, and actionable analytics, delivering measurable impact in production environments.

SKILLS

Data Engineering & Orchestration: Batch & real-time ETL/ELT (dbt Core & Cloud, PySpark, Airflow, Dagster, Apache NiFi, Talend, Fivetran, Airbyte, Matillion), Change Data Capture (CDC), Event-Driven Design, Data Mesh, Data Fabric, Data Contracts

Programming & Scripting: Python, SQL/T-SQL, Scala, Java, Bash, YAML, JSON, REST, R, Postgres, Node.js, React.js, Angular

BI, Analytics & AI Enablement: Power BI, Tableau, Looker, SQL Reporting, KPI Dashboards, Self-Service Analytics, Real-Time Dashboards, ML Pipeline Integration

Big Data & Real-Time Streaming: Apache Kafka (Streams, KSQL, Kinesis), Apache Flink, Apache Beam, Pub/Sub, Hadoop, Delta Lake, Iceberg, Hudi, ClickHouse, Druid, DuckDB, BigQuery, Redshift, Snowflake, Redpanda, Trino, Presto

Cloud & Data Platforms: AWS (S3, Lambda, Glue, Redshift), Azure (Synapse, Data Lake, Databricks, AKS), GCP (BigQuery, Dataflow), Snowflake, Epic Clarity, Palantir Foundry, Multi-Cloud Deployments, Cross-Region Replication, Cloud Cost Optimization, Unity Catalog, Data Lineage (OpenLineage, Marquez)

Security, Compliance & Governance: HIPAA, SOC2, GDPR, IAM, Encryption, Tokenization, HashiCorp Vault, Access Control, Role-Based Policies, SIEM (Splunk, ELK), Compliance Reporting, Risk Management

Observability & Monitoring: Prometheus, Grafana, Datadog, Monte Carlo, Metadata Management, Operational Metrics, Alerting & Anomaly Detection, Logging & Tracing

Data Architecture & Modeling: Star Schema, Snowflake Schema, Dimensional Modeling, Data Vault, High-Volume Data Processing, Data Pipeline Optimization, Real-Time Analytics, Fault-Tolerant Design

AI & Machine Learning: ML Pipeline Development, MLOps, Databricks Feature Store, MLflow, Model Deployment, Real-Time Scoring, Retrieval-Augmented Generation (RAG), Vector Databases (Pinecone, FAISS), GenAI Workflows

DevOps & Infrastructure Automation: CI/CD (GitHub Actions, Jenkins, Azure DevOps), Terraform, Docker, Kubernetes (AKS, EKS, GCP GKE), Helm, Kubernetes Operators, Cloud Cost Management

Leadership & Collaboration: Agile/Scrum, Cross-Functional Leadership, Strategic Planning, Architecture Reviews, Mentorship, Standards & Guidelines, Stakeholder Engagement

PROFESSIONAL EXPERIENCE

Lead Data Engineer

10/2023 – Present

ProCogia

- Built high-volume, HIPAA-compliant data platforms on Azure Databricks and Snowflake, integrating FHIR/HL7 for secure clinical data exchange, processing millions of patient records daily, reducing SLA breaches by 45%.
- Implemented real-time ETL/ELT pipelines using PySpark, Delta Lake, Kafka, Flink, and dbt, powering AI/ML workflows, operational dashboards, and analytics triggers across enterprise healthcare systems.
- Modernized multi-cloud platform with AWS, Azure, and GCP; automated CI/CD with GitHub Actions and Terraform, reducing deployment cycles by 60% and optimizing cloud costs.
- Delivered AI-powered dashboards and reporting tools with Power BI, Tableau, and Looker, enabling self-service analytics and reducing manual reporting by 60%.
- Optimized data governance and compliance using Unity Catalog, IAM, encryption, tokenization, and automated lineage, ensuring HIPAA, SOC2, and GDPR adherence.
- Scaled high-performance lakehouse architecture using data mesh and MLOps principles, enabling enterprise-wide analytics and AI workflows with measurable operational impact.
- Mentored a team of 6 engineers, establishing architecture standards, dbt testing frameworks, and data contracts, improving code quality, team efficiency, and cross-domain collaboration.

Senior Cloud Data Engineer

01/2020 – 09/2023

VentureDive

- Engineered cloud-native analytics platforms on AWS, Azure, and GCP, integrating Snowflake to process millions of enterprise records daily, enabling scalable reporting and AI/ML workflows.
- Developed real-time and batch ETL/ELT pipelines using PySpark, Airflow, dbt, and Kafka, optimizing ingestion, transformation, and orchestration across multi-cloud systems.
- Implemented feature store solutions on Databricks, accelerating ML model training and deployment, reducing data latency by 50%.
- Led cloud migration and modernization initiatives with automated CI/CD using GitHub Actions and Terraform, achieving 40% faster release cycles and improved platform reliability.
- Delivered self-service dashboards and KPI monitoring using Power BI, Tableau, and Looker, cutting manual reporting by 70% and improving business decision speed.
- Designed scalable data architectures and optimized pipelines for high-volume processing, ensuring reliable analytics and operational triggers across enterprise domains.
- Mentored and guided junior engineers, establishing coding standards, data contracts, and best practices, improving team efficiency and cross-functional collaboration.

ETL & Data Engineer

07/2016 – 12/2019

Kanda Software

- Developed robust ETL/ELT pipelines using Airflow, dbt, and SQL to process and transform enterprise data for reporting and analytics.
- Implemented batch and incremental workflows to support KPI dashboards and operational reporting, improving data freshness by 40%.
- Optimized database performance through advanced indexing, partitioning, and caching, reducing query times by up to 50% for BI teams.
- Delivered interactive dashboards and self-service analytics using Power BI, Tableau, and Looker, enabling cross-functional teams to track metrics in near real-time.
- Strengthened data quality, security, and access controls with IAM and HashiCorp Vault, ensuring compliant data handling across enterprise pipelines.
- Collaborated with analytics, product, and operations teams to enable low-latency, actionable insights through streamlined data delivery and operational triggers.
- Designed and maintained scalable data workflows, ensuring operational reliability and consistency across multiple data domains.

EDUCATION

Bachelor's in Computer Science