

Muhammad Hafeez

Data Infrastructure & Platform Lead | Big Data & Analytics Specialist | Data Engineer

databytes.dev134@gmail.com

(347) 745-4997

Cleveland, Ohio, 44101

Experienced Data Engineer with 5+ years of expertise in designing, building, and optimizing scalable data pipelines, cloud-based data architectures, and high-performance data processing systems. Skilled in Python, SQL, PySpark, and modern data tools to transform complex, raw data into reliable, actionable datasets for analytics, reporting, and machine learning. Adept at collaborating with cross-functional teams to deliver efficient, maintainable data solutions that support strategic business decisions.

Skills

Programming Languages: Python, SQL, PySpark, Scala

Big Data & Data Engineering Tools: Apache Spark, Hadoop, Apache Kafka, Hive, Airflow, Prefect, Luigi, dbt, Spark SQL

Data Modeling: Kimball, Data Vault 2.0, Star Schema, Anchor Modeling, Slowly Changing Dimensions

Data Observability: Prometheus, Grafana, OpenTelemetry, ELK, Splunk

Data Processing & Pipelines: ETL, ELT, Data Modeling, Data Warehousing, Data Integration, Batch Processing, Stream Processing, Pipeline Optimization

Databases & Data Warehousing: PostgreSQL, MySQL, Snowflake, Redshift

Data Formats: Parquet, Avro, ORC, JSON, JSONL, Protobuf, Apache Thrift, Iceberg Metadata

Cloud Platforms & Services: AWS (S3, Redshift, Glue, Lambda, EMR), GCP (BigQuery, Dataflow)

DevOps & Workflow Tools: Docker, Kubernetes, Git, CI/CD, Terraform

Analytics & Visualization: Tableau, Power BI, Looker

Testing Tools: Great Expectations, Apache Deequ, dbt Tests, PyTest, Spark Testing Base, Monte Carlo, Talend Data Quality, Bigeye

Professional Experience

Senior Data Engineer, Tebra

01/2022 – Present

- Designed and maintained scalable ETL pipelines to ingest and unify clinical, billing, and operational healthcare data.
- Built automated workflows with Airflow for real-time and batch processing of patient, appointment, and revenue datasets.
- Developed data models to support healthcare reporting, revenue analytics, and practice performance dashboards.
- Collaborated with product, analytics, and clinical operations teams to deliver insights for improving provider workflows.
- Implemented data quality validation and monitoring frameworks to ensure compliance and accuracy of sensitive healthcare data.
- Enabled API integrations for HL7/JSON data exchange between Tebra's ecosystem and third-party systems.

Data Engineer, Streamline Analytics

06/2019 – 12/2021

- Developed and maintained robust ETL workflows for multi-source operational and transactional datasets.
- Transformed raw data into structured models to support client dashboards, KPIs tracking, and BI reporting.
- Automated data ingestion processes using Python and AWS tools to improve reliability and reduce manual tasks.
- Partnered with analysts, product owners, and stakeholders to translate business requirements into data solutions.
- Monitored data quality and pipeline health to ensure accurate delivery of analytics products to end users.

Junior Data Engineer, NextEra Analytics

08/2017 – 05/2019

- Assisted in building ETL pipelines to ingest, clean, and aggregate diverse energy operational datasets (e.g., grid data, asset data).
- Supported integration of smart meter, forecasting and sensor feeds for analytic workflows.
- Performed data quality checks and monitored pipeline outputs to ensure timely and reliable insights for energy analytics platforms.
- Collaborated with senior engineers to troubleshoot and optimize existing data transformation workflows.

Education

Bachelors Degree