

Dexter Anderson

Madison, NJ | (347) 766-7172 | dexteranderson319@outlook.com

Professional Summary

- Senior Data Engineer with **10+ years of experience** designing **scalable cloud data platforms, distributed data pipelines, and real-time streaming architectures**. Expert in building **high-performance ETL/ELT systems** using **Python, Apache Spark, Kafka, and modern cloud platforms** processing **multi-terabyte datasets daily**.
- Designed and operated **high-volume data systems processing 2.5TB+ of data daily** using Python, Spark, Kafka, and cloud architectures.
- Proven track record delivering **high-performance ETL/ELT pipelines and real-time streaming platforms**, improving **data processing performance by up to 40%** and enabling scalable analytics and machine learning workloads.

Core Technologies: Python | SQL | Apache Spark | Kafka | Airflow | AWS | Snowflake | dbt | Delta Lake | Docker | Kubernetes

Skills

- **Languages**
Python, SQL, PySpark, Scala
- **Data Processing**
Apache Spark, Kafka, Airflow, Flink
- **Cloud Platforms**
AWS (S3, Glue, Athena, Redshift, EMR, Lambda)
- **Data Warehousing**
Snowflake, Amazon Redshift, BigQuery
- **Infrastructure & DevOps**
Docker, Kubernetes, Terraform, CI/CD, Git
- **Monitoring**
Grafana, Prometheus, CloudWatch

Key Achievements

- Built distributed Spark pipelines processing **2.5TB+ of data daily**
- Reduced ETL pipeline latency by **30%**
- Improved data processing performance by **40%**
- Designed enterprise **AWS data lake architecture**
- Implemented monitoring systems improving pipeline reliability

Education

Rutgers University, New Brunswick, New Jersey

Bachelor of Science in Computer Science

Sep 2012 – May 2016

Experience

Senior Data Engineer | Feb 2022 – Present

KPMG – New York, NY

- Architected and maintained **enterprise-scale data pipelines processing 2.5TB+ of structured and unstructured data daily**, supporting analytics and machine learning workloads used by **50+ internal stakeholders**.
- Designed and implemented **scalable ETL and ELT pipelines** using **Python, Apache Spark, and Apache Airflow**, reducing pipeline latency by **30%**.
- Built **enterprise data lake architecture** on **AWS S3, Glue, and Athena**, enabling analytics-ready datasets for **BI and ML teams across 10+ business units**.
- Built **real-time streaming pipelines** using **Apache Kafka and Spark Streaming**, processing **millions of events daily**.
- Optimized Spark workloads and SQL queries through **partition tuning, caching strategies, and query optimization**, improving processing performance by **40%**.
- Implemented **automated data quality frameworks**, reducing pipeline failures and data errors by **30%**.
- Designed **CI/CD pipelines for data workflows** using Git, Docker, and Jenkins, reducing deployment time by **40%**.
- Built monitoring and observability systems using **CloudWatch and Grafana**, improving incident detection time by **35%**.
- Led architecture discussions for **scalable enterprise data infrastructure** and mentored **4+ junior engineers**.
- Implemented **data governance frameworks**, including schema validation and access control policies, improving **platform security and compliance**.
- Collaborated with **data scientists, product teams, and analysts** to design reliable datasets supporting **predictive analytics and executive dashboards**.

Data Engineer | Aug 2019 – Feb 2022

DataCore Analytics (Data & Analytics Consulting Firm) – Newark, NJ

- Designed and developed **scalable ETL pipelines integrating data from 20+ internal and external sources**.
- Built **real-time data streaming architectures** using **Kafka and Spark Streaming**, reducing data processing latency by **25%**.
- Designed **enterprise dimensional data models (star and snowflake schemas)** supporting BI platforms used by **30+ analysts**.

- Implemented **automated workflows using Apache Airflow**, reducing manual processing by **50%**.
- Developed **scalable cloud data warehouse solutions** using **Amazon Redshift and Snowflake**.
- Optimized SQL queries and warehouse schemas, improving reporting performance by **35%**.
- Built **monitoring dashboards** for pipeline health using Grafana and CloudWatch.
- Participated in **migration of legacy on-premise data platforms to AWS cloud architecture**.
- Collaborated with cross-functional teams to deliver **scalable enterprise data solutions**.

Junior Data Engineer | Jun 2016 – Aug 2019

Insight Data Systems (Private Data Engineering Consultancy) – Jersey City, NJ

- Developed **batch processing pipelines using Python and SQL** to transform enterprise datasets into centralized data warehouses.
- Built **scalable data ingestion frameworks** collecting data from APIs, databases, and flat files, improving ingestion efficiency by **25%**.
- Maintained **Hadoop and Spark clusters** supporting distributed processing across **multi-terabyte datasets**.
- Implemented **automated data validation and quality checks** across ETL pipelines.
- Optimized SQL queries and database structures, improving analytics query performance by **30%**.
- Assisted in migration of **legacy systems to AWS cloud infrastructure**.
- Authored **technical documentation and ETL development standards** used by engineering teams.
- Collaborated with engineers and analysts to improve **enterprise data accessibility and reporting**.

Additional

- Open to fully remote opportunities across the U.S.
- U.S Citizen – Authorized to work for any employer
- Available for coding challenges and interviews