

Mark Oyer

Belvidere, IL 61008 | (815) 975-6488 | markoyer0806@outlook.com

Summary

Senior Data Engineer with 10+ years of experience designing scalable data pipelines, cloud data platforms, and real-time streaming systems. Expert in Python, Apache Spark, Kafka, and cloud ecosystems including AWS, Azure, and Google Cloud. Proven ability to build enterprise data lakes, optimize large-scale data processing, and enable analytics and machine learning at scale.

Core Technologies: Python | SQL | Apache Spark | Kafka | Airflow | AWS | Snowflake | dbt | Delta Lake | Docker | Kubernetes

Core Skills

- **Languages**
Python, SQL, PySpark, Scala
- **Data Processing**
Apache Spark, Kafka, Airflow, Flink
- **Cloud Platforms**
AWS (S3, Glue, Athena, Redshift, EMR, Lambda)
- **Data Warehousing**
Snowflake, Amazon Redshift, BigQuery
- **Infrastructure & DevOps**
Docker, Kubernetes, Terraform, CI/CD, Git
- **Monitoring**
Grafana, Prometheus, CloudWatch

Key Achievements

- Built distributed Spark pipelines processing **2.5TB+ of data daily**
- Reduced ETL pipeline latency by **30%**
- Improved data processing performance by **40%**
- Designed enterprise **AWS data lake architecture**
- Implemented monitoring systems improving pipeline reliability

Experience

Senior Data Engineer | Feb 2022 – Present KPMG – New York, NY

- Architected and maintained enterprise-scale data pipelines processing 2.5TB+ daily, supporting analytics and machine learning workloads used by 50+ stakeholders.
- Designed scalable ETL/ELT pipelines using Python, Apache Spark, and Airflow, reducing pipeline latency by 30%.
- Built centralized AWS data lake architecture enabling analytics and reporting solutions across multiple client teams.
- Developed real-time streaming pipelines using Kafka and Spark Streaming processing millions of events per day.
- Optimized Spark workloads and SQL queries through partitioning, caching, and query tuning, improving performance by 40%.
- Implemented automated data quality and validation frameworks, reducing pipeline failures by 30%.
- Built CI/CD pipelines with Docker, Git, and Jenkins, reducing deployment time by 40%.
- Implemented monitoring systems with CloudWatch and Grafana, improving incident detection by 35%.
- Led data platform architecture discussions and mentored 4+ junior engineers.
- Collaborated with data scientists and analysts to design datasets powering machine learning models and executive analytics dashboards.

Data Engineer | Aug 2019 – Feb 2022 DataCore Analytics (Data & Analytics Consulting Firm) – Newark, NJ

- Designed and developed scalable ETL pipelines integrating 20+ internal and external data sources.
- Built real-time streaming systems using Kafka and Spark Streaming, reducing processing latency by 25%.
- Designed dimensional data warehouse models (star and snowflake schemas) supporting BI systems used by 30+ analysts.
- Automated workflows using Apache Airflow, reducing manual data processing by 50%.
- Developed cloud data warehouses using Snowflake and Amazon Redshift for multi-terabyte analytics workloads.
- Optimized SQL queries and warehouse schemas, improving reporting performance by 35%.
- Built monitoring dashboards with Grafana and CloudWatch to improve operational visibility.
- Assisted in migration from on-premise infrastructure to AWS cloud architecture.

Junior Data Engineer | Jun 2016 – Aug 2019 Insight Data Systems (Private Data Engineering Consultancy) – Jersey City, NJ

- Developed **batch ETL pipelines using Python and SQL** supporting enterprise analytics platforms.
- Built scalable **data ingestion frameworks** for APIs, relational databases, and flat files.

- Supported **Hadoop and Spark clusters** processing multi-terabyte datasets.
- Implemented automated **data validation and quality checks** across ETL pipelines.
- Optimized SQL queries and database structures improving analytics performance by **30%**.
- Assisted with **AWS cloud migration projects** improving system scalability.

Education

Bachelor of Science – Computer Science
Rutgers University, New Brunswick, NJ
Sep.2012 – May.2016

Additional

- Open to fully remote roles across the U.S.
- U.S Citizen – Authorized to work for any employer.
- Available for coding challenges and interviews.