

Vamsi Krishna Bhashyam

4159964847 | app.vamsi7@gmail.com | Austin | [Linkedin](#)

Senior Data Engineer with 7 plus years of hands on experience delivering Azure lakehouse solutions across ADF, ADLS Gen2, Databricks Spark, Synapse, Azure SQL, and Fabric, with deep expertise in advanced SQL, metadata driven orchestration, and dimensional modeling. Proven client facing leader who builds secure, reliable pipelines with OAuth 2 API ingestion and JSON normalization, strengthens data quality and observability, and drives adoption through workshops, documentation, and repeatable best practice patterns.

SKILLS

Azure Data Platform: Azure Data Factory, Azure Data Lake Storage, Azure Databricks, Apache Spark, Azure Synapse Analytics, Azure SQL Database, Microsoft Fabric

Data Engineering and Orchestration: Metadata driven pipeline design, ETL and ELT pipeline development, Data integration and migrations, Automation and alerting, Scalable and modular data pipeline patterns, Batch and near real time processing

SQL and Dimensional Modeling: Advanced SQL, Relational database design, Dimensional modeling, Star schema design, Transform logic and performance tuning, Data warehouse and lakehouse modeling

API Ingestion and Semi Structured Data: REST API ingestion, OAuth 2.0 authentication, JSON parsing and flattening, Schema normalization to relational structures, Incremental loads and change handling, Data validation and standardization

Architecture, Security, and Governance: Enterprise data architecture, Lakehouse architecture, Data security architecture, Access control and least privilege, Data quality and reliability engineering, Documentation and diagramming

Leadership, Client Delivery, and Project Management: Technical leadership and mentoring, Client workshops and training, Cross functional collaboration, Project planning and delivery, Stakeholder communication, Best practice adoption and roadmap development

EDUCATION

Lindsey wilson college

Feb 2021 - Dec 2020

Masters, Information technology

EXPERIENCE

Cloud data engineer, Jpmorgan chase

Apr 2024 - Present

- Partnered with senior architects and BI stakeholders to translate client style analytics objectives into an Azure lakehouse aligned roadmap, documenting current state and target state architecture across ingestion, storage, transformation, and semantic layers, and producing a prioritized backlog of 20 plus epics with clear acceptance criteria, data contracts, and operational readiness requirements.
- Contributed to enterprise data strategy execution by implementing Azure Data Factory ingestion patterns into ADLS Gen2 using parameterized, metadata driven pipelines, enabling consistent onboarding of 8 source feeds across batch and near real time schedules, with standardized naming, folder conventions, and environment promotion practices for dev, test, and prod.
- Built and hardened API ingestion for OAuth 2 protected endpoints, implementing token acquisition, refresh handling, and retry logic, then normalizing nested JSON payloads into relational structures using Azure Databricks and Spark, delivering curated bronze to silver transformations that reduced downstream parsing effort by about 40 percent for analytics developers.
- Designed and iterated dimensional models for reporting, including star schema fact and dimension tables with conformed dimensions, surrogate keys, and slowly changing dimension patterns, and validated model usability with BI developers by mapping 30 plus business measures to a consistent semantic layer and documenting grain, lineage, and calculation logic.
- Authored advanced SQL for Azure SQL Database and Synapse style workloads, including complex joins, window functions, incremental merge patterns, and data quality assertions, and supported query tuning by reviewing execution plans and indexing strategies, improving median report query latency by about 15 percent on a representative dashboard workload.

data engineer, Smart and final

Aug 2023 - Apr 2024

- Partnered with analytics, BI, and cloud administration stakeholders to translate retail reporting and integration needs into an Azure aligned data engineering roadmap, documenting current state and target state architecture for a lakehouse style platform using ADLS and Synapse, and defining a prioritized backlog of 12 pipeline and modeling enhancements deliverable within the 6 week engagement window.
- Architected and implemented metadata driven orchestration patterns in Azure Data Factory, standardizing parameterized datasets, linked services, and reusable pipeline templates, which reduced new source onboarding time from about 2 days to under 6 hours for 4 initial subject areas, while improving operational consistency through naming standards, environment promotion steps, and runbook documentation.
- Built and hardened API ingestion for 3 external and internal services using OAuth 2 flows, implementing token management, retry logic, and incremental extraction, then flattening and normalizing nested JSON into relational landing tables in Azure SQL and Synapse staging, enabling downstream dimensional modeling with consistent keys, data types, and audit columns.
- Developed and optimized complex SQL transformations for conformed dimensions and facts, designing star schema models for sales, inventory, and promotions, and implementing surrogate keying, slowly changing dimension type 2 patterns, and late arriving fact handling, which improved Power BI dataset refresh reliability to 99 percent across daily loads during the period.

- Implemented data quality and reliability controls across ingestion and transformation layers, including row count reconciliation, schema drift detection, null and domain checks, and quarantine tables, plus automated notifications via Azure Monitor and Logic Apps, reducing time to detect load failures from hours to under 15 minutes and cutting repeat incidents by 30 percent through root cause fixes.

Data engineer, Apple

Mar 2022 - Aug 2023

- Partnered with analytics, BI, and cloud platform stakeholders to shape a pragmatic data engineering roadmap for a lakehouse style platform on Azure, translating business outcomes into a prioritized backlog of ingestion, modeling, and governance epics, and aligning delivery to Microsoft reference architectures and repeatable implementation patterns.
- Architected and implemented metadata driven orchestration in Azure Data Factory, standardizing pipeline templates for ingestion, validation, and publishing across 20 plus source feeds, improving delivery consistency and reducing new source onboarding time from about 2 weeks to 3 to 5 days through parameterization, configuration tables, and reusable activities.
- Built secure, scalable ingestion from external and internal APIs using OAuth 2, implementing token management, retry and backoff, and incremental extraction patterns, then flattening and normalizing nested JSON into relational structures in Azure Databricks using Spark, enabling reliable downstream SQL analytics for 10 million plus records per day with clear lineage from raw to curated layers.
- Designed and optimized dimensional models for executive reporting, delivering star schemas with conformed dimensions and slowly changing dimension handling in Azure Synapse Analytics and Azure SQL Database, and improving key dashboard query performance by 35 to 55 percent through partitioning strategy, selective materialization, and advanced SQL tuning.
- Established best in class data quality and reliability controls, including automated schema drift detection, row count and reconciliation checks, and freshness SLAs with notification workflows, reducing data incident triage time by 40 percent and increasing on time dataset availability from 92 percent to 98 percent within the first two quarters.

Data engineer, T-Mobile

Mar 2021 - Feb 2022

- Partnered with analytics, BI, and Azure platform teams to translate business reporting needs into an enterprise data engineering roadmap, aligning near term deliverables with Microsoft ecosystem best practices for lakehouse style architectures, including curated zones in ADLS and governed serving layers in Synapse and Azure SQL, while documenting decisions, assumptions, and tradeoffs in architecture diagrams and runbooks to accelerate stakeholder alignment.
- Designed and implemented metadata driven orchestration patterns in Azure Data Factory, using parameterized pipelines, configuration tables, and reusable templates to standardize ingestion, transformation, and publishing across multiple subject areas, improving delivery consistency and reducing new pipeline build time by about 30 percent through repeatable patterns and shared components.
- Built and enhanced ingestion pipelines from external and internal APIs using OAuth 2 flows, handling pagination, throttling, and incremental loads, then standardized semi structured payloads by flattening and normalizing JSON into relational structures for downstream analytics, enabling reliable consumption in Synapse and Azure SQL with clear lineage and documented schemas.
- Developed and optimized dimensional models for reporting, including star schemas with conformed dimensions and well defined grain, implementing surrogate keys, slowly changing dimension patterns, and fact table load strategies, which improved dashboard query performance by about 20 percent through better join paths, selective indexing, and partition aware design.
- Authored advanced SQL for complex transform logic, reconciliation, and data quality checks, including window functions, CTE based modular transformations, and incremental merge patterns, while tuning critical queries using execution plan analysis and targeted indexing, reducing recurring batch runtimes by about 15 percent within the first release cycle.

software engineer, Squirle Capital

Jun 2018 - Dec 2020

- Set data engineering vision and roadmap for analytics and reporting at Squirle Capital, aligning stakeholders across finance, operations, and engineering on a lakehouse style target architecture in Azure, then translating that strategy into a prioritized backlog of 30 plus epics covering ingestion, modeling, governance, and self service BI enablement.
- Architected and implemented end to end data integrations using Azure Data Factory for orchestration, ADLS Gen2 for landing and curated zones, and Azure Databricks with Spark for transformation, delivering 60 plus automated pipelines with parameterized templates, environment promotion, and standardized logging that reduced manual data prep effort by 25 plus hours per week.
- Built metadata driven orchestration patterns, including configuration tables, dynamic pipeline generation, and reusable notebook frameworks, enabling rapid onboarding of new sources in 2 to 3 days instead of 2 to 3 weeks, while improving reliability through idempotent loads, watermarking, and restartable processing for incremental ingestion.
- Led API ingestion initiatives for third party and partner systems, implementing OAuth 2 authentication flows, token refresh handling, and resilient retry logic, then standardizing semi structured payloads by flattening and normalizing JSON into relational structures in Azure SQL Database and Synapse, supporting 15 plus endpoints and improving downstream query consistency for BI teams.
- Designed and optimized dimensional models for executive and analyst reporting, delivering 12 star schemas with conformed dimensions, slowly changing dimension patterns, and clearly defined grain, then partnering with BI developers to validate measures and performance, resulting in 40 percent faster dashboard queries and a 20 percent reduction in metric disputes.

CERTIFICATIONS

Databricks certified data engineer, Databricks

May 2027