

Manasa Konaganti

USA, TX | +1 (469)-587-9507 | manasakonaganti226@gmail.com | [LinkedIn](#)

SUMMARY

Data Engineer specializing in scalable ETL pipelines using PySpark, Spark, AWS Glue, and Airflow. Builds reliable batch and streaming systems with Kafka and Kinesis. Strong in data modeling, Redshift, and Snowflake. Experienced with CI/CD automation, performance tuning, and production support. Focused on data quality, pipeline reliability, and stakeholder needs. Comfortable across cloud environments, collaborating with analytics teams, and delivering maintainable solutions that support business reporting needs.

TECHNICAL SKILLS

Programming Languages: Python, SQL, Shell Scripting
Big Data Processing: Apache Spark, PySpark, Hadoop, Hive, MapReduce
ETL & Data Orchestration: Apache Airflow, AWS Glue, Azure Data Factory, Talend, Informatica, Fivetran, dbt, NiFi
Streaming & Real-Time Processing: Kafka, AWS Kinesis, AWS Lambda Triggers
Cloud Platforms: AWS (S3, EC2, Glue, Redshift, Lambda), GCP (BigQuery, Dataflow), Azure
Data Warehousing: Snowflake, BigQuery, Redshift, Delta Lake, Data Lake Architecture
Databases: Oracle, MySQL, PostgreSQL, MongoDB
Data Modeling & Governance: Dimensional Modeling, Schema Design, Data Quality, Data Validation, TDD, Governance,
DevOps & CI/CD: Docker, Kubernetes, Jenkins, GitHub Actions, Terraform, Git, Linux/Unix
Analytics & Visualization: Tableau, Power BI, Looker Studio, Pandas, NumPy
Monitoring & Testing: Unit Testing, Integration Testing, Data Anomaly Detection
Project Tools & Methodologies: Jira, Agile, Scrum

PROFESSIONAL EXPERIENCE

AT&T, USA	Nov 2023 – Current
Data Engineer	
- Built ETL pipelines using PySpark and AWS Glue processing 12 TB daily reducing runtime by 3 hours and supporting 15 analytics datasets across teams. - Orchestrated 40 workflows with Apache Airflow eliminating 20 hours weekly monitoring and ensuring reliable delivery for batch and near real-time ingestion pipelines across systems. - Implemented Kafka and AWS Kinesis pipelines handling 2M events daily enabling Lambda processing and supporting 20 downstream systems with consistent low latency data delivery today. - Designed dimensional models in Redshift and Snowflake supporting 120 dashboards reducing average query time by 18 seconds and lowering monthly compute usage significantly for teams. - Automated CI/CD using Docker, Terraform, and GitHub Actions preventing 15 quarterly incidents and adding 50 data quality checks across production pipelines and release workflows.	
Accenture, India	Jan 2021 – Jul 2022
Data Engineer	
- Built batch ETL pipelines using Spark Hive Hadoop and AWS Glue processing 8 TB daily and supporting 25 business reports across finance operations teams globally. - Scheduled and monitored 30 workflows with Apache Airflow reducing 40 support tickets monthly and improving on time delivery for enterprise data refresh cycles across platforms. - Developed dbt transformation models with tests for 200 tables reducing reported issues and improving consistency across curated analytics layers used by business stakeholders daily reports. - Implemented Kafka ingestion pipelines handling 1M records daily enabling near real time reporting for marketing operations and finance stakeholders across multiple regions and teams globally. - Optimized data models and warehouse schemas improving query response by 12 seconds and reducing monthly infrastructure spend across 10 high usage analytical workloads enterprise platforms.	
Accenture, India	Jan 2019 – Dec 2020
Jr Data Engineer	
- Supported SQL and Python ETL jobs loading 3 TB daily reducing failures and improving reporting accuracy across 12 business critical datasets for stakeholders and users. - Assisted with Spark and Hive transformations on Hadoop improving runtimes by 25 minutes and stabilizing 15 scheduled batch pipelines in production environments daily operations teams. - Maintained Informatica and Talend workflows resolving 60 job failures monthly and improving overall pipeline stability through monitoring incident analysis and documented remediation procedures and automation. - Performed data validation and reconciliation across 100 tables reducing reporting mismatches and improving trust in dashboards used by operations and leadership teams daily reporting reviews. - Collaborated on schema tuning and query optimization improving response times for 20 critical reports and supporting faster ad hoc analysis for analytics users and managers.	

EDUCATION

Masters in Computer Science	Oct 2023
New England College, USA	
Bachelor of Engineering in Computer Science	Jul 2019
DVR College of Engineering, India	

PROJECTS

Real-Time Customer Analytics Pipeline
- Built streaming ingestion using Kafka and Spark Structured Streaming, processing five million events daily, reducing latency sixty percent, delivering curated tables into Snowflake for dashboards. - Orchestrated Airflow pipelines, added data quality checks, implemented schema evolution, and created session and conversion models, improving dashboard freshness by forty percent and analyst trust.
Data Lakehouse + Feature Engineering Platform
- Designed lake house on S3 with Spark and Delta, ingesting batch and CDC sources, creating Bronze Silver Gold layers, enabling reproducible transformations and scalable analytics workloads. - Built feature pipelines using Airflow, standardized definitions, versioned datasets, and served features to training jobs, reducing model preparation time by thirty five percent for teams.