

Sai Vineeth Neeli

+1 4694342575 | vineethn972@gmail.com | United States | Open to Relocate | [LinkedIn](#) | [GitHub](#)

SUMMARY

Data Engineer with 7+ years of hands-on experience designing and delivering production data platforms across Azure and AWS. Advanced SQL practitioner skilled in dimensional modeling and star-schema design, metadata-driven orchestration, and automated pipelines. Experienced with API ingestion (OAuth2) and JSON normalization into relational structures. Passionate about ensuring data quality, security, scalability, and close collaboration with BI, cloud, and analyst teams to produce reliable analytics solutions.

SKILLS & INTERESTS

Programming & SQL	: Python, Advanced SQL, PySpark
Cloud Platforms	: Azure Data Factory, ADLS (Azure Data Lake Storage), Azure Databricks, Synapse Analytics, Azure SQL Database, AWS S3, AWS Glue, Azure Databricks, ADLS (Azure Data Lake Storage), Azure Data Factory, Synapse Analytics, Azure SQL Database
Orchestration & ETL	: Metadata-driven pipelines, Airflow, Azure Data Factory orchestration
Modeling & Storage	: Dimensional modeling, Star schema design, Data lakes, Data warehouses, Snowflake
Integration & APIs	: API ingestion, OAuth2, JSON processing, flattening, and normalization
Data Quality & Security	: Data quality frameworks, Data lineage, HIPAA/PHI handling, Encryption and IAM
Tools & DevOps	: Git, Docker, Terraform, CI/CD
Soft skills	: Cross-functional collaboration, Stakeholder management, Project delivery, Mentoring

EXPERIENCE

Senior Data Engineer, U.S. Bank, Chicago, IL, United States

Apr 2023 — Present

- Build modular PySpark ETL pipelines running on Azure Databricks that read/write to ADLS, processing 2TB+ daily and improving real-time financial analytics capacity by 40%.
- Implement metadata-driven orchestration using Azure Data Factory alongside Airflow DAGs for API and JSON ingestion, reducing manual monitoring and improving pipeline reliability.
- Design OAuth2-enabled API integrations that land raw data into ADLS via streaming or ADF ingestion, improving fraud-detection throughput and lowering alert latency.
- Lead modernization efforts to Azure SQL Database and refine data architecture with Azure Synapse Analytics for reporting and scalability, improving query performance and system scalability.
- Establish metadata and data-lineage tracking to strengthen documentation, auditability, and data security for regulated financial data; lineage tracked across ADLS and Azure SQL to improve documentation and security.

Data Engineer, Sentry Insurance, Stevens Point, WI, United States

Jan 2021 — Mar 2023

- Built PySpark and Talend pipelines to ingest policy and claims data into Snowflake and AWS S3, applying normalization to improve reporting accuracy for actuarial teams.
- Automated Airflow orchestration to streamline data workflows, saving 40+ monthly hours previously spent on manual oversight and improving cross-functional delivery.
- Integrated Kafka streaming for telematics and claims feeds to enable near-real-time fraud scoring and reduce scoring latency.
- Optimized Spark jobs and advanced SQL queries on relational datasets to improve runtime efficiency and reduce cluster costs.
- Led cloud migration efforts to Snowflake and AWS data lake (S3) to support scalable data warehousing and faster reporting across business units.
- Built Python-based validation and anomaly-detection frameworks to increase data reliability for downstream analytics and actuarial models.

Data Engineer, CommonSpirit Health, Phoenix, AZ, United States

Jul 2018 — Dec 2020

- Engineered PySpark and AWS Glue ETL pipelines integrated with AWS Glue workflows to improve data quality and pipeline stability by over 35%.
- Implemented Kafka-based ingestion into AWS S3 / AWS data lake for downstream processing in Databricks, increasing availability for clinical analytics.
- Created HIPAA/PHI validation frameworks in Python to ensure compliance and secure handling of sensitive healthcare data.
- Automated Airflow workflows with recovery and alerting logic to reduce downtime and strengthen operational resilience.
- Contributed to Agile teams with workshop facilitation, code reviews, and technical guidance to support continuous improvement and stakeholder engagement.

EDUCATION

Central Michigan University, Mount Pleasant, MI, United States - *Master's, Information Systems*

PROJECTS

Intelligent Chatbot with BERT and FastAPI (End-to-End Deployment Project), [Link](#)

- Developed and deployed a BERT-based intelligent chatbot capable of understanding contextual user queries using advanced NLP and intent recognition techniques. The system was built with Python, FastAPI, and Streamlit, containerized using Docker for seamless deployment, and integrated with MLflow for experiment tracking and performance monitoring. Implemented FAISS vector search for efficient semantic retrieval, achieving improved response accuracy and faster query resolution while maintaining a scalable, modular architecture suitable for real-time enterprise applications