

Kevin Wang

Senior Data Engineer

✉ kevin.wang4613@gmail.com

📍 El Paso, TX

☎ 346 400 5517

🌐 <https://www.linkedin.com/in/kevin-wang-05170a186/>

👤 Profile

Senior Data Engineer with 8+ years of experience architecting and scaling cloud-native data platforms across e-commerce, marketplace, and advertising ecosystems. Proven track record designing distributed ETL/ELT systems processing billions of events daily using Spark, Airflow, dbt, and AWS. Deep expertise in data modeling, governance, infrastructure-as-code, and ML enablement. Recognized for driving platform reliability, cost efficiency, and cross-functional data innovation.

🧠 Skills

Programming

Python, SQL, PySpark, Java, Scala

Distributed Processing

Apache Spark, Spark SQL, Hive, EMR

Streaming & Event Processing

Apache Kafka, AWS Kinesis, Spark Streaming, Flink

Orchestration

Apache Airflow, AWS Step Functions

Data Platforms

Databricks, Snowflake, Redshift, Hadoop, Lake Formation

ML & Analytics

MLflow, Feature Store, Tableau, PowerBI, Looker

Cloud & Infra

AWS (S3, EMR, Lambda, Glue, CDK), Azure Synapse, Docker

Transformation

DBT

👜 Professional Experience

Chewy

03/2023 – Present

Senior Data Engineer

- Led architecture of Chewy's Unified Commerce Intelligence Platform, integrating clickstream, order, fulfillment, and supply chain systems into a centralized S3-based lakehouse supporting real-time personalization and demand forecasting.
- Designed scalable batch and streaming ingestion pipelines using Spark, PySpark, Airflow, and Kafka/Kinesis, processing billions of events and terabytes of transactional data daily.
- Implemented modular dbt Core transformation layer with incremental models, snapshotting, automated tests, and documentation for analytics-ready marts in Redshift.
- Developed high-performance Spark jobs leveraging partition pruning, bucketing, and optimized file formats in Parquet, reducing compute runtime by 35%.
- Built ML feature pipelines powering churn prediction, recommendation ranking, and inventory forecasting models; collaborated with Data Scientists using feature engineering best practices.
- Established observability framework with Airflow SLAs, CloudWatch metrics, logging, and alerting pipelines to reduce data incidents by 40%.
- Mentored junior engineers, led architectural reviews, and defined platform standards across orchestration, modeling, CI/CD, and secure data access patterns.

Pinterest

10/2020 – 03/2023

Senior Data Engineer

- Architected distributed ETL pipelines for Ads performance analytics and engagement modeling using Spark, Hive, Airflow, and AWS EMR, processing 5+ TB daily.
- Built scalable ingestion workflows from S3 and Kafka sources supporting campaign attribution and real-time reporting.
- Designed optimized data models for revenue analytics, user segmentation, and funnel analysis using advanced SQL and partitioned Hive tables.
- Improved Spark job efficiency via memory tuning, broadcast joins, and partition optimization, cutting execution time by 30%.
- Implemented Airflow DAG orchestration with dependency management, retries, SLA monitoring, and backfill strategies across 300+ workflows.
- Contributed to CI/CD pipelines using Git-based workflows and automated validation testing before production releases.
- Partnered with ML engineering teams to deliver curated feature datasets for targeting, ranking, and conversion prediction models.
- Strengthened governance by enforcing schema validation, metadata tracking, and structured data cataloging standards.

Lyft

06/2018 – 10/2020

Data Engineer

- Designed and maintained scalable ETL pipelines using Python, AWS Glue, and Spark, supporting ride analytics, marketplace pricing, and driver performance insights.
- Built ingestion frameworks from transactional systems into S3 data lake and Redshift warehouse.
- Developed dimensional models (star schemas) for business intelligence and operational dashboards.
- Optimized Redshift clusters through sort keys, distribution styles, and query refactoring, reducing dashboard latency by 35%.
- Implemented monitoring and logging with CloudWatch to maintain SLA compliance for high-priority data assets.
- Enhanced data lake partition strategies and compression formats to improve cost-efficiency and query performance.

Amazon

02/2017 – 06/2018

Data Analyst

- Developed advanced SQL-based analytics pipelines supporting operational performance and supply chain metrics.
- Automated reporting workflows and dashboard generation to improve stakeholder decision-making speed.
- Designed complex analytical queries using window functions and aggregations on large relational datasets.
- Partnered with engineering teams to improve upstream data validation and ETL reliability.

Education

Bachelor's Degree, Computer Science

2012 – 2016

Texas Tech University