

George Iqbal

Data Platform Engineer | Data Infrastructure Engineer | Data Solutions Engineer

✉ georgeiqbal85@gmail.com ☎ (716) 273-5147 📍 Dallas, Texas 75001

Summary

Senior DataEngineer with 8+ years of experience designing scalable cloud native data platforms, enterprise data warehouses, and real time streaming architectures. Strong expertise in AWS-based modern data stack including Databricks, Spark, Airflow, Kafka, Snowflake, and Lakehouse implementations and strong focus on governance. Experienced in healthcare data engineering (HIPAA-aligned systems), dimensional modeling, distributed processing, DataOps automation, and cost optimized cloud infrastructure. Proven ability to lead teams, modernize legacy systems, and build analytics ready platforms supporting BI, reporting, and machine learning.

Skills

Programming

Python, Advanced SQL, Bash scripting

Big Data & Processing

Apache Spark, PySpark, Spark SQL, Spark Structured Streaming, Distributed processing, Batch & real-time pipelines, Performance tuning, Partition optimization

Monitoring & Optimization

CloudWatch, Logging frameworks, SLA tracking, Cost optimization, Storage lifecycle policies

Data Warehousing & Modeling

Dimensional Modeling (Star/Snowflake Schema), Fact & Dimension design, SCD Type 1/2, Data marts, Query optimization

Streaming & Messaging

Apache Kafka, AWS Kinesis, Event-driven architecture, Real-time data ingestion, Low-latency processing

Reverse ETL & Data Activation

Skilled in synchronizing data from warehouses to applications, integrating operational systems, and enabling event-driven customer data workflows.

AI & Machine Learning Data Enablement

Skilled in preparing feature-ready datasets, harmonizing structured and semi-structured data, managing data versioning, and staging data for ML pipelines.

ETL/ELT & Data Engineering

ETL/ELT pipeline design, Incremental loading, CDC, Data transformation, API ingestion, Data reconciliation & validation

Data Observability & Reliability

Skilled in monitoring pipeline health, detecting data drift, tracking freshness, and performing root cause analysis for failures.

Cloud & Data Platforms

AWS (S3, Glue, EMR, Lambda, Redshift, Kinesis, IAM), Databricks, Delta Lake, Snowflake, Lakehouse Architecture, Data Lake design (Bronze/Silver/Gold)

Data Quality & Governance

Great Expectations, Data validation rules, Schema enforcement, Data lineage, PII masking, HIPAA-aligned controls

Orchestration & Automation

Apache Airflow (DAG development, scheduling, SLA monitoring), Terraform, CI/CD (GitHub Actions), Git

FinOps & Cost Governance

Experienced in tracking cloud resource usage, implementing storage tiering, optimizing workloads, and managing budgets for analytics environments.

Data Security Architecture

Skilled in secure data sharing, managing encryption keys, designing access policies, and ensuring regulatory audit readiness.

Professional Experience

Senior Data Engineer, *Arctic Wolf*

09/2022 – Present

- Lead design and development of cloud based data platforms across AWS and Azure for healthcare, financial, and enterprise clients.
- Build and maintain scalable data pipelines using Databricks Spark, and cloud native storage for analytics and operational reporting.
- Develop and optimize ETL and ELT workflows integrating SaaS platforms, financial systems, APIs, and relational databases.
- Design logical and physical data models supporting reporting, analytics, and regulatory requirements.
- Optimize SQL queries and improve performance across high volume transactional and analytical datasets.
- Implement data governance practices including lineage tracking, data validation, and quality controls.
- Design secure data access frameworks aligned with compliance standards in regulated industries.
- Collaborate with cross functional teams and mentor junior engineers to deliver reliable, production ready data solutions.

Data Platform Engineer, *Nexla*

07/2019 – 08/2022

- Designed and maintained a cloud native data platform on AWS enabling self service pipeline development for analytics and engineering teams.
- Built Infrastructure as Code using Terraform to provision and manage S3, EMR, Redshift, and IAM across multiple environments.
- Implemented containerized data services using Docker and Kubernetes to support Spark workloads and streaming applications.
- Developed orchestration standards using Apache Airflow and dbt for scheduling, dependency management, and workflow validation.
- Established monitoring and observability using CloudWatch and centralized logging to detect data quality and system issues.
- Implemented role based access controls, encryption, and data masking aligned with HIPAA and PCI requirements.
- Optimized Spark and storage workflows with CI/CD and automated testing.

Data Integration, *ACC3*

03/2017 – 06/2019

- Architected cloud native ELT pipelines on Amazon Web Services integrating APIs, SaaS platforms, and internal systems into a scalable data platform designed for analytics and AI readiness.
- Built real time event driven data flows using Apache Kafka to enable low latency data availability for operational reporting and machine learning use cases.
- Developed production grade transformation models in Snowflake using dbt, implementing a Medallion architecture to deliver clean, governed, business ready datasets.
- Implemented reverse ETL workflows through RudderStack to activate warehouse insights inside operational systems, supporting data driven product and marketing decisions.
- Engineered automated orchestration frameworks in Apache Airflow with CI/CD integration, version control, and monitoring to ensure reliable and self managed data pipelines.
- Integrated data governance, PII masking, encryption standards, and audit logging directly into data pipelines to support compliance requirements and secure domain oriented data sharing.

Education

Bachelor in Computer Science