

SUMMARY

Data Engineer with 4+ years of hands-on experience designing, building, and maintaining scalable data pipelines and cloud-based analytics platforms. Strong expertise in **Python**, **SQL**, **Spark**, and **modern** data stacks across **AWS** and **Azure**. Proven ability to deliver reliable **ETL/ELT** solutions, optimize **data models**, and support **business intelligence** and **ML workflows**. Known for writing clean, **maintainable code** and **collaborating** closely with cross-functional teams to deliver data-driven solutions.

SKILLS

- **Cloud Platforms:** AWS (S3, Redshift, Glue, Lambda, EMR), Azure (Data Factory, Synapse, ADLS Gen2), GCP (BigQuery)
- **Data Engineering & Analytics:** Apache Spark, PySpark, Hadoop Ecosystem, Databricks, Delta Lake, Presto / Trino
- **Data Warehousing:** Snowflake, Redshift, BigQuery, Azure Synapse
- **ETL / Orchestration:** Apache Airflow, dbt, Talend, Fivetran
- **Streaming & Messaging:** Apache Kafka, Amazon Kinesis
- **Databases:** PostgreSQL, MySQL, MongoDB, Cassandra
- **BI & Visualization:** Power BI, Tableau, Looker
- **DevOps & CI/CD:** Git, Docker, Kubernetes, Jenkins, Azure DevOps
- **Monitoring:** Prometheus, Grafana, ELK Stack
- **ML Data Support:** Feature engineering, data validation, ML pipeline support, AWS SageMaker, MLflow

EXPERIENCE

CitiGroup | USA

Jan 2025 –Present

Data Engineer

- Designed and maintained enterprise-scale data pipelines on Azure using Azure Data Factory, Azure Synapse, and ADLS Gen2, supporting analytics and reporting use cases.
- Developed ELT workflows using PySpark and SQL to ingest, transform, and validate large financial and transactional datasets.
- Built and optimized star and snowflake schema data models, improving Power BI query performance and reducing report refresh time by approximately 40%.
- Developed Power BI dashboards for risk, compliance, and operations teams, enabling near real-time insights for senior stakeholders.
- Implemented CI/CD pipelines using Azure DevOps and Git to automate deployments and enforce version control best practices.
- Collaborated with data science teams to prepare feature-ready datasets supporting machine learning models on Azure ML.
- Applied data quality checks, logging, and governance standards to ensure accuracy, reliability, and regulatory compliance.

Tech Stack: Azure Data Factory, Azure Synapse, ADLS Gen2, PySpark, SQL, Python, Power BI, Azure DevOps, Git, Azure ML

EClerx | India

Aug 2020 – Jul 2023

Data Engineer / Data Analyst

- Built and supported AWS-based data pipelines using S3, AWS Glue, Lambda, and Redshift to deliver analytics-ready datasets.
- Developed batch ETL jobs in Python and SQL to process operational and financial data for reporting and business analysis across multiple business teams.
- Designed optimized Redshift data models, improving dashboard query performance and reducing report latency.
- Created and maintained Tableau dashboards used by finance and operations teams for KPI tracking and performance analysis.
- Automated recurring reporting workflows using Python scripts, reducing manual reporting effort and turnaround time.
- Partnered with business stakeholders to translate requirements into scalable and maintainable data solutions.

Tech Stack: AWS S3, Glue, Lambda, Redshift, Python, SQL, Tableau, Git, Linux

EDUCATION

- Master of Science in Data Science | University of Texas at Arlington, Dallas, Texas |

Aug 2023 – May 2025