

Pranjal Iyer

Newark, NJ | 551-263-8382 | pranjaliyer7@gmail.com | [LinkedIn](#)

PROFESSIONAL SUMMARY

Senior Data Engineer (7+ years) specializing in large-scale distributed data platforms using Spark, Python, SQL, Delta Lake, EMR, and Redshift on AWS. Expertise in high-volume pipelines (~4 TB/day), ML-ready datasets, LLM/Text-to-SQL automation, real-time analytics, data modeling, ETL/ELT, DQ frameworks, and performance optimization across Amazon and JPMorgan.

SKILLS

Big Data: Spark (2.x/3.x), EMR, Delta Lake, Iceberg, Redshift (RA3), Databricks, Snowflake

Languages: Python, SQL, PySpark, Java, Scala

Cloud: AWS (S3, Glue, Lambda, Athena, Step Functions, CloudWatch, SQS), Azure

Data Engineering: ETL/ELT, Data Modeling, Distributed Systems, DQ, Metadata, Lakehouse

Orchestration/CI-CD: Autosys, Airflow, Git, GitHub Actions

Tools: Docker, Kubernetes, Qlik, Datorama, Tableau, Terraform

WORK EXPERIENCE

JPMorganChase

Jersey City, NJ

Software Engineer II

Mar 2023 – Present

- Built Person360 pipelines using Spark and Redshift to unify 15+ financial datasets, surfacing \$700M+ high-net-worth client prospects across Private Bank CRM.
- Created Text-to-SQL engine with LLM to convert analyst prompts to Redshift SQL, reducing ad-hoc queries by 50%.
- Optimized Spark jobs on EMR (~4 TB/day) using partition tuning, memory config, and cluster sizing, cutting runtime by 40%.
- Migrated Salesforce Datorama dashboards to Qlik using S3, Athena, and Lambda, saving \$1M/year in licensing and manual refreshes.
- Built data quality checks with lineage, anomaly flags, and CloudWatch alerts to ensure audit-grade financial analytics.
- Designed modular data models enabling ML and reporting use cases, improving experimentation cycle times for analytics teams.
- Partnered with Private Bank and compliance teams to prioritize KPIs and deliver high-impact dashboards and data services.
- Ingested and normalized structured/unstructured data via PySpark and Delta Lake pipelines, enhancing data availability and governance.
- Deployed CI/CD workflows using GitHub Actions and Jenkins for safe, auditable, and reproducible data pipeline deployments.

Amazon

New York City, NY

Data Engineer

May 2022 – Mar 2023

- Built and optimized Spark + Redshift pipelines of retail and personalization data using S3 and Delta Lake.
- Collaborated with DS teams to deliver ML-ready datasets in Databricks, accelerating model training and A/B testing cycles.
- Improved Redshift performance via WLM tuning, schema redesign, and query optimization, cutting compute cost and query latency.
- Implemented observability with CloudWatch, alerts, and job lineage across 40+ pipelines to ensure operational reliability.
- Built Git-managed modular workflows and parameterized notebooks to enable reproducible deployment pipelines.

LTIMindtree(formally Mindtree Ltd.)

Chennai, India

Module Lead

Jun 2018 - Dec 2020

- Architected ETL pipelines in Spark + Python on AWS for HR and Finance, supporting 500M+ records/month with lineage and quality checks.
- Developed ML recommendation engines using Spark MLlib, raising client revenue \$1M+ via cross-sell and personalization models.
- Built ingestion pipelines using Delta Lake and EMR, enabling resilient batch transformations with reduced SLA misses.
- Led requirement gathering, solution design, and testing reviews with cross-functional finance and compliance teams.
- Established CI/CD processes via Jenkins and Git, improving release velocity and unit-test coverage across all ETL modules.

ADDITIONAL EXPERIENCE

Qruize Technologies - Data Research Trainee (2017–2018): Spark-based data mining, Azure pipelines, DQ & cleaning.

HCL Technologies - Data Analyst (2016–2017): Built ingestion & streaming ETL frameworks; implemented automated test cases.

EDUCATION

Stevens Institute of Technology

Hoboken, NJ

Master in Data Science

Aug 2022

Nagpur University

Nagpur, India

Bachelor of Engineering

May 2016