

Doja Abiose

Senior Data Engineer

Email: dojaabiose911@gmail.com | Phone: +1 (346) 814-8984 | Location: Texas

SUMMARY

- Senior Data Engineer (6+ years) delivering **Databricks + Snowflake lakehouse** platforms for high-volume, regulated enterprise data environments.
- Built and optimized **Spark/Delta pipelines at multi-TB scale**, improving throughput, reliability, and repeatable multi-environment deployments.
- Strong in **DataOps/CI/CD** (Azure DevOps YAML, GitHub Actions, Databricks CLI), enabling gated promotions, rollback readiness, and release governance.
- Designed **dimensional models + SCD Type 1/2** in Snowflake, improving historical accuracy and reducing downstream reporting discrepancies.
- Implemented **data quality/validation frameworks** (reconciliation, schema + join checks, freshness controls), reducing defects and increasing analytics trust.

SKILLS

- **Languages:** Python, SQL, Scala (Spark), Java, PL/SQL, Bash/Shell
- **Lakehouse & Warehouses:** Databrick, Delta Lake, Snowflake, Apache Spark, Hadoop (legacy exposure), BigQuery/Redshift
- **Orchestration & Scheduling:** Apache Airflow, Databricks Workflows; Prefect/Dagster
- **Streaming & CDC:** Kafka, Spark Structured Streaming, CDC patterns, Kinesis
- **Cloud Platforms:** Azure (ADLS, Key Vault, Azure DevOps CI/CD, Service Principals, Databricks), AWS (S3, EC2, EMR, EKS, IAM), GCP (GCS/BigQuery exposure)
- **Containers & IaC:** Docker, Kubernetes, Terraform
- **Databases & Data Stores:** PostgreSQL, SQL Server, MySQL, MongoDB, Redis, AtlasDB
- **Analytics Engineering & Modeling:** ETL/ELT, Data Modeling, Dimensional Modeling, SCD Type 1/2, dbt
- **Data Quality & Testing:** Data Quality Frameworks, Great Expectations / Deequ / dbt tests, reconciliation checks
- **Data Observability:** Monte Carlo / Bigeye / Databand, SLA/SLO monitoring, anomaly detection
- **Governance & Security:** Unity Catalog, PII controls, row/column-level security concepts, secrets management
- **CI-CD for Data:** Azure DevOps YAML Pipelines, GitHub Actions, Git branching/release strategies, Databricks CLI, CI/CD automation, Release Management
- **Monitoring & BI:** Grafana, Kibana, Tableau, Power BI, Matplotlib, Seaborn

PROFESSIONAL EXPERIENCE

BECU **Mar 2024 – Present**

Senior Data Engineer

- Architected multi-environment Databricks lakehouse pipelines processing 3TB+ daily, improving stability and reducing failures ~30% via Spark tuning.
- Standardized orchestration patterns (parameterization, retries, idempotency, dependency controls), reducing rerun waste and increasing consistency across dev/test/prod.
- Built Azure DevOps multi-stage CI/CD (YAML) with gated promotions and service principal auth, cutting deployment errors ~30% across 20+ products.
- Automated environment provisioning using Databricks CLI + Git workflows, reducing manual drift and enabling repeatable, auditable releases.
- Designed Snowflake dimensional models with SCD Type 1/2, improving historical tracking and reducing reporting discrepancies ~20% for finance/ops dashboards.
- Implemented data quality framework (schema, reconciliation, join accuracy, freshness), reducing ingestion defects ~25% across shared pipelines.
- Strengthened governance by aligning Unity Catalog RBAC, cluster policies, secrets handling, and workspace standards across 8 Databricks workspaces.
- Built runbooks and failure-handling standards, reducing time-to-triage and preventing repeat incidents across critical datasets.
- Improved observability using Grafana/Kibana dashboards and alert conventions, speeding root-cause analysis and reducing noisy paging.

Elevance Health **Mar 2022 – Mar 2024**

Data Engineer

- Owned claims reporting pipelines processing 2TB+ weekly, transforming MongoDB event data into curated datasets and improving reporting accuracy 35%.
- Built and optimized ETL/ELT workflows using Airflow, improving execution reliability and reducing end-to-end pipeline failures.
- Enhanced CMS/medical ingestion for compliance reporting, improving SLA adherence 85% → 98% and cutting failures 30%.
- Implemented validation + reconciliation controls, reducing inconsistencies 25% and improving reliability 40% across 6 business units.
- Standardized reusable ingestion and transformation patterns, reducing rework and accelerating delivery across multiple consumer teams.
- Partnered with analysts to translate evolving requirements into scalable pipelines, reducing delivery timelines 45%.

- Improved curated-layer modeling practices, strengthening downstream reporting alignment across 10+ consumer teams.
- Improved supportability through documentation (architecture, lineage assumptions, failure-handling), cutting onboarding time 35%.
- Reduced production errors by enforcing release checklists, validation gates, and post-deployment verification aligned to SLAs.
- Supported time-sensitive datasets through event-driven ingestion patterns, improving freshness for operational and reporting use cases.

Thoughtworks

Feb 2020 – Mar 2022

Consultant Software Developer II

- Designed and delivered backend microservices that improved scalability and reduced latency by ~20% across business-critical workflows.
- Built REST APIs with secure access patterns and robust error handling, improving reliability and reducing escalation tickets.
- Implemented automated testing and regression coverage, lowering post-release defects by ~20% across multi-team releases.
- Performed debugging and performance tuning, improving throughput and reducing intermittent production issues.
- Collaborated with UX/UI and product teams to refine user-facing functionality, improving usability and engagement.
- Delivered enhancements in Agile delivery cycles, improving predictability through better estimation and sprint execution discipline.
- Improved code maintainability via refactoring and clean architecture practices, reducing time spent on recurring tech debt.
- Supported CI/CD practices and release readiness activities, improving deployment consistency and reducing rollback frequency.
- Documented technical decisions and implementation notes, improving handoffs and reducing onboarding friction within teams.
- Coordinated with cross-functional stakeholders to align scope and technical feasibility, preventing late-stage rework.

Visa Inc., CA

May 2018 – Aug 2018

Software Engineer Intern

- Built near real-time ingestion pipelines aggregating service logs from 3,000+ Hadoop servers, improving RCA turnaround time by 50%+.
- Normalized and structured high-volume logs for downstream querying, improving diagnostic speed for platform teams.
- Automated log structure detection and pattern classification, reducing manual analysis by ~60% and accelerating diagnostics.
- Improved pipeline performance through batching and processing optimizations, reducing processing delays during peak loads.
- Implemented validation checks to ensure completeness and correctness, reducing noisy datasets and false investigative leads.
- Enhanced observability with dashboards and metrics reporting, improving visibility into ingestion health and failure causes.
- Collaborated with senior engineers to troubleshoot production-like failures and implement stable remediation approaches.
- Documented ingestion logic, runbooks, and troubleshooting steps, improving operational continuity during handoffs.
- Supported platform reliability initiatives by analyzing log anomalies and recurring error signatures.
- Participated in code reviews and best-practice alignment, improving code quality and maintainability.

Education:	
M.Sc Computer Science	
Prairie View A&M University, Texas	G.P.A 3.53 2017 – 2019
B.Sc Computer Information Systems	
Babcock University	G.P.A 3.28 2009 – 2013

Certifications:	
Databricks Certified Data Engineer Professional	Expected: March 2026