

Kyle Evans

Senior Data Engineer

Tel: 856.301.2505 / kyle.evans0328@outlook.com / Tolono, IL 61880

PROFESSIONAL SUMMARY

Senior Data Engineer with 13+ years of experience building scalable data platforms and cloud-native solutions across AWS, Azure, and GCP. Expert in real-time streaming, lakehouse architectures (Databricks, Snowflake, Delta Lake), and modern data orchestration using Airflow, Kafka, and Spark. Proven success in platform modernization, data governance implementation, and infrastructure automation with Terraform and Kubernetes. Skilled in leading DataOps practices, mentoring teams, and delivering production-grade data pipelines and ML feature engineering solutions.

SKILLS

- **Programming:** Python, SQL, Scala, Shell, R, Java
- **Frameworks & Tools:** Apache Spark, PySpark, Databricks, Apache Beam, Kafka, Flink, Hive, Hadoop, dbt, Airflow, Informatica Cloud
- **Data Architecture:** Lakehouse (Delta Lake, Iceberg), Data Mesh, Data Vault, Star/Snowflake Schema
- **Cloud Platforms:** AWS (Glue, Redshift, Lambda, S3, EMR, Kinesis, SQS, KMS), Azure (ADF, Synapse, Databricks, Data Factory), GCP (BigQuery, Dataflow, Composer)
- **Databases:** PostgreSQL, MySQL, MongoDB, Cassandra, Oracle, DynamoDB, Snowflake, Teradata
- **Infrastructure & DevOps:** Docker, Kubernetes, Terraform, Jenkins, GitHub Actions, CI/CD, CloudFormation
- **Data Quality & Observability:** Great Expectations, Monte Carlo, Soda, Prometheus, Grafana, ELK Stack, OpenLineage
- **Visualization:** Power BI, Tableau, Looker, Superset
- **Other:** REST APIs, GraphQL, JSON, Parquet, Avro, XML, DVC, MLflow, TensorFlow Extended (TFX), Prophet, Pandas, pytest, Workday

WORK EXPERIENCE

Advanced Data Systems

Senior Data Engineer

Paramus, NJ

FEB 2020 – PRESENT

- Led data platform modernization projects for healthcare client across **AWS, Azure, and GCP**. Ensured all solutions met **HIPAA compliance** requirements while improving data access for clinical and business teams.
- Built **data lakehouse** architectures using **Databricks, Delta Lake, and Snowflake** to handle clinical data, insurance claims, and patient records. Cut data processing time by 45%.
- Developed real-time and batch pipelines with **Spark, Kafka, Airflow, and Informatica Cloud** to support clinical decisions, population health tracking, and value-based care reporting.
- Migrated legacy on-premise ETL jobs to cloud platforms including **Informatica Cloud, AWS Glue, and Azure Data Factory**. Reduced infrastructure costs by 30% and improved scalability.

- Designed **data mesh** setups that let different teams (clinical, operations, finance) own their data while following **HL7** and **FHIR** standards for healthcare data exchange.
- Put in place data quality monitoring using **Great Expectations**, **Monte Carlo**, and **Python** scripts. This caught data issues before they affected clinical analytics.
- Built **ETL pipelines** to pull data from EHR systems like Epic and Cerner, claims platforms, and tools like **Salesforce Health Cloud** and **Workday**. Made sure data stayed consistent and met compliance rules.
- Automated infrastructure setup with **Terraform** and **Kubernetes**. This gave us consistent, HIPAA-compliant deployments across different client environments.
- Set up **data governance** using Unity Catalog and schema registries. Added role-based access, audit logs, and PHI masking to protect sensitive data.
- Built machine learning feature pipelines on **Databricks** for predicting hospital readmissions, patient health decline, and gaps in care coverage.
- Gave clinical staff and analysts self-service access to data through **Snowflake** and **Looker** dashboards while keeping security tight.
- Built streaming data systems with **Kafka**, **Kinesis**, and **Flink** for real-time patient monitoring, remote care devices, and operational alerts.
- Created **dbt** transformation workflows with version control and automated testing in **CI/CD pipelines**. This kept analytics reliable for healthcare teams.
- Made flexible ingestion frameworks using **YAML** configs and parameterized mappings. This let us quickly add new healthcare data sources with built-in validation.
- Built **data lineage** dashboards combining **Airflow** metadata, **OpenLineage**, and **Grafana**. Gave full visibility for audits and regulatory needs.
- Explored newer tech like **Apache Iceberg**, **Delta Sharing**, and **Lakehouse Federation** to improve how healthcare data gets shared between systems.
- Mentored junior engineers through code reviews and hands-on training. Taught best practices for data engineering, healthcare data standards, and secure coding.
- Worked with executives and clinical leaders to define **KPIs**, **data SLAs**, and quality metrics for analytics projects spanning multiple departments.
- Standardized DataOps practices like automated testing, monitoring, and CI/CD across client projects. Increased delivery speed by 40%.

Logiwa

Senior Data Engineer

Chicago, IL

JAN 2017 – JAN 2020

- Built end-to-end data platform for ecommerce operations handling customer transactions, inventory, and analytics at scale.
- Developed real-time data ingestion using **Kafka**, **AWS Kinesis**, and **Spark Streaming**. Processed 15M+ customer events per hour.
- Created data lake pipelines in **AWS Glue**, **Redshift**, and **Databricks**. Tracked customer behavior, product performance, and supply chain data.
- Automated 500+ ETL workflows using **Airflow**. Added failure recovery and SLA monitoring.
- Migrated **Oracle** and **Teradata** systems to **Snowflake**. Cut query time by 55%.
- Designed star and snowflake data models. Built dashboards for sales, inventory, and customer metrics across 10+ teams.
- Set up data governance using schema registry and catalog tools. Kept product, customer, and order data consistent.
- Built datasets for recommendation engines and demand forecasting. Used **PySpark**, **Pandas**, and **DVC**.
- Tuned ETL jobs with partitioning and incremental loads. Improved SLA compliance by 30%.

- Created data lineage dashboards with **Airflow**, **Great Expectations**, and **OpenLineage**. Showed data flow from source to reports.
- Built event-driven services using **AWS Lambda** and **SQS**. Handled inventory updates, orders, and pricing in real-time.
- Automated infrastructure with **Terraform**, **CloudFormation**, and **Jenkins**. Cut manual work by 95%.
- Optimized **Snowflake** costs with auto-suspend and storage tiering. Handled peak shopping traffic efficiently.
- Built KPI dashboards in **Power BI** and **Tableau** for business teams. Tracked conversions, cart abandonment, and customer value.
- Added data encryption with **KMS** and PII tokenization. Set up access controls for sensitive data.
- Created DataOps standards with **CI/CD**, automated testing, and monitoring.
- Added **dbt** transformations for customer segmentation and product analytics. Made SQL models modular and testable.
- Built forecasting pipelines using **Python** and **Prophet**. Reduced stockouts and overstock issues.
- Mentored engineers on lakehouse design and ML pipelines. Boosted team productivity by 40%.
- Added unit testing with **pytest** and **Great Expectations**. Made pipelines more reliable.

FourKites

Data Engineer

Chicago, IL

MAY 2012 – DEC 2016

- Built ETL workflows using **Informatica**, **Talend**, and **Python**. Ingested customer and network data from multiple sources.
- Designed data marts and warehouses in **AWS Redshift** and **Azure Synapse**. Enabled analytics and reporting for business teams.
- Created **Airflow** DAGs for pipeline orchestration. Managed dependencies across distributed systems.
- Built streaming pipelines using **Kafka** and **Spark Streaming**. Supported network monitoring and real-time alerts.
- Wrote data transformation logic using **SQL**, **SparkSQL**, and **Pandas**. Cleaned and standardized data for analytics.
- Built dashboards in **Tableau** and **Power BI**. Showed operational KPIs for executives.
- Migrated **Teradata** workloads to **AWS Redshift**. Cut infrastructure costs by 25%.
- Set up CI/CD pipelines with **Jenkins** and **GitHub**. Automated build, test, and deploy for ETL jobs.
- Ran data reconciliation between landing, staging, and target zones. Validated data accuracy at each step.
- Documented data lineage using custom metadata repositories. Tracked data flow across systems.
- Defined data retention and lifecycle policies. Ensured compliance with regulatory requirements.
- Implemented **CDC** and **SCD Type-2** for incremental loads. Captured changes efficiently from source systems.
- Built **Python** validation scripts for data auditing. Compared large datasets for accuracy checks.
- Worked with BI and analytics teams to improve data availability. Streamlined data access for visualization tools.
- Set up job monitoring dashboards using **ELK Stack** and **Grafana**. Tracked pipeline health and performance.

EDUCATION

University of Illinois Chicago

Bachelor of Science in Computer Science

Chicago, IL

SEP 2008 – MAY 2012