

Kelly Smith

Senior Data Engineer

✉ kelly.s3827@gmail.com ☎ 2313801210 📍 Richardson, TX 🔗 LinkedIn

Profile

Senior Data Engineer with 11+ years of experience architecting scalable data solutions across enterprise and hyper-growth tech environments. Expert in building robust ETL pipelines and lakehouse platforms using Databricks, Apache Spark, Kafka, Flink, Airflow, and dbt on AWS, Azure, and GCP. Adept at delivering high-quality, ML-ready datasets and real-time analytics while ensuring data governance with tools like Great Expectations. Known for mentoring engineers, automating infrastructure with Terraform, and driving business impact through clean, reliable, and cost-efficient data systems.

Professional Experience

Cardinal Health

04/2022 – Present

Senior Data Engineer

- Directed the modernization of Cardinal Health's enterprise data architecture by migrating SAP workloads to GCP using Terraform, Docker, and Databricks, improving reliability and reducing infrastructure risk.
- Migrated 203 million business records to a 24TB SAP HANA instance on GCP via Terraform and Docker, delivering 95% uptime and seamless analytics integration with BigQuery.
- Constructed scalable PySpark pipelines in Databricks to extract and transform clinical and supply chain data from SAP, significantly reducing reporting lag across business units.
- Developed real-time Kafka and Flink pipelines to process device telemetry and prescriptions, boosting data availability and enabling on-demand insights for analytics teams.
- Provisioned Kafka, Dataproc, and Kubernetes infrastructure using Terraform and CI/CD with Jenkins, reducing deployment errors and enhancing consistency across data environments.
- Engineered Spark-based ETL workflows integrated with Delta Lake and Airflow DAGs, reducing daily pipeline failures by 87% and expediting data delivery for dashboards.
- Streamlined code deployments using Docker and Git-based workflows with Jenkins and Concourse, improving version control and decreasing rollback time for production fixes.
- Led the implementation of modular dbt transformations and star schema models, enabling Tableau and Looker users to derive actionable insights from clean data layers.
- Implemented row-level security in BigQuery using Airflow and Collibra metadata, strengthening PHI access controls and achieving HIPAA compliance audit readiness.
- Mentored junior engineers on Spark, Flink, and dbt; facilitated peer reviews and onboarding sessions to promote quality standards and technical excellence across the data engineering team.

DoorDash

Senior Data Engineer

08/2020 – 03/2022

- Led development of real-time event platform "Iguazu" using Kafka, Flink, Delta Lake, and Snowflake, delivering 99.99% SLA and simplifying event streaming for 174 teams.
- Designed Spark pipelines on Databricks to aggregate S3 data into Snowflake marts, achieving a 37% improvement in compute efficiency through optimized joins and caching.

- Automated Airflow DAG deployments and Kafka cluster provisioning using Docker and Terraform, reducing manual configuration time and enhancing operational stability.
- Migrated 227 terabytes of batch and streaming data to Delta Lake with Flink integration, accelerating downstream reporting and eliminating delays in metrics refresh.
- Guided adoption of standardized dbt models and modular Airflow DAGs; facilitated code reviews and internal workshops to align engineers on scalable pipeline architecture.

Data Engineer

08/2017 – 07/2020

- Constructed end-to-end ETL pipelines in Spark (Scala) and Airflow to ingest S3 data into Snowflake, supporting core business dashboards for operations and marketing.
- Developed Flink and Kafka pipelines for unifying mobile and backend events, enabling real-time user insights and schema validation with Confluent Schema Registry.
- Optimized Snowflake SQL queries for analytics use cases, reducing query runtimes for financial metrics by 42% through targeted refactoring and indexing.

Nike

08/2014 – 07/2017

Data Analyst

- Delivered cross-functional insights by transforming product and sales data using Python and SQL, supporting Tableau dashboards that improved marketing ROI visibility.
- Built robust ETL pipelines to integrate logistics, inventory, and customer feedback data from Oracle and Hive, enabling data-driven supply chain adjustments.
- Modeled fact and dimension tables in Redshift and Hive, aligning analytics workflows with business metrics for merchandising and finance reporting across 47 regions.
- Collaborated with analysts to define KPIs and automate monthly reporting through SQL workflows, reducing manual effort and accelerating delivery by 41%.

Education

Bachelor's Degree in Computer Science

2010 – 2014

Texas Tech University

Skills

Programming & Data Processing

Python, Java, Scala, SQL, Apache Spark, PySpark, Hadoop, Hive, Flink, Pandas

Data Architecture & Storage

Delta Lake, Snowflake, Redshift, BigQuery, Azure Synapse, PostgreSQL, NoSQL, Data Lakehouse, Dimensional Modeling, Medallion Architecture

Data Quality & Observability

Great Expectations, Deequ, OpenLineage, Marquez, Lake Formation, IAM, KMS, CloudWatch, Datadog

ETL & Integration

Apache Airflow, dbt, Kafka, AWS Step Functions, Flink, Azure Data Factory, Glue Jobs, Control-M

Cloud & Infrastructure

AWS (S3, Glue, Redshift, Lambda, EMR, SageMaker, MWAA), Azure (Synapse, ADF, Databricks), GCP (BigQuery, Dataflow), Terraform, Docker, Kubernetes

Analytics & Visualization

Looker, Tableau, Power BI, QuickSight, SageMaker, Feature Stores, Vector Databases, ML Data Pipelines