

KEVIN FERNANDEZ

Chicago, IL | (470) 574-0608 | kfernan.work@gmail.com | <https://www.linkedin.com/in/kevin-94656930/>

SUMMARY

Senior Data Engineer with **12+ years** of experience designing scalable cloud-native data platforms across **AWS** and **Azure**. Expertise in **Apache Spark**, **distributed data processing**, **ETL/ELT pipelines**, and cloud data lake architectures supporting large-scale analytics and AI workloads.

Proven experience building high-performance data pipelines using **PySpark**, **Snowflake**, **Airflow**, **dbt**, and **Databricks**, optimizing data processing for multi-terabyte datasets and real-time streaming systems. Strong background in **AWS** (S3, **Glue**, **Redshift**, **Kinesis**, **EMR**, **Athena**) and **Azure** (**Databricks**, **Synapse**, **Data Factory**, **Azure SQL**, **Microsoft Fabric**).

Experienced in **data warehousing**, query optimization, medallion architecture, and scalable analytics platforms. Known for partnering with cross-functional teams, mentoring engineers, and delivering reliable data systems that enable business intelligence, reporting, and machine learning initiatives.

EDUCATION

Georgia Institute of Technology

Bachelor of Computer Science, 2009 - 2013

SKILLS

- **Programming:** Python, SQL (Advanced), Scala, Java, Shell, PowerShell
- **Big Data & Distributed Processing:** Apache Spark (PySpark, Spark SQL, Structured Streaming), Databricks, Kafka, Kafka Connect, Apache Beam, Apache Flink
- **Cloud Platforms:** AWS (S3, Redshift, EMR, Glue, Kinesis, Athena, Lambda, CloudWatch, IAM), Azure (Azure Databricks, Azure Synapse Analytics, Azure Data Factory, Azure SQL, Microsoft Fabric (OneLake, Lakehouse, Direct Lake))
- **Data Warehousing:** Snowflake, Amazon Redshift, Azure Synapse, BigQuery, Data Modeling (Star Schema, Snowflake Schema, Data Vault, Medallion Architecture)
- **Data Pipeline & Orchestration:** Apache Airflow, dbt, Azure Data Factory, AWS Glue, Cloud Composer
- **Streaming & Real-Time Processing:** Kafka, AWS Kinesis, Spark Streaming
- **DevOps & Infrastructure:** Docker, Kubernetes, Terraform, Jenkins, GitHub Actions, Azure DevOps
- **Data Quality & Governance:** Great Expectations, Data Profiling, Metadata Management, Data Lineage, RBAC, IAM
- **BI & Analytics:** Power BI (DAX, Power Query), Tableau, Looker

EXPERIENCE

Infosys | Senior Data Engineer

Jun 2022 – Present

- Architected and implemented cloud-native data platforms across **AWS** and **Azure**, supporting analytics and ML workloads on multi-terabyte datasets.
- Designed **Spark-based ETL pipelines** using **Databricks**, **PySpark**, and **Airflow**, improving large-scale data processing performance and reliability.

- Built scalable **data lakehouse** architectures using **Microsoft Fabric** and **Delta Lake**, enabling unified data access for BI, analytics, and AI workloads.
- Developed real-time ingestion pipelines using **Kafka** and **AWS Kinesis**, reducing data availability latency by 30–40% for operational analytics.
- Automated data ingestion and transformation pipelines using **Airflow**, **dbt**, **Spark**, and **Python**, improving pipeline reliability and reducing manual intervention.
- Optimized **Snowflake** and **Amazon Redshift data warehouses** through schema design, indexing, clustering, and SQL tuning, improving query performance by 20–30%.
- Implemented enterprise data quality validation using Great Expectations, reducing downstream data defects by 25%+.
- Automated cloud infrastructure deployment using **Terraform**, **Docker**, and **Kubernetes**, improving deployment speed and consistency.
- Implemented centralized monitoring using **CloudWatch**, **ELK Stack**, and **Grafana** to improve observability and reduce incident resolution time by 30%.
- Collaborated with product managers, analysts, and data scientists to deliver **data products** and **analytics** solutions.
- Mentored **4–6 junior engineers** through design reviews, architecture guidance, and code reviews.

Capgemini | Data Engineer

Jul 2018 – May 2022

- Designed and developed **AWS-based big data pipelines** using **Spark**, **Kafka**, **Kinesis**, and **AWS Glue**, processing **hundreds of millions of events per day**.
- Built and maintained enterprise **data lakes** on **Amazon S3**, supporting large-scale analytics and reporting workloads.
- Developed ETL pipelines using **Spark**, **Python**, and **Airflow**, enabling scalable batch and streaming data processing.
- Engineered **Snowflake** and **Redshift data warehouses**, implementing optimized data models and complex SQL transformations.
- Improved warehouse query performance by 20–25% through SQL tuning, indexing strategies, and partition optimization.
- Implemented secure cloud data architectures using **IAM roles**, **encryption**, and **RBAC** to meet compliance requirements.
- Automated pipeline deployment using **Jenkins** and **GitHub Actions**, reducing production incidents by 20%.
- Supported production data pipelines with monitoring, performance tuning, and incident response.

Deloitte | Software Engineer

Jun 2013 – Jun 2018

- Developed backend services and data APIs using **Python**, **PostgreSQL**, **MongoDB**, and **Elasticsearch** supporting enterprise applications.
- Built data ingestion pipelines to integrate and normalize external datasets for analytics and reporting.
- Designed scalable data processing workflows and optimized relational database performance.
- Introduced CI/CD pipelines and containerization, improving deployment reliability and release cycles.