KYLE **KIFER**

San Antonio, TX · 210-926-0860
kylekifer3@outlook.com
LinkedIn Profile

Lead Data Engineer with 11+ years of experience architecting enterprise-scale, cloud-native data platforms and Big Data ecosystems. Expert in designing highly reliable, secure, and scalable data pipelines across AWS and hybrid environments. Proven leader with a strong background in reliability engineering, CI/CD automation, and mentoring high-performing data teams while delivering measurable business impact.

SKILLS

- **Programming Languages**: Python 3.7, 3.8, 3.9, 3.10, 3.11, Java 8/11, Scala 2.11/2.12, SQL (ANSI), C++, C#, R
- **Big Data & Streaming**: Hadoop 2.7, 3.2; Spark 2.4, 3.3; Kafka 2.6, 3.4; Hive 2.3, 3.1; Flink 1.13; Storm; Airflow 2.2–2.7; MapReduce
- **AI / ML Enablement**: Feature pipelines, MLOps data workflows, MLflow, Spark ML data prep, model monitoring, SageMaker
- **Data Engineering & Transformation**: ETL / ELT, dbt (Core & Cloud)
- **Data & Analytics Libraries**: NumPy 1.18–1.26, Pandas 1.0–2.1, Power BI, Looker / LookML, Databricks
- **Databases**: MySQL 5.7, 8.0; PostgreSQL 11, 12, 13, 14, 15; MongoDB 4.4, 5.0; Cassandra 3.11, 4.0; Redis 6.2, 7.0; Elasticsearch 7.17, 8.x
- **Cloud & DevOps**: AWS (Lambda, S3, API Gateway, DynamoDB, EC2, Glue, Redshift, EMR), Azure, GCP(BigQuery, Dataflow), Oracle Cloud, IBM Cloud, Docker 20.10, 24.x; Kubernetes 1.20, 1.22, 1.25

EXPERIENCE

01/2020 – CURRENT

SENIOR DATA ENGINEER, INFOSYS

- Led the redesign of enterprise data infrastructure, improving system efficiency and reducing data retrieval times by **40%.**
- Architected and implemented automated ETL pipelines, decreasing code redundancy by **30%** and improving maintainability.
- Established CI/CD pipelines using GitLab 15, C#, **Python 3.9**, and Shell scripting (Bash) to automate deployments, configuration management, and routine operational jobs.
- Designed and built cloud-native data pipelines on **AWS**, leveraging Lambda (Python 3.9), API Gateway, S3, DynamoDB, and Snowflake (2021) for scalable data ingestion and processing.
- Developed AWS Lambda functions in Python to process complex and nested JSON data, including transformation, comparison, and sorting logic.
- Implemented Apache Hadoop 3.2 and Spark 3.2 for large-scale data processing, increasing overall data accuracy by 20%.
- Utilized Airflow 2.5 for workflow orchestration and scheduling.
- Built Databricks-based lakehouse solutions using Spark and Delta Lake, supporting large-scale batch and incremental data processing for enterprise analytics use cases.
- Containerized services using **Docker** 20.10 and orchestrated deployments with Kubernetes 1.25.
- Strengthened enterprise data security protocols, resulting in zero security breaches over the past year.
- Championed a data-driven culture by training non-technical stakeholders on data interpretation, significantly improving cross-functional collaboration.

02/2017 – 01/2020

SENIOR DATA ENGINEER, CAPITAL ONE

- Managed and mentored a team of 5 data engineers, increasing team productivity by 30% and consistently delivering projects ahead of schedule.
- Designed and implemented scalable data pipelines using **Apache Spark 2.2**, reducing processing workload by 25%.
- Introduced AI-driven data cleansing solutions using Python 3.6, improving enterprise data quality by 50%.
- Developed and maintained complex ETL workflows using Informatica PowerCenter 9.6 and PowerExchange, supporting continuous data capture and integration across Salesforce (Classic/Lightning), Oracle 11g, DB2, and ODS systems.
- Served as an SFDC integration consultant, designing cross-application data integrations and implementing reliable, business-aligned data mapping rules.
- Developed ELT workflows leveraging Spark and dbt-style modeling principles, improving data consistency and enabling faster iteration for analytics and reporting teams.
- Partnered with business and operational stakeholders to gather functional requirements and translate them into detailed technical specifications.
- Built CI/CD runtime environments using Docker 18.09 and Kubernetes 1.14 for automated build, test, and deployment pipelines.
- Designed and supported **AI/ML** feature pipelines using Python and Spark, enabling data scientists to train and deploy machine learning models with clean, versioned, and production-ready datasets.

11/2014 – 02/2017

DATA ENGINEER, COGNIZANT

- Automated and simplified data extraction processes using C#, Python 2.7, Python 3.4, 3.5, saving approximately 10 hours per week in manual effort.
- Streamlined code integration and release processes, reducing release cycles by 15%.
- Designed and developed reusable ETL frameworks using Informatica PowerCenter 8.6, 9.5 to support enterprise-scale data workflows.
- Built relational data models and table structures using Oracle 10g, 11g and SQL Server 2012, 2014 to support customer profiling based on geography, asset classes, credit ratings, and business domains.
- Performed data analysis and ongoing maintenance of production datasets to ensure data quality and reliability.
- Supported enterprise BI initiatives by delivering curated datasets consumed by Power BI and Looker, ensuring metric consistency and data governance across domains.
- Designed ETL strategies for load balancing, exception handling, and high-volume data processing using Hadoop 2.6, 2.7 ecosystems.
- Supported infrastructure expansion by setting up new servers and executing structured data migration plans.

EDUCATION

BACHELOR'S DEGREE, NEW YORK UNIVERSITY - 2014

- Concentration: Database Management Systems

CERTIFICATIONS

- Software Engineer Certificate – Hackerrank
- SQL (Advanced) Certificate – Hackerrank
- Python (Basic) Certificate - Hackerrank