

BRIAN CAPLE

Staff Data Engineer

Peoria, IL | caplebrian2292@gmail.com | 352 225 9111

SUMMARY

Staff Data Engineer with 10+ years of experience designing, building, and optimizing large-scale data platforms on cloud environments (AWS, GCP, Azure). Expert in architecting **data lakehouse and warehouse solutions**, developing **batch and real-time pipelines** using Spark, PySpark, Kafka, and Airflow, and implementing **scalable ELT/ETL frameworks**. Skilled in **data modeling, governance, and performance optimization**, with a proven track record of delivering high-quality, mission-critical datasets for analytics and AI/ML initiatives. Adept at leveraging **cloud-native tools, DevOps practices, and automation** to drive operational efficiency and enable enterprise-scale data solutions.

WORK EXPERIENCE

Staff Data Engineer

02/2022 – Present

InvoZone - Pembroke Pines, FL

- Architected **enterprise-scale batch and streaming data pipelines** on **Azure**, leveraging **Azure Data Factory, Databricks, Event Hubs, Synapse**, and **Data Lake Storage** to support mission-critical analytics platforms.
- Designed and implemented **Databricks Lakehouse architecture** using **Delta Lake** with Bronze/Silver/Gold layers to enable governed, high-performance analytics.
- Built **parameterized, metadata-driven Azure Data Factory pipelines** with event-based triggers for scalable ingestion of structured and semi-structured data (JSON, CSV, XML, APIs, SQL).
- Developed **PySpark and Python pipelines** within Databricks for orchestrating ETL/ELT workflows and stream processing.
- Implemented **data quality frameworks** using Python and SQL validation scripts with automated reconciliation and audit logging.
- Applied **dimensional modeling and SCD Type 2 strategies** in Synapse and Delta Lake to maintain historical data integrity.
- Orchestrated ELT workflows using **Azure Data Factory, Databricks jobs**, and **Airflow DAGs** with dependency management and retry strategies.
- Containerized ETL workloads with **Docker** and deployed pipeline infrastructure on **Kubernetes**.
- Automated CI/CD for data pipelines using **Azure DevOps**, Jenkins, and Terraform infrastructure-as-code.
- Integrated RESTful services using Python and **Azure Functions** for automated ingestion from third-party systems.
- Monitored pipelines using **Azure Monitor**, custom dashboards, and alerts to ensure SLA compliance.
- Mentored engineers on **Databricks best practices**, lakehouse architecture, and distributed data processing.

Senior Data Engineer

11/2017 – 01/2022

United Data Technologies - Orlando, FL

- Designed and delivered enterprise cloud data platforms supporting **multi-terabyte analytics workloads on Azure and GCP**.
- Developed **large-scale ETL/ELT pipelines** using **Azure Data Factory, Databricks, PySpark, Airflow**, and **Talend**.
- Migrated on-premises data warehouses to **Azure Synapse, Snowflake, and BigQuery**, optimizing schemas for performance.
- Built **streaming ingestion pipelines** using **Azure Event Hubs, Databricks Structured Streaming, and Kafka** into data lake and cloud storage.
- Engineered **distributed Spark jobs** with Spark SQL, RDDs, and DataFrames for high-volume data processing.
- Implemented automated ingestion workflows with **Python, Airflow, and Azure Functions**.
- Designed Synapse and Hive data models with **advanced partitioning and clustering strategies**.
- Built audit orchestration and validation microservices using **Python, Spring Boot, and YARN scheduling**.
- Developed **CI/CD pipelines in Azure DevOps and GitHub Actions** for automated ETL deployments.
- Partnered with architecture and DevOps teams to define **best practices for cloud data engineering on Azure**.

Data Engineer

10/2013 – 09/2017

Altoros - Pleasanton, CA

- Built **scalable ETL frameworks** and **distributed Spark applications** on **Azure Databricks and EMR**, using PySpark and Spark SQL.
- Migrated legacy **SQL Server and Oracle warehouses** to **Azure Synapse and Redshift**, optimizing schema and performance.
- Developed **SSIS ETL workflows** integrating SQL Server, flat files, and external sources.
- Modeled **data marts using star schema and dimensional modeling**.
- Implemented automated job scheduling with **Azure Data Factory triggers, shell scripting, and Autosys**.
- Processed **NoSQL datasets** (HBase, Cassandra) for high-throughput analytics.
- Built CI/CD-enabled ETL workflows using **Azure DevOps, Jenkins, and Git**.
- Developed reusable **data validation and governance frameworks** for Azure and cloud environments.

SKILLS

- **Cloud Platforms:** Azure (Data Factory, Databricks, Synapse, Event Hubs, Functions, DevOps), AWS (Glue, EMR, Lambda, S3), GCP (BigQuery, Composer, Cloud Functions)
 - **Data Processing:** PySpark, Spark SQL, Scala, Python, SQL, ETL/ELT workflows, Stream & Batch processing
 - **Data Architecture:** Lakehouse (Bronze/Silver/Gold), Medallion architecture, Dimensional Modeling, SCD Type 2, Data Marts
 - **DevOps & CI/CD:** Azure DevOps, Jenkins, GitHub Actions, Terraform, Docker, Kubernetes
 - **Data Integration:** APIs, RESTful services, heterogeneous data sources (JSON, XML, CSV, RDBMS)
- Other Tools:** Airflow, Talend, Collibra, Hive, Synapse, Snowflake

EDUCATION

Princeton University – Princeton

Bachelor's Degree in Computer Science

05/2013