

MOSES ONAMADE

Senior Data Engineer

bmo.cybertech@gmail.com | +1 (470) 504-2715

PROFESSIONAL SUMMARY

Experienced **Senior Data Engineer** with over **9 years** of hands-on expertise in designing and implementing scalable **data pipelines**, **real-time data processing**, and optimizing data systems across industries like **music streaming**, **gaming**, and **financial services**. Proven track record in building and maintaining **large-scale data platforms** and leading cross-functional teams to drive **data-driven insights**. Strong experience in integrating **machine learning (ML)** models, **cloud services**, and **database technologies** to solve complex business problems. Skilled in streamlining workflows, reducing costs, and increasing performance, with a focus on **ML integration**, **cloud services**, and **database technologies**.

SKILLS

Data Engineering Tools: Scio, Dataflow, Luigi, Apache Spark, Kafka, dbt, Airflow, Databricks

Machine Learning: Supervised & Unsupervised Learning, Feature Engineering, Model Deployment, Time Series Forecasting, ML Frameworks (Scikit-learn, H2O.ai, TensorFlow, PyTorch)

Cloud Platforms: Google Cloud, AWS (EC2, S3, Lambda, RDS, Glue), Azure

Programming: Python, SQL, Java, Scala, Bash

Real-time Data Processing: Scio Streaming, Kafka, Spark Streaming

Data Warehousing: BigQuery, Snowflake, Redshift

Database Management: PostgreSQL, MySQL, NoSQL (MongoDB, DynamoDB)

DevOps: Docker, Kubernetes, CI/CD, GitHub Actions, Helm, Docker Swarm, Terraform

BI & Visualization: Looker, Tableau, Sigma, PowerBI, Metabase

PROFESSIONAL EXPERIENCE

Acer America – Senior /Staff Data Engineer | AI Automation Lead **02/2025 – Present**

- Architected a multi-tenant, cloud-native data platform (GCP, Dataflow, BigQuery, Kafka, Vertex AI) processing 5+ TB/day, enabling real-time fan engagement analytics across 1,500+ artists and labels.
- Defined and led the technical roadmap for AI-driven fan behavior prediction systems, deploying production ML pipelines that improved engagement CTR by 28% and boosted release-day conversions by 34%.
- Designed an enterprise-scale analytics framework (Spotify Wrapped-style for creators) delivering automated KPI intelligence to executive stakeholders, reducing manual insights generation by 85%.
- Built low-latency streaming architectures (Kafka, Pub/Sub, Spark Streaming) achieving sub-3-second end-to-end SLA for live release tracking.
- Introduced AI-powered anomaly detection for streaming metrics, preventing revenue leakage and improving data accuracy by 45%.
- Established data governance standards (dbt, Great Expectations, CI/CD), improving data reliability to 99.9% pipeline SLA.
- Optimized cloud cost allocation and workload orchestration, reducing infrastructure spend by 32% annually (~\$1.8M savings).
- Mentored 6+ engineers and led architecture reviews, enforcing scalable design patterns and platform-first engineering culture.

- Partnered with Product, AI Research, and Marketing leadership to align data platform capabilities with monetization strategy.

Cardinal Health -Senior Data Engineer | AI & Streaming Architecture Lead 06/2022 – 01/2025 | New York, NY

- Led modernization of legacy batch systems into real-time distributed streaming architecture (Kafka, Dataflow), cutting processing time by 50% and saving \$2M+ annually.
- Designed and scaled royalty computation engines processing billions of transactions/month, ensuring financial accuracy across global reporting systems.
- Productionized ML forecasting pipelines for royalty prediction and anomaly detection, reducing payout discrepancies by 22%.
- Architected enterprise data marts (BigQuery, partitioning, clustering) reducing query time by 60% for executive reporting.
- Automated financial reconciliation workflows using AI-assisted validation frameworks, reducing manual audit effort by 70%.
- Defined DevOps and CI/CD standards for data pipelines, reducing deployment failures by 40%.
- Led cross-functional architecture councils driving platform scalability, compliance (SOX), and governance initiatives.
- Built distributed data ingestion and transformation frameworks (Databricks, Scio, Dataflow) processing terabytes daily into production analytics environments.
- Engineered real-time data delivery systems reducing reporting latency from 24 hours to <1 hour.
- Standardized dimensional modeling and partitioning strategies improving BI dashboard performance by 50%.
- Developed ML integration pipelines to support predictive financial reporting and operational forecasting.
- Automated observability and monitoring (custom Python alerting, SLA tracking), improving reliability from 92% to 99.7%.

Sunrise Community of Tennessee – Data Engineer | AI Systems Engineer 03/2019 – 04/2020

- Architected high-volume telemetry pipelines (Hadoop, Spark, Databricks) processing 10B+ events/month for real-time player analytics.
- Migrated batch architectures to streaming (Kafka, Spark Structured Streaming), reducing analytics latency by 75%.
- Designed ML-based churn and monetization prediction systems, increasing in-game revenue by 18%.
- Built automated feature engineering and training data pipelines reducing ML experimentation cycle by 35%.
- Implemented fraud detection models and anomaly monitoring systems improving platform security.
- Led development of enterprise-scale streaming pipelines serving 45M+ customers (Capital One ecosystem) with 99.9% SLA.
- Designed personalization engines powered by ML and streaming data, improving customer engagement by 27%.
- Architected chatbot analytics pipelines (Eno), increasing automation accuracy by 30%.
- Established Spark optimization and data partitioning standards reducing compute costs by 25%.

EDUCATION

The University of Texas Permian Basin, Odessa, TX

May 2018