

James Jones

Data Platform Engineer

jamesjonesok2@gmail.com | 94102

SUMMARY

Senior Data Engineer with 10+ years of experience designing, building, and optimizing **cloud-native data platforms**. Expert in **Azure Data Factory, Databricks, Synapse, Event Hubs, and Lakehouse architectures**, with strong skills in **PySpark, Python, SQL, ETL/ELT pipelines, streaming, and CI/CD automation**. Proven track record delivering **high-performance batch and real-time data pipelines**, improving data quality, reducing latency, and enabling actionable analytics. Adept at translating complex business requirements into scalable data solutions, mentoring teams, and implementing best practices for governance, security, and operational excellence.

SKILLS

- **Programming & Query Languages:** Python, SQL (PostgreSQL, SQL Server, Oracle, T-SQL, PL/SQL, SQLScript), Scala, Java, R, TypeScript, C#
- **Data Engineering & Processing:** Apache Spark, PySpark, Pandas, dbt (Core & Cloud), Apache Airflow, **Azure Data Factory**, SSIS, Informatica, Talend, Fivetran, Matillion, SAP Datasphere Flows, SAP SLT, Batch & Streaming Pipelines, CDC Pipelines
- **Data Modeling & Warehousing:** Dimensional Modeling (Star & Snowflake), Slowly Changing Dimensions (SCD Type 1 & 2), Fact & Dimension Design, Data Vault, **Lakehouse Architecture**, Medallion Architecture, Schema Design, Enterprise Data Warehousing, Analytical Modeling, Semantic Layer Design, OLAP/OLTP Design
- **Cloud Data Platforms:** **Azure** (Databricks, ADLS, Synapse, Data Factory, Functions, SQL, Storage, Event Hubs, Purview), AWS (S3, Glue, Redshift, Lambda, Kinesis), GCP (Dataproc, BigQuery, Dataflow, GCS), Databricks Lakehouse Platform, Snowflake
- **Streaming & Big Data Ecosystem:** Apache Kafka, **Azure Event Hubs**, Pulsar, Apache NiFi, Spark Streaming, Flink, Hadoop, Hive, Delta Lake, Iceberg, Hudi, Trino, Presto
- **Databases & Storage:** Snowflake, **Azure Synapse**, Databricks Delta Lake, PostgreSQL, SQL Server, Oracle, MongoDB, DynamoDB, CosmosDB, Elasticsearch, ADLS, S3, GCS, HDFS
- **Machine Learning Platform Enablement:** **Azure ML**, Databricks ML, AWS SageMaker, MLflow, Feature Stores, Model Monitoring, Batch Inference Pipelines, Experiment Tracking
- **DevOps & Infrastructure:** **Azure DevOps**, Terraform, ARM Templates, Docker, Kubernetes, GitHub Actions, GitLab CI/CD, Jenkins, Infrastructure as Code, CI/CD for Data Pipelines
- **Governance, Security & Compliance:** Data Governance, Data Quality Frameworks, Metadata Management, Lineage, Data Catalogs, RBAC, Encryption at Rest/In Transit, HIPAA, GDPR, SOC2, CCPA, PII/PHI Protection, Unity Catalog, SAP Authorization Concepts
- **Business Intelligence & Analytics:** Power BI, Tableau, Looker, SAP Analytics Cloud, SSRS, SSAS, Qlik, Domo, KPI Modeling, OLAP Cubes, Semantic Models, Self-Service BI

PROFESSIONAL EXPERIENCES

Data Platform Engineer | *Blend, San Francisco, CA* | Nov 2021- Present

- Led development of high-performance **Azure Databricks** ETL pipelines using Python and PySpark for structured and semi-structured datasets, achieving 40% faster processing times.
- Designed and implemented **real-time streaming pipelines with Azure Event Hubs and Databricks Structured Streaming**, enabling continuous financial data ingestion.
- Architected **cloud-native data architectures on Azure** using ADLS, Synapse Analytics, and Databricks for scalable, secure, and cost-efficient storage and querying.
- Migrated legacy on-prem Hadoop and SQL workloads to **Azure Databricks**, refactoring Hive and SQL jobs into optimized Spark transformations.
- Applied performance tuning on Spark jobs and SQL queries in Synapse and Databricks, reducing ETL runtime and memory usage by 35%.
- Designed and managed **Azure Data Lake** architecture, integrating batch and streaming workflows from ADLS, Synapse, and third-party APIs.
- Built and orchestrated complex ETL workflows with **Azure Data Factory** and Airflow.
- Developed machine learning pipelines using **Azure ML and Databricks ML**, enabling model training and batch inference within ETL processes.
- Implemented **data governance, encryption, masking, and auditing in Azure Purview**, meeting GDPR, SOC 2, and other compliance standards.
- Built reusable dimensional and data vault models for financial reporting, improving downstream analytics efficiency.
- Mentored a team of 5–7 engineers on Azure, Spark, and CI/CD best practices.
- Implemented **CDC pipelines** for near real-time data synchronization between on-prem and cloud systems.

Data Architect | *Allstate, Austin, TX* | Oct 2018 – Sep 2021

- Designed and developed batch and real-time ETL pipelines using **Azure Data Factory, Databricks, and Synapse**, replacing legacy MapReduce and Pig workloads with Spark SQL, improving processing speed by 30%.
- Built predictive and clustering pipelines using Python, Spark MLlib, and Azure ML for customer segmentation and churn prediction.
- Led **real-time data ingestion** using **Azure Event Hubs and Databricks**, enabling timely analytics.
- Migrated large-scale datasets from on-premise to **ADLS and Synapse**, optimizing for performance and cost.
- Developed **infrastructure automation on Azure** using ARM Templates and Terraform for consistent deployment and scaling.
- Implemented dimensional modeling and Data Vault approaches for enterprise data warehouse design.
- Ensured data quality, governance, and compliance across all Azure ETL processes.
- Optimized Spark, Synapse, and Databricks queries for performance and cost efficiency.
- Mentored a team of 4–6 engineers on Azure, Spark, and big data best practices.

Data Engineer | *Mindsource, Mountain View, CA* | Sep 2013 – Sep 2018

- Developed scalable **ETL pipelines on Azure Databricks and ADLS** for distributed data processing.
- Designed and implemented logical and physical data models for OLTP and OLAP systems.
- Built and optimized **Azure Data Lake** and Synapse warehouses, integrating on-premise and cloud datasets.
- Created complex Spark SQL and Databricks transformations, optimizing for large-scale dataset processing.
- Developed automated data ingestion pipelines with **Azure Data Factory** and Airflow.
- Migrated legacy datasets to **Azure cloud storage**, ensuring integrity and compliance.
- Applied partitioning, bucketing, and dynamic partitions in Spark to optimize large dataset processing.
- Built custom Spark applications for RDD and DataFrame operations, including ML workflows.
- Deployed CI/CD pipelines using **Azure DevOps and Jenkins** for automated ETL deployments.

EDUCATION

Bachelor of Science in Computer Science University of California, Berkeley -May 2013