

Robin Patel

Address: Chicago, IL 60626 **Phone:** (678) 369-1066 **Email:** robinpatel3022@outlook.com

LinkedIn: <https://www.linkedin.com/in/robin-patel-4989293b1/> **Visa Status:** US Citizen

SUMMARY

Senior Data Engineer with over 10 years of experience architecting and implementing enterprise data strategies, specializing in Azure and Spark-based solutions.

Adept at designing scalable ETL/ELT pipelines, optimizing data quality, reliability, and automation, and leading cross-functional teams to deliver robust data platforms.

Demonstrated expertise in dimensional modeling, data migrations, and integrating data lakes and data warehouses to support advanced analytics and reporting.

Proven leadership in mentoring data engineers, driving best practices, and aligning technical solutions with business objectives.

EDUCATION

University of Florida

Gainesville, FL

August 2016 - May 2018

Master's degree, Data Science

University of Miami

Coral Gables, FL

September 2008 - June 2012

Bachelor's degree, Computer Science

SOFT SKILLS

- Effective Communication and Articulation
- Teamwork and Collaboration
- Empathy and Emotional Intelligence
- Creativity and Innovation
- Continuous Learning and Skill Development
- Client-Centric Approach and Customer Service
- Analytical mind

TECHNICAL SKILLS

- **Programming:** SQL, Python
- **Data Engineering & Big Data:** ETL/ELT Pipelines, Apache Spark, PySpark, Spark Streaming, Hive, Pig, DBT, Talend
- **Streaming & Messaging:** Kafka, RabbitMQ
- **Cloud Platforms:** Azure Data Factory, Databricks, Blob Storage, AWS EC2, EKS, S3, Glacier, Redshift, Glue, Athena, CloudFormation, Elastic Beanstalk, CloudWatch
- **Databases & Warehouses:** Snowflake, Redshift, SQL Server, Oracle, PostgreSQL, MongoDB, Redis, Elasticsearch
- **Orchestration & Containers:** Docker, Kubernetes
- **APIs & Services:** Flask, FastAPI, REST APIs
- **BI & Analytics:** Power BI, Tableau, yBI
- **Monitoring & Observability:** Prometheus, Grafana, Elasticsearch, CloudWatch
- **Automation & DevOps:** CI/CD, Selenium, Chef, Power Automate, WeAutomate, Ant Robot

- **Data Quality & Governance:** Data Validation, Reconciliation Frameworks, Incremental Models, Snapshots, Testing
- **Extra:** Azure, ADLS, Azure Databricks, Spark, Synapse Analytics, Azure SQL Database, Fabric, OAuth 2, Dimensional modeling, Data lakes, Data warehouses

EXPERIENCE

Senior Data Engineer (Tech Lead)

Gainwell Technologies

July 2019 – Current

Irving, Texas

- Architected and optimized scalable ETL/ELT pipelines using Apache Airflow, Spark (PySpark), Talend, and DBT, enhancing data accessibility, pipeline reliability, and analytics performance across diverse business domains.
- Designed and implemented cloud-based data platforms with AWS S3, Glacier, Redshift, Glue, and Athena, applying lifecycle management and partitioning strategies to reduce storage costs by 20% and improve query efficiency.
- Developed incremental DBT models, snapshots, and materializations, establishing best practices for transformations, lineage, and testing, resulting in a 40% reduction in processing time.
- Orchestrated complex Airflow DAGs to manage Spark, Hive, and Pig jobs with robust dependency management, retries, and SLA monitoring, ensuring reliable data availability.
- Engineered Python-based REST APIs (Flask, FastAPI) to automate data ingestion, validation, and quality checks, reducing manual intervention by 50% and improving data accuracy.
- Led large-scale data migrations across SQL Server, Oracle, PostgreSQL, MongoDB, and Snowflake, maintaining 99.9% data integrity through comprehensive reconciliation and validation frameworks.
- Optimized Spark-based batch and streaming pipelines by tuning memory usage, partitioning, and serialization, increasing processing efficiency and reducing operational overhead.
- Built automated ingestion frameworks integrating Snowflake, Redis, Elasticsearch, Kafka, and S3, enabling near real-time data availability for analytics and search applications.
- Implemented real-time streaming pipelines with Kafka and Spark Streaming to support low-latency analytics and downstream microservices.
- Designed and managed cloud-native architectures on AWS (EC2, EKS, S3, Glue) and Azure (Data Factory, Databricks, Blob Storage), optimizing scalability, fault tolerance, and cost efficiency.
- Containerized and orchestrated workloads using Docker and Kubernetes, improving deployment consistency, environment parity, and achieving 55% better scalability.
- Enhanced BI pipelines for Redshift and Snowflake by optimizing schema design, clustering, and query tuning to improve dashboard response times.
- Established centralized logging, monitoring, and alerting with Elasticsearch, Prometheus, Grafana, and CloudWatch, increasing observability of data pipelines.
- Integrated ML-driven log analytics with Chef and custom Python services to detect anomalies and accelerate incident resolution.
- Mentored and led a team of data engineers and analysts, setting standards for code quality, documentation, and CI/CD, and ensuring timely project delivery.

Technologies: Apache Airflow, Apache Spark / PySpark, DBT, AWS Data Platform (S3, Redshift, Glue, Athena), Kafka & Spark Streaming, Snowflake, Docker & Kubernetes, Elasticsearch, Prometheus, Grafana

Data Engineer / Automation Engineer

Dentsu

October 2015 – June 2019

New York, NY

- Managed end-to-end data projects from ingestion to analytics consumption, delivering predictive maintenance and operational analytics solutions that generated \$1.2M ROI.

- Designed and implemented complex ETL/ELT workflows using Talend, SQL, and Python, reducing processing time from 8 hours to 2 hours and improving data freshness.
- Built data automation pipelines with Python (pandas, openpyxl, pywin32) and SQL to streamline data transformation and reporting, increasing efficiency by 30%.
- Developed process automation frameworks using WeAutomate, Power Automate, and Ant Robot, eliminating 40% of manual operational tasks.
- Integrated and optimized data pipelines on Marshalls Data Lake to improve query response times and analytics performance.
- Implemented advanced transformation logic and data validation rules, enhancing data consistency and reducing downstream errors.
- Built and enhanced dashboards with Power BI, Tableau, and yBI, enabling self-service analytics for business stakeholders.
- Expanded automated test coverage using Selenium and Django, improving deployment stability and reducing defects.
- Collaborated with cross-functional teams to design automation tools and web-based solutions that addressed business requirements.
- Led a small team in data automation initiatives, streamlining workflows and improving delivery efficiency.
- Achieved 99% deployment success rates across HIS, cloud, and Marshalls platforms through robust automation and testing.

Technologies: *ETL/ELT with Talend, Python Data Automation, Power Automate, WeAutomate, Ant Robot, Data Validation & Transformation Logic, Power BI & Tableau, Pandas-based Reporting*

Software / Data Engineer

MilesIT

July 2012 – August 2015

Austin, Texas

- Executed over 50 AWS CloudFormation templates and orchestrated Elastic Beanstalk deployments, reducing deployment time by 30% and enabling scalable Python and Docker-based applications.
- Optimized AWS storage architectures using S3 and EBS, improving retrieval speed by 25% and reducing costs by 15%.
- Built and maintained Python REST APIs (Flask, SQLAlchemy) to process 500K+ daily transactions on PostgreSQL databases.
- Implemented monitoring and alerting for Elasticsearch, improving query response times by 40% and ensuring 99.9% uptime.
- Developed automated data retrieval programs for AWS-hosted systems, increasing client operational efficiency.
- Automated machine alarm data ingestion and timestamping using Python and optimized SQL to streamline reporting.
- Built data analysis and reporting workflows with Python (Pandas), reducing reporting time by 50%.
- Led Agile teams and managed delivery for more than 10 projects, achieving high client satisfaction.

Technologies: *AWS CloudFormation & Elastic Beanstalk, PostgreSQL, REST APIs Flask, SQLAlchemy, Docker, SQL Optimization, AWS S3 & EBS*

CERTIFICATIONS

- AWS Certified Data Engineer – Associate
- Google Cloud Professional Data Engineer
- Databricks Certified Data Engineer Professional
- SnowPro Core / SnowPro Advanced Data Engineer

AWARDS

- Arctic Code Vault Contributor (Github)