

Siva Krishna Bandi

Data Engineer

USA | +1 (508) 603 9624 | sivakrishnabadi20@gmail.com | [LinkedIn](#)

PROFESSIONAL SUMMARY

Senior Data Engineer with **5+ years** of experience designing and delivering scalable batch and real-time data platforms across financial services and healthcare domains, including roles at Goldman Sachs, CVS Health, and Morgan Stanley. Proven track record of processing multi-terabyte datasets using Apache Spark, Kafka, and cloud-native services on AWS and Azure, reducing pipeline failures by ~40% and improving query performance by 30%+. Highly skilled in ETL/ELT pipeline development, Snowflake and Redshift data modeling, Airflow orchestration, and data governance, with strong expertise in Python, SQL, and CI/CD-driven data engineering to enable secure, reliable, and business-ready analytics.

PROFESSIONAL EXPERIENCE

Senior Data Engineer | Goldman Sachs | Boston, MA, USA

Jul 2024 – Present

- Designed and scaled Spark (Scala/PySpark) data pipelines on AWS, processing multi-terabyte financial datasets and publishing curated Snowflake tables for regulatory and enterprise reporting use cases.
- Enabled batch and near real-time ETL workflows using AWS Glue, S3, Lambda, and Redshift to support enterprise banking data ingestion pipelines.
- Orchestrated 100+ daily production pipelines using Apache Airflow with retries, SLA alerts, and dependency management to support business-critical workflows.
- Enabled real-time streaming ingestion via Apache Kafka, supporting high-volume event processing and delivering low-latency datasets for fraud detection, BI, and compliance analytics.
- Led Teradata-to-Snowflake migration efforts, managing schema conversion, data validation, clustering strategies, and query tuning for high-concurrency analytical workloads.
- Established data security and governance standards using Snowflake row-level security, column masking, IAM, and Ranger policies to ensure GDPR-compliant data access.

Data Engineer | CVS Health | Bengaluru, India

Dec 2020 – Aug 2023

- Architected a cloud data platform on Azure (ADLS Gen2, Azure Data Factory, Databricks) using medallion architecture to support enterprise healthcare analytics.
- Developed Spark-Scala ETL pipelines in Azure Databricks, transforming large healthcare datasets (AVRO, Parquet, CSV) from claims, eligibility, and provider systems.
- Orchestrated parameterized Azure Data Factory pipelines for ingestion, transformation, validation, and alerting across multiple source systems and environments.
- Modernized legacy Informatica workflows by migrating pipelines to ADF and Databricks, reducing operational complexity and improving pipeline reliability.
- Designed dimensional Star Schema models in Snowflake, enabling high-performance Power BI dashboards for actuarial and operations teams.
- Embedded data quality, governance, and security controls using Great Expectations, Azure Purview, Snowflake masking policies, and Key Vault to ensure HIPAA-compliant data processing.

Data Engineer | Morgan Stanley | Bengaluru, India

Aug 2019 – Nov 2020

- Built a hybrid AWS data platform using S3, AWS Glue, Redshift, and Athena, organizing raw, staging, and curated layers for financial analytics workloads.
- Engineered PySpark ETL jobs on AWS Glue, ingesting transactional data from SFTP and REST APIs into optimized Redshift tables.
- Delivered real-time streaming pipelines using Kafka and Spark Structured Streaming to enable near real-time risk and compliance reporting.
- Optimized Redshift data marts using dimensional modeling, sort keys, distribution styles, and WLM tuning to support high-performance analytical queries.
- Automated infrastructure provisioning using Terraform, deploying Glue jobs, S3 buckets, IAM roles, and Lambda functions as code.
- Automated CI/CD pipelines using GitHub Actions to enforce testing, schema validation, and controlled deployments across dev, QA, and production environments.

SKILLS

- Data Engineering & Architecture:** Large-Scale Data Pipelines, Batch & Real-Time Processing, Distributed Systems, Data Lake & Lakehouse Design, Medallion Architecture, Event-Driven Architectures
- Big Data & Streaming:** Apache Spark (PySpark, Spark-Scala), Spark Structured Streaming, Apache Kafka, Apache Flink, Hadoop Ecosystem (HDFS, Hive, YARN), Schema Evolution
- Cloud Platforms & Services:** AWS (S3, Glue, EMR, Lambda, Redshift, Athena, Step Functions), Azure (ADLS Gen2, Data Factory, Databricks, Synapse), GCP BigQuery (POC)
- ETL / ELT & Orchestration:** ETL & ELT Pipeline Development, Apache Airflow DAGs, Azure Data Factory Pipelines, AWS Glue Jobs, Workflow Scheduling, SLA & Dependency Management
- Data Warehousing & Modeling:** Snowflake Architecture, Redshift Optimization, Dimensional Modeling (Star, Snowflake, Data Vault), OLAP / OLTP Design, Query & Cost Optimization
- Programming & Query Languages:** Python, SQL (Advanced & Performance Tuning), Scala, Java, Bash / Shell Scripting, REST API Integration
- Data Quality, Governance & Security:** Data Validation Frameworks, Great Expectations, Metadata & Lineage Management, RBAC, Row-Level & Column-Level Security, GDPR & HIPAA Compliance
- DevOps, CI/CD & Infrastructure:** CI/CD for Data Pipelines, GitHub Actions, Jenkins, GitLab CI, Docker, Kubernetes (AKS/EKS), Terraform, CloudFormation
- Monitoring & Reliability Engineering:** Pipeline Monitoring & Alerting, CloudWatch, Grafana, Prometheus, ELK Stack, SLA Tracking, Incident Resolution
- Visualization & Analytics Enablement:** Power BI, Tableau, Looker, AWS QuickSight, KPI Dashboards, Business Reporting, Data-Driven Decision Support
- Process & Optimization:** Performance Tuning, Cost Optimization, Workflow Automation, Compute Resource Optimization, Scalable System Design
- Collaboration & Delivery:** Agile & Scrum Methodologies, Cross-Functional Collaboration, Stakeholder Communication, Technical Documentation, Knowledge Transfer

EDUCATION

Clark University, USA

Aug 2023 – May 2025

Master's in Information Technology