

Jordan Zaman

Senior Data Engineer

+ 1 (904)-580-6008

jordanzaman196@outlook.com

Canton, NY

SUMMARY

Senior Data Engineer with over 10 years of experience in the **software industries**, specializing in **data architecture, pipeline development, and large-scale data processing**.

Expertise in designing and optimizing **scalable, high-performance data solutions**, leveraging technologies such as **Apache Spark, Hadoop, Kafka, AWS, Azure, GCP, Snowflake, DBT, and PostgreSQL**. Proficient in **ETL/ELT workflows, real-time and batch data processing, and cloud-native data engineering**.

Skilled in leading **cross-functional teams**, collaborating with **stakeholders**, and driving **data-driven decision-making** by developing efficient, **secure, and reliable data infrastructures**. Passionate about building **scalable data ecosystems** to support **business intelligence, AI/ML models, and enterprise analytics**

EDUCATION

Stanford University | Bachelor of Science (BS), Computer Science | 2011- 2015

TECHNICAL SKILLS

- **Programming:** Python, SQL, Java, Scala, Shell Scripting, JavaScript, PHP
- **Cloud Platforms:** AWS (S3, Lambda, Step Functions, SageMaker), GCP (BigQuery, Vertex AI), Azure
- **Data Engineering & Pipelines:** ETL/ELT, Apache Kafka, Apache Spark, PySpark, Databricks, DBT, AWS Glue, Snowflake, Redshift, Delta Lake, Data Lakes, Kafka Streams, AWS Kinesis
- **Databases:** PostgreSQL, MongoDB, Cassandra, Oracle, NoSQL
- **Machine Learning & Data Science:** Scikit-learn, TensorFlow, PyTorch, Pandas, NumPy, MLOps, Databricks, Model Deployment
- **DevOps & Tools:** Docker, Kubernetes, Terraform, Jenkins, Git, CI/CD pipelines
- **Monitoring & Logging:** Prometheus, Grafana, Datadog, CloudWatch, Airflow
- **Data Visualization:** Tableau, PowerBI, Looker, Seaborn
- **Data Modeling:** Dimensional Modeling, Star Schema, Snowflake Schema, Data Vault
- **Soft Skills:** Cross-Functional Collaboration, Mentorship, Agile Methodologies, Stakeholder Communication

PROFESSIONAL EXPERIENCE

Senior Data Engineer

Accenture, New York City, NY | June 2022 – present

- **Engineered and managed** highly scalable data platforms by integrating **Apache Spark, Databricks, and Kafka** with **AWS** and **GCP** services, enhancing operational intelligence across multiple business functions.
- Spearheaded the design and deployment of **Databricks Lakehouse** architectures, optimizing **Delta Lake** and **Unity Catalog** for data reliability, efficient schema evolution, and reduced latency in delivering real-time analytics.

- Developed **streaming pipelines** using **Spark Structured Streaming** and **Kafka**, powering interactive dashboards and automating event-driven systems to facilitate real-time decision-making and operational agility.
- Led the implementation of **Airflow-based orchestration** with robust features like SLA enforcement, retries, backfills, and alerting, ensuring mission-critical data workflows were executed reliably across environments.
- Played a key role in **multi-cloud initiatives** across **AWS** and **Azure**, overseeing seamless cloud migrations and optimizing hybrid ingestion strategies, enabling flexible, cross-cloud access for diverse datasets.
- Built and optimized **ETL/ELT pipelines** in **Python** and **SQL**, automating the ingestion of structured, semi-structured, and unstructured data from multiple sources (APIs, relational databases, flat files, streams).
- Delivered curated **analytics-ready datasets** to fuel BI reporting, enhance operational KPIs, and enable advanced machine learning and AI workloads, supporting data-driven decision-making.
- Collaborated with **product engineering** teams to expose high-value data through **REST** and **gRPC** APIs, enabling secure, schema-driven integrations to meet cross-functional needs.
- Established and enforced **data quality frameworks**, proactively identifying pipeline issues and leveraging **anomaly detection** to ensure data integrity across the workflow.
- Integrated **enterprise-grade security controls**, including **RBAC**, encryption, and **PII masking/tokenization**, to ensure regulatory compliance and protect sensitive data assets.
- Pioneered the development of **CI/CD pipelines** using **GitHub Actions** and **Azure DevOps**, automating the deployment and validation of data infrastructure changes, enhancing the speed and reliability of updates.
- Optimized **Spark performance** and **SQL queries**, focusing on improving resource efficiency and minimizing computational costs, while meeting high-throughput, low-latency processing requirements.
- Supported **AI/ML teams** by preparing feature-rich datasets, providing key metrics for experiments, and ensuring smooth integration with production-level data feeds for continuous model training.
- Documented system architecture, data lineage, and operational processes to ensure transparency, reliability, and audit compliance across the data platform.
- Led technical initiatives by reviewing designs, conducting **code reviews**, and driving best practices in collaboration with cross-functional data teams to maintain high standards in development and operational efficiency.

Data Engineer

Salesforce, San Francisco, CA / *Mar 2018 – May 2022*

- Led the design and automation of **ETL workflows** using **Apache Airflow** and **AWS CloudWatch**, significantly improving the reliability, performance, and scalability of data pipelines. This reduced manual intervention and streamlined data processing workflows.
- Integrated real-time data feeds from multiple sources into **Tableau** and **other BI platforms**, ensuring that reports and dashboards were available with up-to-date information for immediate business decisions. This reduced report generation time by 30%.
- Collaborated with the cloud engineering team to migrate traditional data processing workflows to cloud-based solutions, improving data processing time and scalability across **AWS Redshift** and **Amazon S3** platforms.
- Built predictive models using **Scikit-learn** to support customer segmentation and churn prediction efforts, leading to more targeted marketing campaigns and improved customer retention strategies.
- Implemented automation scripts using **Python** to streamline the repetitive tasks of data collection, transformation, and reporting, saving over 20 hours per week across the engineering team.
- Created and optimized **SQL queries** to improve database performance for large-scale reporting, resulting in more efficient data retrieval and faster decision-making.
- Led the integration of third-party APIs and external data sources, enabling seamless data onboarding and processing for clients, improving business insights and reducing the time to insight.

- Optimized data pipelines for better efficiency and faster processing by improving the architecture of **SQL-based data models**, implementing indexing strategies, and improving schema design to enhance query performance.
- Established real-time monitoring dashboards using **AWS CloudWatch** and **Logstash**, allowing teams to quickly identify and resolve bottlenecks in **ETL workflows**.
- Worked closely with cross-functional teams, including data scientists, analysts, and business stakeholders, to ensure accurate, timely data delivery and alignment on project goals.
- Delivered high-impact insights by building data visualizations and dashboards in **Tableau**, enabling senior leaders to track KPIs and make data-driven decisions across various departments.

Junior Data Developer

Dropbox, San Francisco, CA | May 2016 – Feb 2018

- Designed and maintained **ETL/ELT pipelines** that processed large-scale datasets from multiple sources, ensuring the delivery of clean, accurate data for reporting and analysis.
- Automated key data analysis and processing tasks using **Python** and **SQL**, resulting in a 40% increase in team efficiency by reducing manual interventions and allowing the team to focus on higher-value analysis.
- Developed and deployed **regression models** in **R** and **Python** to predict credit risk and improve financial decision-making processes, enhancing risk assessment models used in the loan underwriting process.
- Built and maintained **interactive dashboards** and **data visualizations** in **Tableau** and **QlikView** to track key business metrics, helping non-technical stakeholders quickly interpret data and make informed decisions.
- Enhanced **SQL query performance** by optimizing database schema designs, improving report generation times by 50%, and facilitating better data access for end-users.
- Led the integration of **third-party data sources** and external APIs to provide additional context for financial risk analysis, improving decision-making and reducing risk exposure.
- Worked with cross-functional teams to align on project goals and ensure that the data infrastructure was scalable and adaptable to future business requirements.
- Facilitated collaboration between data engineers and product managers to define and implement data integration solutions for new digital features, improving customer engagement and retention.
- Improved cloud infrastructure by migrating legacy systems to **AWS**, enhancing the scalability and flexibility of data processing workflows.
- Ensured consistent data quality and governance by developing standardized data pipelines and enforcing data management best practices.
- Trained junior engineers on data transformation techniques and automation frameworks, improving team performance and fostering a culture of continuous learning and development.