

Hunter Davis

Phone: (321) 290-1047

Mail: h.davis.unter@gmail.com

Location: Remote

SUMMARY

Senior Data Engineer with 9+ years designing and operating production-grade data ingestion, transformation, and delivery pipelines. Expert in scalable ETL/ELT workflows, data lakehouse and data warehouse modeling, and data quality/observability for analytics and ML. Proficient in SQL, Python, Spark, and modern orchestration and infra tooling. Proven leader driving cross-functional collaboration, architecture decisions, and reliable, scalable data platforms.

SKILLS

Programming

- SQL, Python

Data Processing

- Apache Spark, PySpark, Pandas

Data Warehousing & Lakehouse

- Delta Lake, Snowflake, Redshift, PostgreSQL
- Dimensional modeling (star/snowflake)

Orchestration & Scheduling

- Flyte, Kubernetes, Terraform, Docker

Streaming & Messaging

- Kafka, Spark Structured Streaming

Cloud Platforms & Storage

- AWS: S3, EMR, EKS, RDS
- Multi-cloud familiarity; infrastructure as code

Data Quality & Observability

- Schema validation, de-duplication, data dictionaries, monitoring (Prometheus/Grafana)

Security & Collaboration

- Data lineage, auditability, cross-team collaboration, stakeholder engagement
-

EXPERIENCE

IntelliTech AI – Senior Data Engineer

Remote / San Diego, CA

June 2020 – February 2025

- Led end-to-end design and delivery of a centralized data platform ingesting event, transactional, and third-party datasets into a Delta Lake, enabling analytics and ML use cases.
- Built scalable PySpark pipelines processing 500M+ records/day with transformations, de-duplication, and enrichment before loading to Snowflake for reporting.
- Implemented data quality checks and automated validation (schema/NULL/range checks), reducing production data incidents by 60%.
- Containerized pipeline tasks and deployed on Kubernetes (EKS) with Helm charts and CI/CD, reducing manual deployments and environment drift.
- Automated infrastructure provisioning with Terraform for EKS clusters, S3 storage, and IAM policies to standardize environments across dev/stage/prod.

- Introduced Flyte for long-running ETL workflows, improving retry, lineage, and reproducibility for data jobs.
- Partnered with Data Science and Product teams to onboard 12+ datasets, defining SLAs, schemas, and access patterns to ensure downstream readiness.
- Mentored 3–4 junior data engineers; led design sessions and internal trainings to accelerate adoption of best practices.
- Instrumented pipeline monitoring and alerting (Prometheus/Grafana) to track job success rates, latency, and data freshness.

Tech Innovations Inc. – Data Engineer

Miami, FL

February 2018 – May 2020

- Designed and implemented ETL frameworks in Python and Spark to ingest CRM, application logs, and third-party feeds into a centralized data warehouse (Postgres -> Redshift).
- Created reusable ingestion templates and transformation libraries to harmonize fields across 3+ sources, reducing onboarding time by 40%.
- Wrote advanced SQL models for data marts supporting product analytics and recommender systems; optimized queries and partitioning to cut query cost and latency.
- Implemented automated de-duplication and dedup key strategies to improve data integrity for downstream ML models.
- Managed Kafka-based streaming ingestion for near-real-time event processing, integrating with Spark Structured Streaming and S3-backed storage.
- Coordinated with analysts and ML engineers to document schemas, data dictionaries, and data lineage for auditability and reproducibility.

Tech Innovations Inc – Data Engineer

Miami, FL

June 2015 – October 2017

- Built batch ETL pipelines using Python and Spark to convert raw operational files into analytics-ready tables consumed by stakeholders.
- Implemented file ingestion workflows (SFTP, scheduled pulls, API-based pulls) and standardized file validation and error handling.
- Developed and enforced data modeling standards (naming, partitioning, time-series handling) used across the analytics organization.
- Assisted in migrating legacy ETL jobs into containerized workloads on Kubernetes and established CI/CD pipelines for deployment.
- Performed root cause analysis on data failures, documented runbooks, and improved process resilience and recovery times.

EDUCATION

Florida International University

BSc. in Computer Science | September 2011 – May 2015

Notes: This resume emphasizes data platform architecture, large-scale data processing, data warehousing/lake-house modeling, data quality and observability, orchestration, and collaboration with analytics/ML teams—aligned with the senior data engineering needs at JourneyTeam. It highlights measurable outcomes (records/day, incident reduction, onboarding time) and leadership responsibilities (mentoring, design sessions, cross-functional alignment). While Azure-specific hands-on experience isn't listed, the candidate's cloud-agnostic data platform experience and familiarity with scalable, enterprise-grade data architectures are showcased to support rapid alignment with Microsoft data solution initiatives.
