Jack Jones

Senior Data Engineer

jjones8593@gmail.com | Jacksonville, FL

SUMMARY/OBJECTIVE

Senior Data Engineer with 10+ years of experience designing and scaling cloud-native data platforms across retail, manufacturing, and enterprise analytics environments. Proven expertise in building high-performance batch and real-time pipelines using **AWS, Spark, Kafka, Airflow, and modern lakehouse architectures**. Strong track record of optimizing data warehouses, improving pipeline reliability, reducing costs, and enabling near real-time business insights. Adept at translating complex business requirements into scalable, secure, and production-ready data solutions while driving data quality, governance, and CI/CD automation in enterprise environments.

EDUCATION

Bachelor Degree of **Computer Science** at University of Florida – May, 2013

TECHNICAL SKILLS

- **Programming & Query Languages:**
  Python, SQL, PySpark, Scala, Java, T-SQL, PL/SQL, R, TypeScript, C#
- **Data Engineering & Processing:**
  Apache Spark, PySpark, Batch & Streaming Data Pipelines, Kafka, AWS Kinesis, Spark Streaming, Airflow, dbt (Core), Pandas, ETL/ELT Frameworks
- **Data Modeling & Warehousing:**
  Dimensional Modeling (Star & Snowflake), Slowly Changing Dimensions (SCD Type 1 & 2), Fact & Dimension Design, Kimball Methodology, Data Vault Concepts, Schema Design, Enterprise Data Warehousing, Analytical & Reporting Models, Lakehouse Architecture, Medallion Architecture
- **Cloud Data Platforms:**
  AWS (S3, EMR, Glue, Lambda, Redshift, IAM, CloudWatch),
  Azure (Databricks, ADLS, Synapse Analytics, Data Factory),
  GCP (BigQuery, Dataproc, GCS),
  Databricks Lakehouse Platform, Snowflake
- **Streaming & Big Data Ecosystem:**
  Apache Kafka, AWS Kinesis, Hadoop, Hive, Delta Lake, HDFS
- **Databases & Storage:**
  Amazon Redshift, Snowflake, PostgreSQL, SQL Server, Oracle, MongoDB, Amazon S3, ADLS
- **BI, Analytics & Reporting:**
  Power BI, Tableau, OLAP Concepts, KPI Modeling, Semantic Layers, Self-Service BI Enablement
- **DevOps, CI/CD & Infrastructure:**
  Git, GitHub Actions, Jenkins, Docker, Terraform (IaC), CI/CD for Data Pipelines, Cloud Monitoring & Logging
- **Governance, Security & Compliance:**
  Data Quality Frameworks, Data Lineage & Metadata Management, RBAC, Encryption at Rest & In Transit, GDPR, SOC 2, PII Protection

EXPERIENCE

Boll & Branch

Senior Data Engineer **Oct 2020 – Present**

Houston, TX

- Scaled the company's analytics platform by designing **cloud-native data pipelines on AWS**, enabling reliable processing of tens of millions of daily retail and e-commerce events.
- Reduced end-to-end data latency by **40%** by migrating legacy ETL workloads to **Amazon S3, Redshift, and Spark-based processing**.
- Modernized the enterprise warehouse by re-architecting dimensional models **(Kimball)**, significantly improving query performance for executive and merchandising dashboards.
- Enabled near real-time insights by building streaming pipelines with **Kafka, AWS Kinesis, and Spark Streaming**, supporting marketing attribution and customer behavior tracking.
- Cut storage costs by **20%** through **data lifecycle management on S3**, optimizing partitioning and compression strategies.
- Improved data trust across teams by **implementing automated data quality checks and anomaly detection**, reducing downstream reporting errors by **30%**.
- Increased analyst productivity by optimizing **complex SQL transformations and Redshift workloads**, accelerating high-concurrency BI queries.
- Standardized pipeline reliability by introducing **Airflow-based orchestration**, eliminating ad-hoc job failures across multi-environment deployments.
- Supported advanced analytics initiatives by integrating **ML outputs into curated datasets**, enabling customer segmentation and demand forecasting use cases.
- Strengthened compliance posture by enforcing **GDPR and SOC 2 controls**, securing PII through encryption, access controls, and audited data flows.
- Accelerated release cycles by implementing **CI/CD for data pipelines using GitHub Actions and Jenkins**, reducing deployment risk and rollback time.
- Improved platform observability by introducing **data lineage and metadata tracking**, enabling faster root-cause analysis during incidents.
- Partnered with product, marketing, and finance leaders to **translate business questions into scalable cloud data solutions**, driving measurable improvements in decision-making speed.

Oshkosh Corporation

Data Engineer **Mar 2016 – Sep 2020**

Oshkosh, WI

- Built a centralized analytics foundation by engineering **Spark-based ETL pipelines** to process operational and manufacturing data at scale.
- Improved reporting performance by **50%** through **design and optimization of an AWS Redshift data warehouse** supporting enterprise BI workloads.
- Enabled real-time operational visibility by implementing **Kafka-driven event pipelines**, supporting manufacturing monitoring and alerting use cases.
- Reduced data processing bottlenecks by tuning **SQL transformations and warehouse schemas**, accelerating reporting across operations and finance teams.
- Designed scalable data architectures that **supported cross-functional analytics** for production, supply chain, and sales organizations.
- Lowered operational overhead by automating **ingestion and pipeline monitoring**, cutting manual intervention by **60%**.
- Unified siloed data sources by integrating **Oracle, SQL Server, and MongoDB systems**, enabling enterprise-wide analytics.
- Improved data reliability by introducing **structured validation and reconciliation checks** across critical datasets.

- Partnered with business stakeholders to **deliver tailored analytical datasets**, improving data adoption and self-service reporting.
- Strengthened security posture by implementing **encryption and RBAC controls**, protecting sensitive operational data in the cloud.
- Established foundational data engineering best practices that **supported long-term cloud adoption and scalability**.

Melaleuca, Inc.

Data Engineer **Nov 2013 – Feb 2016**

Knoxville, TN

- Supported business growth by building **Python- and SQL-based ETL pipelines** to process high-volume operational and customer data.
- Improved data freshness by introducing **near real-time Kafka workflows**, enabling faster visibility into customer interactions.
- Increased reporting efficiency by optimizing **PostgreSQL schemas and queries**, significantly reducing dashboard load times.
- Enabled smarter product decisions by integrating **predictive model outputs into analytics datasets**.
- Delivered timely business insights by automating **reporting pipelines for Power BI and Tableau users**.
- Reduced infrastructure costs by leading the migration **from on-prem systems to AWS S3 and Redshift**, cutting operational spend by **30%**.
- Improved workflow stability by implementing **Airflow-based scheduling and retry mechanisms**.
- Partnered closely with analysts to **translate reporting needs into reliable data models**.
- Established early cloud and big-data practices that **laid the foundation for future scalability and modernization**.