

Daniel Ho

Meridian, ID | (208) 203-3761 | daniel.h9025@gmail.com

SUMMARY

Senior Data Engineer with 10+ years designing, building, and optimizing large-scale data platforms across finance, healthcare, and enterprise SaaS. Expert in end-to-end data pipelines (ETL/ELT), data lakehouse architectures, and data quality/monitoring. Proven track record delivering high-throughput, low-latency analytics infrastructure on Spark/Databricks, Delta Lake, and cloud data warehouses. Strong architect, mentor, and cross-functional collaborator driving data-driven business outcomes with measurable improvements in data reliability, performance, and analytics adoption.

SKILLS

Programming: Python, SQL, Scala, Bash

Data Processing: Apache Spark, Databricks, Airflow, dbt

Data Warehousing / Lakehouse: Delta Lake, Parquet, Redshift, Snowflake, PostgreSQL

Orchestration: Airflow

Streaming / Data Integration: CDC, ETL/ELT pipelines

Cloud Platforms / DevOps: AWS (S3, Glue, EMR, Lambda, EC2, RDS), Terraform, Docker, CI/CD (GitHub Actions, Jenkins)

Data Quality / Monitoring: Great Expectations, Prometheus, Grafana

Data Modeling / Governance: Entity Resolution, Data Modeling, Data Lineage

EXPERIENCE

SingleStore | Senior Data Engineer | Feb 2021 – Feb 2025 | San Francisco, CA

- Designed and implemented Spark-based ETL pipelines on Databricks to process over 5 TB/day of multi-source data into Delta Lake, enabling scalable analytics.
- Built an entity resolution framework that merged 1B+ user records with 98% accuracy using PySpark and AWS Glue, improving customer analytics fidelity.
- Deployed data quality validation layers using Great Expectations and automated anomaly detection via Airflow DAGs, increasing data reliability and trust.
- Migrated legacy batch jobs to event-driven pipelines on AWS, reducing pipeline latency by 40% and compute costs by 25%.
- Collaborated with data scientists to productionize ML-ready datasets and streamline model deployment workflows, accelerating analytics outcomes.
- Led design discussions and contributed to architecture reviews to align data solutions with enterprise standards and client objectives.

Talentpair | Data Engineer | Nov 2017 – Jan 2021 | Remote

- Developed incremental and partitioned data ingestion pipelines using Spark and PostgreSQL for patient analytics systems, enabling near-real-time insights.
- Built and managed a data lakehouse architecture on AWS S3 and Databricks to support analytics on 10M+ records, improving data accessibility for analysts.
- Created custom Airflow operators to orchestrate ETL pipelines and automate schema evolution across data sources, reducing maintenance overhead.
- Enhanced data reliability with monitoring dashboards (Prometheus, Grafana) and robust schema validation tooling, increasing trust in analytics outputs.
- Led a data migration initiative consolidating disparate systems into a unified cloud-based data warehouse (Redshift), improving reporting consistency.

Seer Interactive | Data Integration Engineer | Nov 2014 – Oct 2017 | San Diego, CA

- Engineered data transformation and enrichment pipelines in Python and SQL for regulatory and analytics workflows, delivering timely insights.
-

Implemented CDC (Change Data Capture) strategies to synchronize operational databases with analytical stores, enabling up-to-date analytics.

- Designed and documented data schemas, lineage, and design specs, improving maintainability and cross-team visibility.
- Collaborated with analysts and BI teams to standardize reporting datasets, achieving a 30% improvement in query performance.

EDUCATION

Bachelor of Science, Boise State University | May 2014

Note: Willing to align with Microsoft data platform initiatives and Azure-based patterns as needed, leveraging existing Databricks-based lakehouse experience to accelerate adoption of Microsoft data solutions and integration patterns.