

Divyakumar Patel

Chapel Hill, NC 27514 | +1(984) 383-1528 | pateldivyakumar.77@gmail.com | [LinkedIn](#)

PROFESSIONAL PROFILE:

- **Data Engineer** with 7 years of experience designing, developing, and optimizing large-scale data solutions by integrating **cloud-based platforms, big data technologies**, and real-time streaming systems.
- Proficient in **Python, SQL**, and **Scala**, with hands-on expertise in building **ETL pipelines, data warehouses**, and real-time data processing solutions using **Apache Spark, Kafka**, and **Snowflake**.
- Experienced in **AWS (S3, Redshift, Glue, EMR, Lambda, Athena)** and **Azure (Databricks, Synapse, Data Factory)** for scalable cloud data solutions.
- Strong knowledge of data modeling, data governance, and performance tuning for structured and semi-structured data using **PostgreSQL, MongoDB**, and **Cassandra**.
- Expertise in streaming and batch data processing with **Apache Kafka, Spark Streaming**, and **Flink**, ensuring low-latency, real-time analytics solutions.
- Skilled in **DevOps** for data pipelines, utilizing **Docker, Kubernetes, Terraform**, and **CI/CD** tools like **Jenkins** and **GitHub Actions** for scalable deployments.
- Adept at designing **data lake** and **data warehouse architectures**, implementing data pipelines that support **BI tools** like **Tableau, Power BI**, and **Looker**.
- Experience in orchestrating workflows using **Apache Airflow, Luigi**, and **Prefect**, ensuring reliable and scheduled execution of complex **data workflows**.
- Strong background in **data security** and compliance (**GDPR, HIPAA**), implementing role-based access control (**RBAC**), **encryption**, and audit logging for sensitive data.
- Effective in **Agile environments**, actively collaborating in sprint planning, data strategy discussions, and cross-functional teams to deliver high-quality, scalable solutions.

TECHNICAL SKILL:

Programming:	Python, SQL, Scala, Java.
Data Processing:	Apache Spark, PySpark, Spark SQL, Kafka Streams, Apache Flink.
Databases:	PostgreSQL, Snowflake, Redshift, MongoDB, Cassandra, MySQL.
Cloud & DevOps:	AWS (S3, Redshift, Glue, EMR, Lambda, Athena), Azure (Databricks, Synapse), Docker, Kubernetes, Terraform.
ETL & Orchestration:	Apache Airflow, AWS Glue, Azure Data Factory, dbt, Luigi.
Big Data Technologies:	Hadoop (HDFS, Hive, Pig), Spark, Presto, Druid.
BI & Analytics:	Tableau, Power BI, Looker, Superset.
Testing & Automation:	PyTest, Great Expectations, DataDiff, dbt tests.
Version Control & CI/CD:	Git, Jenkins, GitHub Actions, ArgoCD.

PROFESSIONAL EXPERIENCE:

Company: Deutsche Bank, Cary, NC
Role: Senior Data Engineer

Sept 2022 – Present

- Developed ETL pipelines using Apache Airflow and AWS Glue, ensuring efficient ingestion and transformation of high-volume financial transactions.
- Engineered a real-time fraud detection system using Kafka Streams and Apache Flink, reducing fraud detection time by 40% and improving transaction security.
- Migrated PostgreSQL-based data warehouses to AWS Redshift, improving query performance by 35% while reducing storage and maintenance costs.
- Designed a scalable Data Lake on AWS S3, integrating Delta Lake and Presto, enabling seamless data retrieval for business analytics.
- Optimized Spark jobs, reducing batch processing times by 50%, leveraging partitioning, bucketing, and Parquet file formats.
- Integrated Apache Druid for real-time dashboards, providing instant visibility into financial transaction trends and fraud detection insights.
- Implemented RBAC and encryption for data security compliance (GDPR, PCI DSS), securing PII-sensitive datasets with AWS KMS.
- Automated Terraform scripts for cloud infrastructure provisioning, ensuring consistent deployment across AWS environments.
- Enhanced data pipeline monitoring using Datadog and Prometheus, proactively identifying latency bottlenecks and system failure.

Company: Truist Bank, Charlotte, NC

Jul 2020 – Aug 2022

Role: Data Engineer

- Designed and implemented data pipelines for a real-time credit scoring system, integrating Apache Spark ML for predictive analytics.
- Developed batch data pipelines using AWS Glue and Apache Flink to enhance financial transaction processing efficiency.
- Built a high-availability payment fraud detection system using Kafka Streams and Cassandra, ensuring low-latency transaction analysis.
- Integrated Airflow DAGs for ETL orchestration, improving workflow automation and reducing manual interventions by 60%.
- Migrated on-premise Hadoop workloads to Azure Databricks, reducing compute costs and improving processing efficiency.
- Deployed streaming analytics dashboards using Power BI and Looker, enabling real-time visibility into loan application statuses.
- Designed and implemented RESTful APIs for real-time transactional data retrieval, improving backend performance and scalability.
- Implemented Apache Iceberg for time-travel queries, enhancing auditability and compliance reporting.
- Configured Terraform-based infrastructure to automate deployment of data lakes and processing clusters across AWS environments.

Company: Plaid, San Francisco, California

Apr 2018 – Jun 2020

Role: Data Engineer

- Developed scalable ETL workflows using AWS Glue, Apache Airflow, and dbt, ensuring efficient data ingestion and transformation from multiple investment data sources.
- Built a predictive analytics platform using GraphQL and Snowflake, improving data retrieval performance by 30% for investment forecasting models.
- Designed batch and real-time data processing workflows using PySpark and Kafka, reducing latency in trade settlement reporting by 45%.
- Integrated Snowflake with Tableau and Looker, allowing real-time insights for portfolio risk assessment and performance tracking.

- Optimized queries in PostgreSQL and Redshift, reducing complex analytics query execution times from 20 minutes to 45 seconds.
- Deployed containerized microservices using Docker and Kubernetes, improving data API performance and scalability.
- Implemented data validation automation using Great Expectations and PyTest, ensuring data integrity across ETL workflows.
- Established CI/CD pipelines for data pipelines using Jenkins and GitHub Actions, improving deployment automation and release efficiency.
- Integrated Apache Hudi for incremental data processing, minimizing redundant computations and improving storage utilization.

EDUCATION:

- Bachelors in Computer Science | S.P. University

March 2018

CERTIFICATIONS:

- AWS Certified Data Analytics
- Microsoft Certified: Azure Data Engineer Associate
- Snowflake SnowPro Core Certification
- Certified Apache Spark Developer