

Md Javid

Senior Data Engineer / Data Architect

Bellerose, New York • md21.javid@gmail.com • (952) 592-2944 • <https://www.linkedin.com/in/mohammad-javid-0a4005203/>

SUMMARY

Senior Data Engineer / Data Architect with over 10+ years of experience leading the design and implementation of enterprise data platforms and modern data architectures. Expertise in AWS, Azure, Snowflake, Databricks, and big data technologies, with a strong focus on data modeling, governance, and performance optimization. Proven track record of delivering scalable, high-quality data solutions, mentoring engineering teams, and driving data strategy aligned with business goals.

SKILLS

- **Data Engineering:** ETL/ELT, Batch & Real-Time Data Pipelines, Data Modeling (Star/Snowflake, SCDs), Data Warehousing, Data Lakes/Lakehouse, Data Integration, Data Quality, Metadata Management, Data Governance, Data Lineage
- **Programming:** Python (PySpark, Pandas, NumPy), SQL (Advanced – CTEs, Window Functions, Query Tuning), Scala, Java, Shell Scripting (Bash)
- **Big Data:** Apache Spark (Core, SQL, Streaming), Hadoop (HDFS, Hive, MapReduce), Kafka, Apache Flink, Delta Lake, Apache Iceberg, Apache Hudi
- **Cloud Platforms:** AWS (S3, Redshift, Glue, EMR, Lambda, Kinesis, Athena, Step Functions, CloudWatch), Azure (ADF, Synapse, ADLS, Databricks, Event Hubs), GCP (BigQuery, Dataflow, Pub/Sub, Composer, Dataproc)
- **Databases & Warehousing:** Snowflake (Snowpipe, Streams, Tasks), Amazon Redshift, Google BigQuery, Azure Synapse, PostgreSQL, MySQL, SQL Server, Oracle, MongoDB, Cassandra, DynamoDB
- **Orchestration & Transformation:** Apache Airflow (DAGs), dbt (Core/Cloud), Azure Data Factory (ADF), AWS Glue Workflows, Prefect, Luigi
- **Streaming & Messaging:** Apache Kafka (Producers/Consumers, Kafka Streams), Spark Structured Streaming, AWS Kinesis (Streams/Firehose), Azure Event Hubs, Google Pub/Sub
- **DevOps & CI/CD:** Git (GitHub, GitLab, Bitbucket), Jenkins, GitHub Actions, Azure DevOps, Docker, Kubernetes, Terraform, CloudFormation
- **Monitoring & Logging:** AWS CloudWatch, Azure Monitor, GCP Operations Suite, ELK Stack (Elasticsearch, Logstash, Kibana), Prometheus, Grafana
- **BI & Analytics:** Power BI, Tableau, Looker, QuickSight, Data Visualization, Semantic Layer Design, Metrics Definition
- **Security & Governance:** RBAC, IAM (AWS/Azure/GCP), Data Encryption, GDPR/HIPAA Compliance, Data Catalogs (AWS Glue Catalog, Collibra, Alation)
- **Methodologies:** Agile/Scrum, SDLC, DataOps, CI/CD Pipelines, Performance Optimization, Cost Optimization, Stakeholder Management, Technical Leadership

EXPERIENCE

Senior Data Engineer / Data Architect | Huron

November 2022 - Present, New York, NY

- Designed and implemented scalable ETL/ELT pipelines using Python, PySpark, and SQL, processing over 5TB of structured and semi-structured data daily. Built reusable frameworks to standardize ingestion, transformation, and loading across multiple business domains.
- Architected a cloud-native data lakehouse solution on AWS (S3, Glue, EMR, Redshift, Lambda, Kinesis, Athena), enabling centralized data access and reducing data retrieval time by 40%.
- Developed and optimized Snowflake data warehouse, leveraging Snowpipe, Streams, and Tasks for continuous data ingestion and near real-time processing. Improved data availability and freshness for analytics teams.
- Implemented dbt (Data Build Tool) for transformation workflows, creating modular and version-controlled data models. Enhanced data testing, documentation, and lineage tracking across the pipeline.
- Orchestrated and monitored workflows using Apache Airflow, managing 200+ DAGs with dependencies, retries, SLA tracking, and alerting mechanisms to ensure high reliability.

- Designed enterprise-grade data models (Star/Snowflake schema, SCD Type 1 & 2) to support BI tools such as Tableau and Power BI, improving reporting performance and scalability.
- Built real-time streaming pipelines using Apache Kafka and Spark Structured Streaming, enabling event-driven data processing and reducing latency from hours to minutes.
- Improved query performance by 50% through advanced SQL optimization, partitioning, clustering, and indexing strategies in Snowflake and Redshift environments.
- Led migration of legacy on-prem Hadoop systems to AWS Lakehouse architecture using Delta Lake, ensuring minimal downtime and improved scalability.
- Established CI/CD pipelines using Jenkins, GitHub Actions, Docker, and Terraform, automating deployments and infrastructure provisioning while enforcing best practices.

Lead Data Engineer | Fractal

April 2020 - November 2022, New York, NY

- Architected and implemented enterprise-scale data solutions using Azure (ADF, Synapse Analytics, ADLS Gen2, Databricks), enabling centralized data processing for multiple business units. Designed scalable architectures to support both batch and near real-time workloads.
- Developed high-performance distributed data pipelines using PySpark on Azure Databricks, processing terabytes of structured, semi-structured, and streaming data. Applied optimization techniques such as partitioning, caching, and broadcast joins.
- Built and maintained robust ELT workflows in Azure Data Factory, integrating data from APIs, SQL Server, flat files, and external systems. Implemented parameterized pipelines and reusable components to reduce development effort.
- Designed and optimized enterprise data warehouse in Azure Synapse Analytics, implementing star and snowflake schemas. Improved query performance by 35% through indexing, distribution strategies, and materialized views.
- Implemented real-time ingestion pipelines using Azure Event Hubs and Apache Kafka, enabling event-driven data processing and reducing latency from hours to minutes.
- Developed transformation frameworks using dbt and Spark SQL, ensuring modular, testable, and version-controlled data models aligned with business logic.
- Built reusable data quality, validation, and auditing frameworks, incorporating logging and alerting mechanisms to proactively identify and resolve data issues.
- Optimized big data workloads in Databricks and Synapse by tuning Spark configurations, managing cluster resources, and implementing efficient file formats like Parquet and Delta Lake.
- Implemented CI/CD pipelines using Azure DevOps, Git, and Docker, automating build, test, and deployment processes for data pipelines and infrastructure.
- Collaborated closely with business stakeholders, analysts, and architects to define KPIs, data models, and reporting layers, ensuring alignment between technical solutions and business objectives.
- Led and mentored a team of data engineers, conducting code reviews, enforcing best practices, and driving adoption of modern data engineering standards and tools.

Senior Data Engineer | Itransition

July 2015 - April 2020, New York, NY

- Developed and maintained scalable ETL pipelines using Python, SQL, and Informatica, integrating data from multiple enterprise systems into centralized data warehouses. Ensured data consistency and reliability across various business domains.
- Worked extensively with the Hadoop ecosystem (HDFS, Hive, MapReduce) to process and analyze large datasets. Designed Hive tables and optimized queries for improved performance.
- Built efficient data ingestion pipelines using Apache Sqoop and Kafka, enabling seamless movement of data between relational databases and distributed systems.
- Designed and implemented dimensional data models (Star and Snowflake schemas) to support reporting and analytics. Collaborated with business teams to translate requirements into scalable data structures.
- Developed complex SQL queries, stored procedures, and views in Oracle and SQL Server, supporting reporting and downstream data consumption needs.
- Supported and enhanced enterprise data warehouse environments (Teradata/Oracle) by performing performance tuning, indexing, and query optimization.
- Created interactive dashboards and reports using Tableau and Power BI, providing actionable insights and improving decision-making capabilities.
- Implemented data quality checks, validation rules, and reconciliation processes, ensuring accuracy, completeness, and integrity of data across systems.

- Automated data workflows using Shell scripting and scheduling tools (Autosys, Control-M), reducing manual intervention and improving operational efficiency.
 - Monitored and troubleshoot production data pipelines, resolving failures and performance bottlenecks to ensure high availability and reliability.
 - Collaborated within Agile/Scrum teams, participating in sprint planning, daily stand-ups, and code reviews to deliver high-quality data solutions on time.
-

EDUCATION

Bachelor's Degree in Computer Science

UCLA – Los Angeles, CA, United States, June 2015