Michel Dorta

Senior Data Engineer

michel.dorta1@gmail.com

(727) 605-5734

NAPLES, FL 34116

michel-dorta-526491140

PROFILE

Senior Data Engineer with 9+ years building scalable data platforms and analytics products. Specialized in event-driven pipelines, streaming analytics, and customer-facing data products using Databricks, Snowflake, Spark, dbt, and Iceberg. Proven track record improving data freshness, observability & reliability by applying strong DataOps practices, partnering with cross-functional teams, and mentoring engineers.

SKILLS

Languages: Python, SQL, Scala, Java, TypeScript, Ruby, Bash

Streaming & Event-Driven: Apache Kafka, Apache Flink, Apache Spark, AWS Kinesis, Spark Structured Streaming, Apache Beam, Google Pub/Sub

Warehouses: Databricks, Snowflake, BigQuery, Amazon Redshift, Azure Synapse

Ingestion & CDC: Airbyte, Fivetran, Debezium, Kafka Connect, Apache NiFi

Orchestration & DataOps: Apache Airflow, Dagster, Prefect, dbt, Docker Kubernetes, Jenkins, Terraform

Data Quality & Governance: Great Expectations, Soda, OpenLineage, DataHub, Unity Catalog, data contracts, data lineage, PII handling

BI & Visualization: Tableau, Looker, Power BI, Apache Superset, Metabase

Storage & Table Management: Amazon S3, ADLS, Google Cloud Storage, HDFS, Parquet, Avro, Apache Iceberg

EDUCATION

Master of Science degree in Computer Science, *University of Florida*
2014 – 2016

PROFESSIONAL EXPERIENCE

Plaid, *Senior Data Engineer*
06/2022 – Present

- Led Plaid's evolution into a real-time data intelligence platform, delivering scalable financial insights for thousands of apps by building customer-facing data products.
- Improved transaction freshness for developers, redesigning core update workflows using event-driven pipelines with Go, Kafka, and Spark Streaming, reducing data latency by 10%.
- Enabled credit and risk decisioning products, producing high-quality financial signals for scoring and trust use cases through dbt, Snowflake, and Databricks, increasing feature freshness 2x.
- Supported real-time fraud detection capabilities, powering low-latency data flows with Kafka and Spark that allowed downstream systems to block risky transactions before completion.
- Accelerated investment data experiences, optimizing sync and processing pipelines using Spark and Snowflake, reducing account refresh times and improving end-user confidence significantly.
- Increased reliability of customer-facing datasets, implementing observability, SLAs, and data quality checks via Airflow and internal tooling.
- Collaborated cross-functionally with product, platform, and business stakeholders to translate customer needs into scalable data solutions, while mentoring and influencing roadmap decisions across multiple pods.

Walmart, *Senior Data Engineer*
07/2019 – 06/2022

- Led enterprise real-time data platform transformation across Triplet Cloud Model (public, private, edge) processing 10 PB/day for 10,000 stores.
- Architected Real-Time Inventory & Replenishment platform using Kafka, Spark Streaming, Cassandra reducing inventory signal latency.
- Engineered event-driven ingestion for Black Friday–scale traffic via Kafka, custom Messaging Proxy Service handling trillions of daily events with nearly zero data loss.
- Modernized Data Café analytics platform by migrating batch workloads to BigQuery, ADLS, Hive enabling analysis over 40+ PB of transactional data.
- Built Project Luminate supplier analytics pipelines under Walmart Data Ventures using Spark, Airflow, BigQuery delivering near–real-time insights to thousands of suppliers.
- Implemented multi-cloud orchestration and DataOps standards with Airflow, Azure Data Factory, IaC improving pipeline reliability and cutting failures by 15%.
- Enabled edge-based in-store analytics for associate tools by deploying localized processing synced to cloud lakes reducing inventory lookup latency.
- Partnered with product, infra, and ML teams to drive standardization and Agile delivery supporting 100,000 network devices with scalable, reusable pipelines.

Root Inc., *Data Engineer*
07/2016 – 07/2019

- Supported telematics data pipeline by using AWS, Ruby on Rails, PostgreSQL, enabling weekly pricing/analytics datasets from mobile test-drive data for thousands of trips.
- Built ingestion and validation jobs by using Python, Redshift SQL, S3, CloudFormation, Redis, processing 200 sensor variables per driver during 2–3 week test-drive windows.
- Implemented feature datasets for risk scoring by using GPS/accelerometer transforms, delivering clean behavioral signals (hard braking, phone handling, safe hours) used in model iterations.
- Standardized state-specific ETL outputs by using schema versioning and automated exports, supporting multi-state rollouts with consistent reporting and baseline auditability across dozens of states.