

Sarat Alman Senior Data Engineer | Multi-Cloud Specialist

✉ saratalman77@gmail.com ☎ (774) 436-7597 📍 Plano, TX 75252 🔗 serat_alman_website

PROFESSIONAL SUMMARY

Results-driven Data Platform Architect with 10+ years of experience designing scalable, multi-cloud data ecosystems across AWS, Azure, and GCP. Specialized in high-throughput streaming, complex ETL/ELT pipelines, and modern Data Lakehouse architectures. Adept at integrating Data Observability and GenAI workflows to drive business value. Expert in spearheading end-to-end digital transformations by bridging the gap between legacy systems and cloud-native solutions to ensure seamless scalability. Committed to fostering a culture of technical excellence through strategic roadmap development and the implementation of robust DataOps and MLOps practices that accelerate time-to-market for enterprise-grade products. Proven track record of optimizing cloud infrastructure, ensuring regulatory governance, and mentoring engineering teams to deliver cost-effective, real-time analytics solutions that align directly with enterprise goals

SKILLS AND TECHNOLOGIES

PROGRAMMING & SCRIPTING:

Python, Advanced SQL, Scala, Java, Golang (Go), C# / .NET, Bash / Shell Scripting.

CLOUD WAREHOUSING & DATA MODELING:

Snowflake, Amazon Redshift, Google BigQuery, Azure Synapse, Teradata, SQL Server, Oracle, NoSQL, Apache Iceberg, Dimensional Modeling (Star Schema), Data Vault 2.0, Data Mesh, Data Fabric, Semantic Modeling, Data Contracts, Reverse ETL, Parquet, ORC.

BIG DATA, STREAMING & AI PIPELINES:

Apache Spark (Batch & Streaming), Apache Kafka, Flink, Delta Lake, Presto, Beam, Pulsar, Hadoop, Hive, HBase.

CLOUD ECOSYSTEMS (AWS/AZURE/GCP):

S3, Glue, Lambda, Redshift, Kinesis, Lake Formation | ADF, Synapse, ADLS, Event Hub, Azure SQL | BigQuery, Dataflow, Dataproc, Pub/Sub, Vertex AI.

DOMAIN EXPERTISE:

Healthcare (FHIR, HL7, Claims, EHR, Eligibility, LOINC, CPT, ICD-10), Financial Services (Risk Management, Forecasting, Financial Modeling, Variance Analysis), Insurance, Retail & E-commerce, Manufacturing, Media & Entertainment.

DATA PIPELINES & ORCHESTRATION:

ETL/ELT, dbt, Apache Airflow, Azure Data Factory (ADF), Fivetran, Prefect, Dagster, Apache NiFi, Matillion, Talend, Informatica, SSIS, Data Migration, Real-Time Processing, Pipeline Optimization.

GOVERNANCE, SECURITY & OBSERVABILITY:

HIPAA, GDPR, CCPA, PHI Handling, IAM, RBAC, ABAC, Encryption, Data Masking, OpenLineage, Databricks Unity Catalog, Alation, Collibra, AWS Glue Catalog, Microsoft Purview, Metadata Management, Responsible AI, Audit Logging, Great Expectations, Monte Carlo.

DEVOPS, INFRASTRUCTURE & CI/CD:

Docker, Kubernetes (K8s), Terraform, Helm, GitHub Actions, GitLab CI/CD, Azure DevOps, Jenkins, Ansible, GitOps.

BI, REPORTING & ANALYTICS:

Power BI, Tableau, Looker, Apache Superset, Mode Analytics, KPI Modeling, Interactive & Operational Dashboards.

ETL & DATA PIPELINE DEVELOPMENT:

ETL/ELT, dbt, Apache Airflow, Azure Data Factory (ADF), Fivetran, Matillion, Apache NiFi, Prefect, Dagster, Pipeline Optimization, Data Migration, Real-Time Data Processing.

PROFESSIONAL EXPERIENCE

Senior Data Engineer

04/2022 – Present

KPG99 INC

- Architected Scalable Data Pipelines & Lakehouse Migrations using dbt, Airflow, and Snowflake, successfully transitioning legacy Informatica/SSIS workflows into modern, cloud-native architectures for high-volume healthcare data (FHIR, HL7, Claims).
- Engineered High-Performance Ingestion & Streaming Frameworks using Go, Python, and Apache Kafka to automate the flow of EHR and IoT telemetry, ensuring decoupled, reliable data layers and real-time insights via Power BI/ Tableau.

- Championed DataOps & Infrastructure Excellence by implementing Terraform and GitHub Actions for CI/CD, effectively streamlining resource allocation to reduce monthly cloud spend and operational overhead.
- Solidified Enterprise Governance & Observability via Databricks Unity Catalog, Monte Carlo, and Great Expectations, automating HIPAA/PHI compliance tracking and monitoring 300+ batch jobs for consistent SLA adherence.
- Optimized Distributed Compute & SQL Performance for massive clinical datasets exceeding 500M rows, leveraging Spark to drastically reduce ETL execution windows and accelerate downstream analytics for the business.

Data Engineer

11/2018 – 03/2022

Clarify Health

- Modernized Enterprise Data Architecture by transitioning legacy SSIS and Informatica workflows into high-scale Cloud Lakehouses using dbt, Airflow, and Snowflake, specifically tailored for complex healthcare datasets like FHIR and HL7.
- Engineered High-Performance Ingestion Frameworks using Go, Python, and Apache Kafka to automate the flow of EHR and IoT telemetry, effectively eliminating manual intervention and delivering real-time operational insights via Power BI/Tableau.
- Championed Data Reliability & Performance by optimizing distributed compute and SQL execution for massive clinical datasets, which drastically reduced processing windows and accelerated downstream analytics for the business.
- Solidified Enterprise Governance & Security through Databricks Unity Catalog, where I automated HIPAA/PHI compliance and encryption protocols, significantly reducing data-related incidents across the platform.
- Cultivated a DataOps Culture by implementing Terraform, GitHub Actions, and Monte Carlo/Great Expectations, ensuring consistent SLA adherence while keeping a sharp focus on cloud cost-efficiency and observability.

Associate Data Engineer

03/2016 – 11/2018

Lumenalta

- Modernized Enterprise ETL & Cloud Workflows by migrating legacy Informatica, SSIS, and Talend processes to Snowflake and SQL Server, ensuring seamless data movement and optimized pipeline scheduling across multi-cloud environments.
- Engineered High-Performance Ingestion Frameworks using Go-based Microservices and Automation Scripts to streamline data transformations, significantly improving pipeline reliability for IoT and manufacturing telemetry data.
- Championed Data Integrity & Compliance by automating HIPAA/PHI tracking and implementing robust Data Quality Frameworks, ensuring secure handling of healthcare datasets and 100% readiness for regulatory audits.
- Optimized Platform Performance & Analytics by refining complex SQL/PL-SQL queries for high-volume datasets, ensuring consistent SLA adherence and delivering actionable insights via Power BI and Tableau dashboards.

PROJECTS

Enterprise Data Mesh & Governance Platform:

Architected a 5-domain Data Mesh (Snowflake, Databricks, dbt, AWS), automating lineage and RBAC for HIPAA-compliant onboarding.

GenAI-Driven Clinical Insights Engine:

Developed Real-time RAG via Databricks, LangChain, Pinecone, and Kafka for automated healthcare risk prediction.

Petabyte-Scale Delta Lakehouse Migration:

Led end-to-end Azure-native Lakehouse migration (Delta Lake, ADF, Spark, Airflow), significantly cutting infrastructure costs and optimizing ETL performance.

CERTIFICATES

- | | | |
|--------------------------------------|---|------------------------------------|
| • Databricks Certified Data Engineer | • Google Cloud Professional Data Engineer | • Azure Solutions Architect Expert |
|--------------------------------------|---|------------------------------------|

EDUCATION

Bachelor of Science in Computer Science