

David Hammond

Senior Data Engineer

4016 Lake Underhill Rd Apt H, Orlando, FL 32803

david.mobile.hammond@gmail.com | +1 239- 439 - 5888 | <https://www.linkedin.com/in/david-j-hammond>

SUMMARY

Senior Data Engineer with **10+ years of experience** designing, building, and optimizing scalable data platforms and pipelines across cloud and big-data ecosystems. Expert in batch and real-time data processing, streaming architectures, and cloud-native data solutions using Spark, Kafka, Airflow, Databricks, AWS, GCP, and Azure. Proven track record of enabling analytics, ML, and business intelligence through reliable, high-performance data systems. Strong collaborator who bridges engineering and business to deliver measurable impact.

SKILLS

Data Engineering & Pipelines: Apache Spark (PySpark, Spark SQL, Spark Streaming), Apache Beam, Apache Flink, Apache Airflow, Kafka, Kafka Connect, Kafka Schema Registry, AWS Glue, dbt, ETL/ELT, Batch and Real-Time Data Processing

Streaming & Event-Driven Systems: Kafka, Spark Streaming, Apache Flink, AWS Kinesis, Google Pub/Sub, Low-Latency and Event-Driven Architectures

Cloud & Data Platforms: AWS (S3, EC2, EMR, Glue, Redshift, Athena), GCP (BigQuery, Dataflow, Pub/Sub, Vertex AI), Azure (Databricks, Synapse, Data Factory)

Data Warehousing & Databases: Snowflake, Amazon Redshift, Google BigQuery, PostgreSQL, MySQL, Oracle, SQL Server, Cassandra, MongoDB, Redis, Hive, Presto

Data Modeling & Architecture: Dimensional Modeling (Star & Snowflake Schemas), Medallion Architecture (Bronze/Silver/Gold), Data Vault 2.0, OLTP/OLAP, Data Lineage, Metadata Management

Programming Languages: Python, SQL, Scala, Java, Shell Scripting

DevOps, MLOps & Reliability: Docker, Kubernetes, Terraform, CI/CD (GitHub Actions, Jenkins, GitLab CI), Data Quality (Great Expectations), Observability (ELK Stack, CloudWatch, Grafana), ML Feature Pipelines

Analytics & Visualization: Tableau, Power BI, Looker, Matplotlib, Seaborn Governance & Security: IAM, RBAC, Encryption, Data Masking, HIPAA, GDPR, PCI-DSS Compliance

PROFESSIONAL EXPERIENCE

WTE(Wichita Tribal Enterprises)
Senior Data Engineer

Houston, TX
Aug 2022 - Present

- Promoted from Mid-Senior role within 9 months based on consistent delivery of **high-impact, scalable data solutions** across cloud and hybrid environments.
- Architected and deployed **ETL/ELT pipelines** using **Apache Spark, Airflow, and Python**, processing **2+ TB/day** of structured and semi-structured data.
- Built **low-latency streaming pipelines** using **Kafka, Kafka Connect, and Apache Flink**, reducing executive dashboard latency by **70%**.
- Unified analytics across **Redshift, BigQuery, and Snowflake**, improving cross-team data trust and accessibility.
- Automated **Kubernetes infrastructure provisioning** with Terraform, reducing setup time by **80%** and improving deployment consistency.
- Implemented **50+ Airflow DAGs** with **Great Expectations**, improving SLA compliance and pipeline reliability.
- Integrated **ML-adjacent workflows** (feature generation, model inputs, versioning) using **DVC and Azure ML**, reducing handoff delays by **40%**.
- Designed **dimensional data models** (Star & Snowflake schemas), cutting query times by 35% and enabling self-service analytics.
- Built event-driven architectures with **Kinesis and Flink** for fraud detection and behavioral analytics.

- Established centralized **observability and monitoring** (ELK, CloudWatch, Grafana), reducing MTTR by **45%**.
- Optimized **Redshift and BigQuery workloads**, lowering costs by **30%** while improving performance.

Providence
Senior Data Engineer

Renton, WA
Nov 2017 - May 2022

- Partnered with PMs, analysts, and data scientists to deliver **scalable data platforms**, increasing analytics adoption by **20%**.
- Managed **CI/CD pipelines** (GitHub Actions, Jenkins) for data infrastructure, reducing deployment errors by **40%**.
- Led migration of **200+ legacy jobs** to **Apache Beam on GCP Dataflow**, improving performance **3×**.
- Implemented **data governance frameworks** including schema registries, lineage, and RBAC to ensure **HIPAA and GDPR compliance**.
- Built executive dashboards and KPI monitoring using **Looker and Power BI**, eliminating **50+ hours/month** of manual reporting.
- Mentored **6+ junior engineers**, improving ramp-up time by **30%**.
- Collaborated with compliance teams to implement **end-to-end auditability and lineage** across critical datasets.

Bank Of America
Software Engineer (Data Engineer)

Plant City, FL
Jun 2013 - Oct 2017

- Developed **real-time data pipelines** using **Kafka, Kinesis, and Spark Streaming**, processing **2M+ events/sec**.
- Built and managed a centralized **AWS data lake** (S3, Glue, Athena), reducing retrieval latency by **25%** and infra costs by **20%**.
- Optimized **Snowflake warehouses** using clustering and materialized views, improving query performance by **40%**.
- Implemented secure **RBAC and IAM-based access controls** for healthcare and financial datasets.
- Automated CI/CD for data pipelines, cutting manual errors by **50%** and release cycles by **60%**.
- Built distributed ETL frameworks with **Spark and Airflow**, processing **320+ GB/day**.
- Developed REST and GraphQL APIs to expose analytics data to internal platforms.
- Improved data quality and reliability through standardized validation and testing frameworks.
- Built scalable backend services and **data APIs** using Python (Django/DRF), PostgreSQL, MongoDB, and Elasticsearch.
- Introduced **CI/CD and Kubernetes**, reducing release cycles by **40%**.
- Designed data ingestion pipelines for external sources, increasing data coverage by **45%**.
- Improved system performance and reliability through tuning, testing, and automation.

EDUCATION

Bachelor's degree in computer science
University of Florida - Gainesville, FL
Sep 2009 - Apr 2013