ANOOP VAMSI MEDURI
Email: anoopvamsi98@gmail.com | Phone: 814-812-0163
LinkedIn: linkedin.com/in/anoop-vamsi-meduri-7913a4109
Location: Open to Relocation | Full-Time Opportunities

PROFESSIONAL SUMMARY

- Results-driven Data Engineer with 7 years of experience designing enterprise-scale data platforms, building cloud-native pipelines, and enabling real-time analytics across AWS, Azure, and GCP ecosystems.
- Expertise in ETL/ELT development, big data processing, and distributed computing using Spark, PySpark, Hadoop, and Airflow to process multi-terabyte structured and unstructured datasets.
- Strong experience implementing modern data warehousing solutions using Snowflake, Redshift, and BigQuery, enabling high-performance reporting and enterprise BI analytics.
- Proven ability to architect cloud data lakes, CI/CD automation frameworks, and infrastructure-as-code deployments using Terraform, Jenkins, and CloudFormation.
- Hands-on exposure to AI/ML pipeline integration using Scikit-learn, TensorFlow, and PyTorch to operationalize predictive analytics within production data ecosystems.
-

TECHNICAL SKILLS

Programming: Python, SQL, Scala, Java, Bash
Big Data: Apache Spark, PySpark, Hadoop, Hive, Kafka, HBase, Sqoop
ETL & Orchestration: Apache Airflow, AWS Glue, DBT, Informatica, NiFi
Cloud Platforms: AWS, Azure, Google Cloud Platform
Data Warehousing: Snowflake, Amazon Redshift, BigQuery, Azure Synapse
Databases: MySQL, PostgreSQL, MongoDB, Cassandra
DevOps: Jenkins, Git, Docker, Kubernetes, Terraform, CloudFormation
Streaming: Kafka, Spark Streaming, Kinesis
Visualization: Power BI, Tableau, AWS QuickSight

PROFESSIONAL EXPERIENCE:

T-Mobile — Data Engineer Bellevue, WA | Nov 2024 – Present
Project: Enterprise Cloud Data Modernization & Advanced Analytics Platform

- Led the end-to-end design and implementation of scalable enterprise ETL pipelines using AWS Glue and EMR, enabling ingestion and processing of multi-terabyte telecom and supply-chain datasets while improving data processing performance by forty percent.
- Architected a centralized AWS Data Lake integrating structured and semi-structured data from Oracle, on-prem systems, and external APIs into S3, enabling enterprise-wide analytics and self-service BI reporting.
- Migrated legacy on-prem Oracle warehouse workloads into Amazon Redshift, redesigning schema models, optimizing partition strategies, and reducing query latency while lowering infrastructure costs.
- Engineered high-performance PySpark transformation frameworks on EMR clusters to process clinical, manufacturing, and operational datasets exceeding multiple terabytes per batch cycle.
- Integrated machine learning inference pipelines using Azure Databricks and AWS services, operationalizing predictive analytics models for real-time and batch scoring across production systems.
- Implemented enterprise data governance and security frameworks using AWS IAM, encryption policies, and row-level access controls within Redshift environments.
- Automated CI/CD deployment pipelines for data workflows using Jenkins, Git, and Docker, reducing manual deployment errors by eighty percent and accelerating release cycles.
- Built executive dashboards using Power BI integrated with Redshift and Athena, enabling near real-time operational visibility for R&D and business stakeholders.
- Designed real-time ingestion frameworks leveraging Kafka, Spark Streaming, and HBase to support streaming analytics and event-driven processing use cases.
- Supported hybrid cloud data initiatives by orchestrating Airflow pipelines across AWS and GCP platforms for departmental analytics workloads.

Environment: AWS (Glue, EMR, S3, Redshift, Lambda), PySpark, Airflow, Jenkins, Power BI, Kafka, Docker

Tiger Analytics — Data Engineer India | May 2020 – Nov 2023

Project: Real-Time Customer & Business Intelligence Data Platform

- Designed and deployed scalable ELT pipelines using AWS Glue, Lambda, and Kinesis, enabling ingestion of high-velocity transactional and behavioral data streams from multiple enterprise applications.
- Architected cloud data warehouse ecosystems using Amazon Redshift and S3, implementing partitioning, compression, and indexing strategies to optimize analytics performance and cost efficiency.
- Developed infrastructure-as-code deployment frameworks using Terraform and AWS CloudFormation, standardizing environment provisioning and reducing manual configuration overhead.
- Built automated real-time processing pipelines using Kinesis streams and Lambda triggers to deliver near real-time data availability for reporting and downstream ML models.
- Implemented DBT transformation frameworks integrated with Airflow orchestration, enabling modular SQL modeling, version control, and automated testing of enterprise datasets.
- Optimized Snowflake data warehouse performance by forty percent through clustering strategies, warehouse sizing, and query tuning techniques.
- Developed custom Python ETL utilities for parsing, validating, and cleansing complex structured and semi-structured datasets.
- Deployed machine learning forecasting models using AWS SageMaker, embedding predictive insights into production analytics workflows.
- Built executive visualization dashboards in AWS QuickSight delivering operational KPIs and business performance insights.
- Leveraged Amazon EMR Hadoop clusters for large-scale distributed batch processing workloads.

Environment: AWS, Snowflake, DBT, Airflow, SageMaker, PySpark, Terraform

---

Reliance Industries — Data Engineer India | Jun 2018 – Apr 2020

Project: Enterprise Data Warehouse & Reporting Modernization

- Developed foundational big data and warehouse capabilities supporting enterprise finance, supply chain, and retail analytics.
- Built distributed Spark and Hive ETL pipelines ingesting ERP, CRM, and operational datasets into centralized Hadoop ecosystems.
- Engineered transformation frameworks using PySpark and SQL to standardize enterprise reporting datasets.
- Designed dimensional data models enabling scalable BI reporting and executive dashboards.
- Developed Tableau visualization layers integrating Presto and Hive analytical engines.
- Automated ETL validation and reconciliation pipelines using SSIS frameworks.
- Implemented containerized batch workloads using Docker and Kubernetes clusters to improve compute scalability.
- Established CI/CD automation pipelines via Jenkins and Git for data platform releases.

Environment: Spark, Hive, SQL Server, Tableau, Docker, Kubernetes, Jenkins