

Nagarjun Chidaravalli Shanthakumar

Data Engineer

CA | 901-231-7086 | nagarjunchidaravallis@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Engineer with 4+ years of experience building scalable data platforms and intelligent pipelines on **Azure** and **Databricks**. Proficient in **Python**, **SQL**, **PySpark**, and **Airflow**, with a strong track record in cloud Data Lakehouse design, cost-efficient **ETL**, and analytics enablement. Experienced in developing **RAG** pipelines using **LangChain** and **OpenAI** for **LLM-powered** question answering. Passionate about merging data engineering with AI, embedding document intelligence, and automating modern data stacks.

SKILLS

- Methodology:** SDLC, Agile, Waterfall
- Programming Languages:** C, C++, Python, R, SQL, PySpark, Scala, Java, HTML, CSS
- Big Data Ecosystem:** Hadoop, MapReduce, Hive, Apache Spark, HDFS, Data Lake, Data Warehouse
- Generative AI & LLM Ecosystem:** RAG, LLM, LangChain, Prompt Engineering, Vector DB, Hugging Face, OpenAI API, LlamaIndex
- Visualization and Reporting Tools:** Power BI, Tableau
- ETL Tools:** SSIS, Apache Kafka, Airflow, DBT
- Cloud Technologies and Infrastructure:** Azure, AWS, GCP
- Development and Operations Tools:** Git, GitLab, Jenkins (CI/CD), Docker, Kubernetes, Grafana, Splunk, Jira, GitHub
- Databases:** SQL Server, MySQL, PostgreSQL, MongoDB, RDBMS, Snowflake

EDUCATION

- Masters in Information Systems** May 2025
California State University, Long Beach
- Bachelors in Computer Science and Engineering** July 2021
JSS Science and Technology University, Mysuru

EXPERIENCE

- Morgan Stanley, CA | Data Engineer** Sep 2024 – Current
- Enabled faster and more accurate financial reporting by architecting and deploying **ETL** pipelines using **Apache Airflow** and **SSIS** to integrate disparate financial data sources, improving data availability by **35%** and reducing reporting delays.
 - Accelerated BI queries by designing a star schema in a centralized **Synapse data warehouse**, with optimized data modeling, reducing complex joins and cutting average query execution time by 25% across analytics workloads.
 - Empowered strategic decision-making by delivering executive-level **Tableau dashboards** with real-time insights into revenue, cost, and budget data improving scenario planning speed by 20%.
 - Enhanced financial data consolidation and compliance through scalable **Azure Data Factory** pipelines, ensuring seamless integration of cloud and on-premises **ERP**, billing, and ledger systems, minimizing manual reconciliations and audit risks.
 - Reduced computing costs by 25% and improved job reliability by mitigating data skew in **PySpark** jobs on **Databricks** through optimized partitioning and join strategies, enabling balanced **Spark** task execution on large-scale **financial data**.
 - Built an Azure-native **LangChain RAG** assistant using **GPT-4** on Azure **OpenAI** to enable natural language search across 2,000+ internal financial documents, reducing analyst research time from 40 to under 2 minutes with cited, verifiable results.
- Infinite Infolab, India | Data Engineer** Jan 2020 – July 2023
- Orchestrated parameterized **ETL** pipelines using **Azure Data Factory** to ingest **CSV**, **JSON**, and **REST API** data from diverse sources, enabling **SQL-based**, schema-aware, scalable integration into **Databricks** for downstream analytics and reporting.
 - Engineered a fault-tolerant, high-throughput real-time data ingestion framework using **Apache Kafka** integrated with **Azure Data Lake Storage Gen2**, enabling seamless and reliable streaming from heterogeneous legacy systems.
 - Automated data validation and quality checks using **Python** and **Delta Live Tables (DLT)** with declarative expectations, ensuring schema compliance, referential integrity, and reducing QA overhead by **45%** through seamless **CI/CD** integration.
 - Developed incremental **ETL** pipelines in **Python** using **CDC** and timestamp logic, reducing data processing load by **40%**, and integrated **Git-based CI/CD** for automated, version-controlled deployments.
 - Designed and deployed dynamic, drill-down **Power BI** dashboards with **DAX** measures and role-based security, enabling executive leadership to track operational KPIs in real time and reduce reporting latency.
 - Led cross-functional **Agile** data engineering sprints with **Jira** and **GitLab**, managing backlog grooming, sprint planning, and continuous integration of feature branches, resulting in consistent delivery of scalable data solutions aligned.

CERTIFICATIONS

Databricks Certified Data Engineer Associate | **AWS Certified Cloud Practitioner** | **Tableau Desktop Specialist**