

MEGHANA MEKA
+1 (329) 221-4914
meghanameka04@gmail.com | [Linkedin](#)

PROFESSIONAL SUMMARY

- Data Engineer with 6+ years of experience designing, building, and supporting cloud-based data pipelines across Azure and AWS.
 - Strong hands-on expertise developing production-grade batch and streaming pipelines using PySpark, Databricks, Snowflake, dbt, Airflow, and cloud-native orchestration tools.
 - Proven experience owning end-to-end data pipelines including ingestion, transformation, optimization, monitoring, incident resolution, and root cause analysis in high-availability environments.
 - Demonstrated ability to optimize performance, improve data reliability, and support multi-terabyte daily workloads through efficient Spark tuning, Delta Lake optimization, and scalable architecture design.
 - Adept at collaborating with US stakeholders, product owners, and analytics teams while actively participating in Agile ceremonies, release planning, and continuous improvement initiatives.
-

CORE TECHNICAL SKILLS

Programming & Querying	Python, SQL, PySpark, Scala, Java
Data Engineering & ETL	Apache Spark, Databricks, dbt, Informatica, AWS Glue, Azure Data Factory
Cloud Platforms	Microsoft Azure, AWS, Azure Data Lake Gen Two, Azure Synapse, Azure Event Hubs, AWS S Three, Redshift, IAM, Kinesis
Big Data & Streaming	Hadoop, Hive, Kafka, Spark Streaming, Delta Lake, Apache Airflow
DevOps & Orchestration	Azure DevOps, Jenkins, GitHub Actions, Docker, Kubernetes, ARM Templates
Analytics & BI	Power BI, Tableau, Amazon QuickSight, Snowflake SQL
Databases & Storage	Snowflake, Redshift, PostgreSQL, MySQL, Oracle, Cassandra, MongoDB
Testing & Data Quality	dbt Tests, Pytest, Schema Validation, Data Quality Frameworks
Monitoring & Observability	Azure Monitor, Log Analytics, CloudWatch
Collaboration & Agile Tools	Git, GitHub, GitLab, JIRA, Confluence

PROFESSIONAL EXPERIENCE

Meridian Tech – Oakbrook, IL

Senior Data Engineer | September 2024 – Present

- Design and operate Azure Data Factory pipelines ingesting structured, semi-structured, and streaming data from SQL systems, REST APIs, and Kafka into Azure Data Lake Gen Two with built-in retry logic, lineage tracking, and SLA monitoring.
- Engineer scalable PySpark transformation frameworks in Azure Databricks processing multi-terabyte daily datasets, applying optimized joins, partitioning strategies, and window functions to reduce end-to-end pipeline runtimes.
- Built and managed dbt transformation layers in Azure Synapse with incremental models and automated data quality testing.

- Optimized Delta Lake tables using compaction and Z-Ordering, improving query performance for Power BI and Snowflake reporting.
 - Develop metadata-driven ingestion frameworks using parameterized ADF pipelines and ARM templates, improving deployment consistency and reducing manual configuration across environments.
 - Own CI/CD pipelines in Azure DevOps integrating Databricks notebooks, dbt models, and ADF artifacts through YAML-based workflows supporting frequent, reliable production releases.
 - Design event-driven ingestion patterns using Azure Event Hubs and Azure Functions to enable near real-time data ingestion and operational alerting use cases.
 - Establish observability standards using Azure Monitor and Log Analytics, actively resolving pipeline failures, performing root cause analysis, and restoring data availability within defined SLAs.
 - Design dimensional data marts using star and snowflake schemas, enabling faster analytical queries and improved dashboard responsiveness for business stakeholders.
- Environment: Azure Data Factory, Azure Databricks, Azure Synapse, Azure Data Lake Gen2, Event Hubs, PySpark, dbt, Delta Lake, Power BI, Snowflake, Azure DevOps

Azad Engineering – Hyderabad, India

Data Engineer | November 2019 – July 2023

- Architect and support enterprise data lake solutions on AWS S3 using AWS Glue to ingest data from SAP systems, relational databases, APIs, Kafka, and Kinesis streams.
- Develop Spark-based ETL and ELT pipelines processing tens of millions of records daily, delivering curated datasets for manufacturing, supply chain, and operational analytics.
- Implement Snowflake ingestion pipelines using Snowpipe and S3 event notifications to enable near real-time data availability for business intelligence and reporting teams.
- Build Spark Structured Streaming applications supporting real-time monitoring dashboards and proactive issue detection for operational systems.
- Configure AWS Glue Crawlers and Data Catalog to automate schema discovery, metadata management, and governance across raw and curated data zones.
- Improve pipeline reliability through checkpointing, fault-tolerant Spark design, and automated recovery mechanisms, reducing recurring production issues.
- Automate build and deployment workflows using Jenkins and Ansible, integrating testing steps and versioned artifacts to support controlled releases.
- Design Redshift data marts using dimensional modeling techniques, improving analytical query performance and consistency of KPI reporting.
- Orchestrate multi-layer data workflows using Apache Airflow, enhancing dependency management, traceability, and operational visibility.

Environment: AWS S3, AWS Glue, Redshift, Snowflake, Spark, PySpark, Kafka, Kinesis, Airflow, Jenkins, Ansible, Tableau, QuickSight

Orotron – Hyderabad, India

Data Analyst | September 2017 – October 2019

- Develop SSIS and PL/SQL-based ETL processes migrating data from Oracle and SQL Server into centralized enterprise reporting data marts.
- Implement automated data validation and reconciliation frameworks ensuring accuracy and completeness across migration and reporting pipelines.
- Design incremental load strategies using SQL to reduce refresh windows and improve reporting stability.
- Build SSRS and Tableau dashboards supporting finance and operations teams with actionable, data-driven insights.
- Tune stored procedures, SQL queries, and indexes to improve database performance and resource utilization.

- Support on-premises to AWS RDS migration initiatives including data validation, cutover support, and post-migration verification.
- Collaborate closely with business stakeholders to translate reporting requirements into ETL mappings and analytical outputs.

Environment: SQL Server, Oracle, SSIS, PL/SQL, Tableau, SSRS, Python, AWS RDS

EDUCATION

Master of Science in Information Systems Management

Marist College – Poughkeepsie, New York

CERTIFICATIONS

- Google Cloud Certified Professional Data Engineer
- Microsoft Certified Azure AI Engineer Associate
- Microsoft Certified Fabric Data Engineer Associate