

PROFESSIONAL SUMMARY

- Experienced Data Engineer with 6+ years of hands-on experience designing, building, and supporting scalable data platforms, real-time ingestion pipelines, and cloud-native analytics solutions across Azure, AWS, and hybrid environments.
- Strong expertise in PySpark, Databricks, Azure Data Factory, dbt, Airflow, and Snowflake, with proven success handling high-volume structured and semi-structured data from IoT devices, SAP HANA, APIs, and legacy enterprise systems.
- Demonstrated ability to design lakehouse architectures, implement incremental and CDC-based processing, and optimize performance using partitioning, Z-Order clustering, adaptive query execution, and query tuning techniques.
- Hands-on contributor to CI/CD automation, data quality enforcement, schema governance, and secure cloud deployments, ensuring production reliability and compliance across regulated and enterprise environments.
- Actively leveraging Generative AI and Agentic AI frameworks such as LangChain, FAISS, and OpenAI APIs to enable natural language analytics, Text-to-SQL, and intelligent data access for business and engineering stakeholders.

TECHNICAL SKILLS

- **Programming & Scripting:** Python, SQL, Scala, Java, C++, Bash, YAML, Regex
- **Data Engineering & ETL:** PySpark, Spark (Scala), dbt, Informatica PowerCenter.
- **Cloud & Platforms:** **Azure** (Data Factory, Databricks, Synapse, Data Lake Gen2, Event Hubs, IoT Hub, AKS, Azure Monitor, Functions, Key Vault), **AWS** (S3, Redshift, Glue, IAM, KMS), **GCP**, **Snowflake**, Oracle, BigQuery
- **Big Data & Streaming:** Apache Hadoop, Airflow, Kafka, Hive, Flink, Pig, Delta Lake, Z-Order, AQE
- **Workflow Orchestration & DevOps:** Airflow, Azure DevOps, GitHub Actions, Jenkins, Docker, Kubernetes, CI/CD, ARM
- **Data Visualization & Analytics:** Power BI, Tableau, Snowflake SQL, Oracle SQL, SSIS
- **Machine Learning & AI:** LangChain, OpenAI API, LangSmith, FAISS, PyTorch, TensorFlow, Scikit-learn
- **Generative AI:** Agentic AI, LangChain, RAG, Transformers, Prompt Engineering, Fine-Tuning (PEFT, LoRA), GANs
- **APIs & Integrations:** REST APIs, SAP HANA, SFTP, Salesforce, Veeva Vault, SharePoint.
- **Tools & Observability:** Streamlit, Selenium, SQLAlchemy, Application Insights, Log Analytics, Azure Monitor
- **Databases & Storage:** PostgreSQL, SQLite, Oracle, Redshift, Snowflake, BigQuery, Delta Lake
- **Version Control & Testing:** Git, Pytest, dbt Tests, Unit Testing, Schema Validation
- **Project & Scientific Thinking:** Root Cause Analysis, Assumption Testing, Agile Execution, Documentation & Knowledge Transfer, Data Storytelling, Stakeholder Communication, Critical Thinking, Leadership, Independent Work Capability

PROFESSIONAL EXPERIENCE

CVS Health Care – Irving, TX

Senior Data Engineer | Aug 2024 – Present

- Designed and implemented Azure Data Factory pipelines to ingest high-velocity IoT telemetry from vehicle-mounted devices, on-prem SQL systems, and Azure Blob Storage into Azure Data Lake Gen2, using parameterized pipelines and event-based triggers.

- Developed distributed PySpark workloads on Azure Databricks to clean, enrich, and aggregate time-series IoT data, applying watermarking, window functions, partition pruning, and schema-drift handling for reliable near-real-time processing.
- Built dbt transformation layers integrated with Azure Synapse, leveraging incremental models, automated tests, and documentation to support trusted operational and diagnostic reporting.
- Designed analytics data marts using star-schema modeling, columnstore indexing, and materialized views, enabling sub-second dashboard performance for vehicle health and operations analytics.
- Automated CI/CD deployments using Azure DevOps YAML pipelines to promote Databricks notebooks, ADF pipelines, and dbt models across environments, with secrets managed via Azure Key Vault.
- Implemented AKS-based microservices to process telemetry streams from Azure IoT Hub and Event Hubs, integrating Dockerized Spark jobs and event-driven Azure Functions.
- Improved platform observability by configuring Azure Monitor, Log Analytics, and Application Insights, tracking ingestion latency, message loss, and SLA adherence across streaming systems.
- Optimized data access patterns by refactoring data models and queries, reducing analytical access time and improving downstream BI consumption through Snowflake and Synapse.
- Enabled Power BI dashboards by delivering curated datasets and semantic models for operational and executive reporting.

Environment: Azure Data Factory, Azure Databricks, PySpark, dbt, Azure Synapse, ADLS Gen2, Event Hubs, IoT Hub, AKS, Azure Functions, Power BI, Azure DevOps, Docker, CI/CD, Azure Monitor

Cognizant – Hyderabad, India

Data Engineer | Sep 2020 – Jun 2023

- Built multi-source ingestion pipelines using AWS Glue and Python, integrating data from S3, SAP HANA, FTP, and BigQuery, with schema evolution and deduplication strategies.
- Designed and maintained a multi-zone Amazon S3 data lake, enforcing IAM policies, encryption via KMS, and versioning to support secure enterprise analytics.
- Tuned Amazon Redshift performance using distribution styles, sort keys, and workload management, improving query efficiency for large analytical workloads.
- Led migration initiatives from AWS to Azure and Snowflake, consolidating fragmented ETL logic into standardized dbt models and gold-layer views.
- Re-engineered ingestion workflows using Azure Data Factory, loading structured and semi-structured data into Snowflake through configuration-driven pipelines.
- Developed scalable PySpark transformations on Azure Databricks, processing multi-terabyte datasets with schema validation and lineage tracking.
- Implemented CI/CD pipelines using GitHub Actions and Azure DevOps, integrating Pytest and dbt tests for controlled production releases.
- Supported advanced analytics by enabling Power BI datasets and assisting business teams with exploratory data analysis.

Environment: AWS Glue, S3, Redshift, Azure Data Factory, Azure Databricks, PySpark, Snowflake, dbt, GitHub Actions, Azure DevOps, Power BI, SAP HANA

Health Axis – India

Associate Data Engineer | April 2018– Aug 2020

- Developed batch data pipelines using Apache Spark (Scala) on HDFS to process large-scale transactional and sales datasets.
- Built Informatica PowerCenter workflows to extract and transform data from SAP and FTP sources into staging environments.

- Migrated legacy HiveQL transformations to Spark SQL, improving execution performance and scalability.
- Scheduled and maintained Sqoop jobs to ingest relational data into HDFS for downstream processing.
- Implemented streaming prototypes using Kafka and Azure Event Hubs, supporting near-real-time ingestion use cases.
- Created Pig scripts to parse historical XML data and convert it into analytics-ready formats.
- Assisted in deploying dbt models and Airflow DAGs to orchestrate data transformations in Snowflake.
- Supported Azure Data Factory pipelines and Azure Data Lake integrations with validation and monitoring logic.

Environment: Apache Spark, Scala, Hadoop, HDFS, Hive, Informatica PowerCenter, Sqoop, Kafka, Azure Event Hubs, dbt, Airflow, Azure Data Factory, Snowflake

EDUCATION

Masters in Information Technology Management
Campbellsville University, Kentucky, USA

CERTIFICATIONS

Microsoft Certified: Fabric Data Engineer Associate (2025)
Microsoft Certified: Azure Data Engineer Associate (2023)
Google Cloud Certified: Professional Data Engineer (2023)
Microsoft Certified: Azure AI Engineer Associate (2025)
Microsoft Certified: Azure Data Scientist Associate (2025)
IBM Certified Data Scientist (2025)