BRYAN SCHAEFER
Lead Data Engineer

Houston, Texas | bryancoleschaefer1@outlook.com | linkedin.com/in/bryan-cole-schaefer-9775503b4

Professional Summary

Lead Data Engineer with 8+ years architecting production-grade batch and streaming data platforms using PySpark, Apache Airflow, Kafka, Snowflake, and lakehouse architectures. At Flatiron Health, designed end-to-end pipelines processing 8+ TB of daily healthcare data, reducing latency by 45% and delivering feature-ready datasets for ML models. Expert in data modeling, Great Expectations observability, CI/CD, and data reliability. Open to Lead/Senior Data Engineer roles focused on scalable cloud-native platforms.

Experience

Flatiron Health
Lead Data Engineer

Sep 2021 – Present

- Designed and orchestrated 15+ production Airflow DAGs with PySpark on Snowflake, processing 8+ TB of daily healthcare data → reduced end-to-end latency by 45% and improved data freshness SLA from 24h to 2h.
- Built real-time Kafka streaming pipelines feeding ML models and analytics dashboards → enabled 30% faster experimentation for research teams and cut compute costs by $180K annually.
- Implemented Great Expectations + dbt data quality framework with observability dashboards → dropped production incidents by 65% and achieved 99.8% data reliability.
- Defined Medallion (bronze/silver/gold) lakehouse architecture and mentored 6 data engineers on CI/CD, GitOps, Docker, and modular pipeline design.
- Translated complex clinical requirements into scalable, analytics-ready datasets that accelerated ML experimentation and supported BI reporting for hundreds of internal users.

Flare
Senior Data Engineer

Dec 2018 – Aug 2021

- Engineered scalable ETL/ELT pipelines with Apache Spark and Airflow integrating 20+ operational systems into Snowflake → improved query performance 3x and reduced transformation time by 55%.
- Created standardized Medallion lakehouse architecture with dbt and dimensional modeling → increased data consistency across teams and delivered curated analytics-ready tables for dashboards and ML use cases.
- Introduced Kafka event-driven streaming for near-real-time analytics → achieved sub-5-minute data freshness for key business metrics.
- Optimized data warehouse and data lake environments → strengthened data quality with validation logic and monitoring, reducing pipeline failures by 60%.

Uber
Data Engineer

Feb 2018 – Oct 2018

- Built high-throughput Spark streaming pipelines processing 10M+ events per hour → optimized partitioning strategies and reduced job runtime by 50% for product analytics and ML training.
- Produced analytics-ready datasets and structured data models supporting operational dashboards and machine learning experimentation.
- Implemented data quality validation and anomaly detection → ensured 99.9% data availability and resolved production issues for downstream teams.

Education

Texas Tech University
Bachelor of Science in Computer Science, GPA 3.8/4.0
Lubbock, Texas | Aug 2013 – Dec 2016

Skills

Languages: Python, SQL (Advanced – Window Functions & Optimization), PySpark, Scala, Java, Bash
Pipelines & Orchestration: Apache Airflow (DAGs), dbt, ETL/ELT, Medallion Architecture, Lakehouse
Streaming & Big Data: Apache Kafka, Spark Structured Streaming, Apache Spark
Data Platforms: Snowflake, Databricks, Amazon Redshift, BigQuery, Delta Lake, Parquet
Cloud & DevOps: AWS, Azure, GCP, Docker, Git, CI/CD, Terraform
Data Quality & Reliability: Great Expectations, Observability, Data Governance (RBAC), Monitoring