

Ethan Jones

Senior Data Engineer | AWS | Spark | Real-Time Data Platforms

Remote

Cleveland, TX | ethan.j.jones1@outlook.com | +1 737 339 7107

SUMMARY

Senior Data Engineer with 9+ years designing and delivering enterprise-scale data platforms, robust ETL/ELT pipelines, and cloud data infrastructures. Proven track record in data lakehouse architectures, dimensional modeling (star schemas), data quality and monitoring, and analytics enablement for BI and ML. Strong leadership and cross-functional collaboration with analytics, BI, and data science teams; skilled at driving reliability, performance, and governance in large-scale data environments. Eager to apply data engineering expertise to Microsoft data solutions and multi-cloud data initiatives.

SKILLS

Programming: Python, SQL, Scala, Bash

Data Processing: Apache Spark, Airflow, Kafka, dbt

Data Warehousing: Redshift, Snowflake, PostgreSQL

Orchestration: Apache Airflow, Metadata-driven pipelines

Streaming: Kafka, Spark Streaming

Cloud Platforms: AWS (S3, Glue, Redshift, EMR, Lambda, EC2)

Infrastructure / DevOps: Docker, Terraform, Git, Jenkins

Data Quality / Monitoring: Data quality checks, dashboards, lineage and governance basics

Data Modeling: Dimensional modeling, star schemas, data marts

Security / Compliance: Data access controls, audit-friendly architectures

Collaboration: Cross-functional partnership with analytics, BI, and data science teams

EXPERIENCE

Insperity

Location: Remote / Austin, TX (as applicable)

Senior Data Engineer

March 2021 – Present

- Architected and deployed scalable data pipelines processing 5+ TB of structured and unstructured data daily using AWS Glue, Spark, and Airflow.
- Led the design of a cloud-based data lake and warehouse architecture on AWS (S3 + Redshift), enabling governed analytics and streamlined data access.
- Optimized ETL workflows and distributed processing jobs, reducing end-to-end data latency by 45%.
- Mentored a team of 4 junior data engineers and established best practices for code quality, testing, and documentation.
- Implemented real-time data ingestion pipelines using Kafka and Spark Streaming for near real-time analytics.
- Collaborated with data scientists and BI teams to deliver production-ready datasets for machine learning and executive reporting.
- Introduced infrastructure-as-code (Terraform) to automate deployment of data infrastructure and improve system reliability.
- Enhanced data reliability and monitoring through automated quality checks and lineage tracking.

Tyler Technologies

Location: Plano, TX

Data Engineer

June 2017 – February 2021

- Developed and maintained scalable ETL pipelines using Python, SQL, and Apache Spark for large-scale public sector datasets.
- Designed optimized dimensional models and schemas (star schemas) to support enterprise reporting and analytics.
- Migrated legacy on-premise data systems to AWS cloud infrastructure, improving scalability and cost efficiency.
- Built automated data quality checks and monitoring dashboards to ensure data reliability across pipelines.
- Partnered with software engineering and analytics teams to translate complex business requirements into technical data solutions.
- Improved database and query performance through indexing, partitioning, and query optimization techniques.

EDUCATION

Bachelor of Science in Computer Science

University of Houston, Houston, TX

Aug 2012 – May 2016

(Note: Dates and company names preserved as provided.)
