

Saketh Sudagani

Data Engineer

Arlington Heights, IL | sakethsudagani8@gmail.com | (424) 391-7437 | [LinkedIn](#)

SUMMARY

Results-driven **Data Engineer with 5+ years of experience** designing, building, and optimizing **scalable data pipelines** and **cloud-based data platforms** across **Google Cloud, Azure, and AWS**. Proven expertise in **Big Data processing (Apache Spark, Hadoop, PySpark)**, **ETL/ELT workflows**, **real-time streaming (Kafka, Pub/Sub)**, and **data modeling (SQL, dbt, Snowflake, BigQuery, DynamoDB)**. Skilled in **ML pipeline automation**, **CI/CD for data products**, and **infrastructure-as-code (Terraform)** with a track record of delivering **highly reliable, cost-efficient solutions** that support **petabyte-scale analytics** and **AI-driven applications**. Adept at collaborating in **Agile environments** to enable **real-time insights**, **predictive analytics**, and **data-driven decision-making** for enterprise clients.

PROFESSIONAL EXPERIENCE

Google | CA | Data Engineer

November 2023 – Present

- Architected and implemented **serverless data pipelines** using **Google Cloud Dataflow, Cloud Composer, and Pub/Sub** to process **terabytes of ML training data** daily, achieving **99.99% system reliability** and **40% cost reduction** through **auto-scaling optimization**.
- Developed **AI-driven ETL pipelines** in **Python, TensorFlow, and PySpark** to automate **feature engineering** for **recommendation models**, boosting model accuracy by **22%** and reducing data prep time by **40%**.
- Optimized **ML feature storage and retrieval** by redesigning **BigQuery SQL schemas** for **time-series data**, reducing query latency by **50%** and enabling **real-time model inference** for **10M+ daily users**.
- Implemented **CI/CD pipelines** for **data products** using **Cloud Build and Terraform**, enabling automated **testing and deployment** of **50+ data models weekly**, achieving zero production incidents **over 6 months**.
- Designed and maintained **dbt (data build tool)** workflows to **transform, clean, and model raw data** into **analytics-ready datasets** in **BigQuery**, improving data accuracy by **30%** and enabling self-service reporting across teams.
- Integrated **GenAI models (Gemini, Vertex AI)** to automate **pipeline documentation** and **data lineage tracking**, improving **metadata coverage** for **500+ datasets**, enhancing **data governance and compliance readiness (GDPR/CCPA)**.

Walmart | AR | Data Engineer

September 2022 – October 2023

- Engineered **cloud-native ETL pipelines** on **Azure Data Factory (ADF), Event Hubs, and Synapse Analytics**, leveraging **Terraform IaC** to automate ingestion of **8TB+ daily retail data**, reducing **infrastructure provisioning time** from **days to hours**.
- Developed **event-driven streaming pipelines** using **Apache Kafka and Apache Airflow**, processing **5M+ events/hour** via **PySpark on Databricks** with **99.9% uptime**, cutting latency by **60%** for real-time inventory updates.
- Optimized **complex SQL workloads (Snowflake & Spark SQL)** through **indexing, partitioning, and query restructuring**, reducing query execution time by **65%** and enabling **near real-time supply chain reporting**.
- Designed and maintained **scalable Big Data pipelines** on **Azure Databricks, Snowflake, and Apache Spark**, enabling **end-to-end data integration workflows** across multiple retail business domains.
- Contributed to **Agile transformation** by facilitating **Scrum ceremonies** (sprint planning, retrospectives, daily standups) for **data engineering teams**, improving delivery predictability and reducing sprint cycle times by **30%** through continuous process improvements.

Wipro | India | Data Engineer

August 2018 – June 2021

- Developed **predictive maintenance ML models** in **Python (Scikit-learn, TensorFlow)** and **AWS SageMaker**, analyzing **IoT sensor data** to detect equipment anomalies with **90% accuracy**, reducing unplanned downtime by **35%**.
- Built and optimized **Big Data processing pipelines** using **Hadoop ecosystem technologies (HDFS, MapReduce, Hive, Pig)** across **hybrid on-premise and AWS environments**, improving data processing efficiency by **30%**.
- Designed and implemented **scalable object storage architectures** using **Amazon S3**, leveraging **data tiering, object versioning, and lifecycle policies** to manage large-scale datasets, achieving **20% storage cost reduction**.
- Optimized data retrieval by designing high-performance **NoSQL models** in **DynamoDB**, improving query speed by **27%**.

TECHNICAL SKILLS

Programming & Scripting: Python (Pandas, NumPy, Scikit-learn), SQL (BigQuery, T-SQL, PL/SQL), Scala, Shell Scripting

Big Data & Distributed Computing: Apache Spark, Hadoop (HDFS, MapReduce, Hive, Pig), Amazon EMR, Kafka, Event-Driven Data Streaming

Databases & Storage: DynamoDB, MongoDB, Cosmos DB, PostgreSQL, MySQL, SQL Server, Snowflake, Azure Blob Storage, Delta Lake

Cloud Platforms & Data Warehousing: GCP (BigQuery, Dataflow, Vertex AI, Pub/Sub), AWS (EC2, S3, Redshift, Glue, SageMaker, Athena, CodePipeline), Azure (Data Factory, Synapse, Databricks, Cosmos DB)

Machine Learning & AI: TensorFlow, PyTorch, XGBoost, Vertex AI, BigQuery ML, NLP (spaCy, NLTK), LLM-based Pipelines, Feature Engineering, GenAI models (Gemini, Vertex AI)

Data Engineering & ETL/ELT Tools: dbt, Apache Airflow, AWS Glue, Azure Data Factory, Informatica, SSIS, CI/CD Pipelines for Data Workflows

DevOps & Infrastructure as Code: Docker, Kubernetes, Terraform, Jenkins, Azure DevOps, Git, Google Cloud Build

Data Visualization & Reporting: Power BI, Tableau, Amazon QuickSight, Matplotlib, Seaborn

EDUCATION

Master of Science in Computer Science | University of Central Missouri, Warrensburg, Missouri, USA

Bachelor of Science in Computer Science | Lovely Professional University, Jalandhar, Punjab, India