

Yashwanth Amgoth
+1 813-499-7718
yashwanth.amgoth09@gmail.com
[LinkedIn](#)
Data Engineer



PROFESSIONAL SUMMARY

- Senior Data Engineer with over five years of hands-on experience delivering enterprise-scale data engineering solutions for Fortune-level clients, owning end-to-end data pipelines across ingestion, transformation, streaming, and analytics layers in production environments.
- Proven expertise in designing, operating, and supporting highly available batch and real-time data platforms using AWS, Azure, and big data ecosystems, with strong ownership of production workloads, incident response, root cause analysis, and SLA-driven delivery.
- Demonstrated ability to process hundreds of millions of records monthly, optimize distributed systems for performance and reliability, and partner closely with business, analytics, and engineering stakeholders to enable governed, self-service analytics and operational reporting.
- Highly experienced in Spark-based processing, Kafka streaming, cloud-native ETL frameworks, and CI/CD automation, with a consulting mindset focused on scalability, cost efficiency, and continuous platform improvement.

CORE TECHNICAL SKILLS

Category	Technologies
Programming & Scripting	Python, SQL, NoSQL, Java, Scala (Spark), UNIX Shell
Big Data & Processing	Apache Spark, Spark SQL, Spark Streaming, Hadoop, Hive, Impala, HBase
Streaming & Messaging	Apache Kafka, Kafka Streams, Azure Event Hub, Azure Stream Analytics
Cloud Platforms	AWS, Microsoft Azure, Google Cloud Platform
AWS Services	Lambda, CloudWatch, IAM
Azure Services	Azure Data Factory v2, Azure Data Lake, Azure SQL Database, Azure Synapse, Azure Functions, Blob Storage, Cosmos DB
Data Warehousing	Snowflake, Amazon Redshift, BigQuery, SQL Server, PostgreSQL, MySQL
Analytics & BI	Power BI, Looker, Tableau
DevOps & CI/CD	Jenkins, GitLab, Docker
Databricks & AI	Databricks, PySpark, Databricks Generative AI
Monitoring & Testing	Grafana, TPC-DS, JMeter, YCSB
Data Modeling & Governance	Erwin Data Modeler, Apache Atlas
Version Control	Git
Operating Systems	Linux

EDUCATION & CERTIFICATION

Master of Science in Computer Science
Saint Leo University, Tampa, FL, USA

Aug 2022 – Aug 2024

CERTIFICATION

- Snowflake Hands-On Essentials: Data Warehousing Workshop – Snowflake
- Databricks Fundamentals – Databricks
- Databricks Generative AI – Databricks

PROFESSIONAL EXPERIENCE

FedEx | USA

Role: Sr Big Data Engineer

Own end-to-end development and production support of large-scale batch and streaming data pipelines processing enterprise logistics and operational datasets across distributed Spark and Hadoop clusters.

July 2023 - Current

- Design and implement high-performance ETL frameworks using Apache Spark, Hive, and Spring Batch to integrate data from SAP, SharePoint, and cloud data lakes, significantly improving data availability for downstream analytics.
- Tune Spark DAG execution, memory allocation, and shuffle behavior to stabilize long-running jobs, reduce execution time, and improve overall cluster reliability in production environments.
- Build real-time streaming pipelines using Spark Streaming and Kafka to process event-driven logistics data, enabling low-latency operational insights and automated downstream consumption.
- Lead incident triage, on-call support, and root cause analysis for production failures, resolving Spark contention, thread deadlocks, and JVM performance issues to restore service within SLA.
- Enhance HBase read and write performance through compaction tuning, Bloom filters, and block cache optimization, improving lookup efficiency for high-throughput workloads.
- Implement automated monitoring and alerting using AWS Lambda and CloudWatch to proactively detect job failures, performance degradation, and data anomalies.
- Standardize CI/CD pipelines using Jenkins and GitLab to automate Spark job deployments, reduce release errors, and improve deployment consistency across environments.
- Integrate Apache Atlas for metadata management and lineage tracking, strengthening governance, compliance visibility, and audit readiness across analytics platforms.

Environment: Cloudera Hadoop (HDFS, Hive, Impala, HBase), Apache Spark (Scala, PySpark, Spark SQL, Spark Streaming), Apache Kafka, AWS Lambda, Amazon CloudWatch, Power BI, Grafana, Jenkins, GitLab, Linux, SQL Server, PostgreSQL, MySQL

Cognizant | India

Sep 2021 – Aug 2022

Role: Data Engineer

- Design and deliver scalable AWS Glue and PySpark-based ETL pipelines processing over two hundred million financial transactions monthly for regulated financial data platforms.
- Build real-time streaming solutions using Apache Kafka and Apache Flink to support trade reconciliation and near-real-time risk analytics reporting.
- Develop and maintain CI/CD pipelines using Jenkins and Docker to automate ETL deployments, reduce operational risk, and accelerate release cycles.
- Create reusable Python and SQL-based ingestion frameworks to automate source-to-target mappings and significantly reduce manual development effort.
- Optimize SQL queries and data access patterns across Snowflake and relational databases to improve query performance and analyst productivity.
- Develop Looker dashboards connected to Snowflake to surface compliance anomalies and operational trends for business and leadership stakeholders.
- Implement Databricks PySpark jobs for high-volume table-to-table transformations, ensuring data quality, consistency, and performance optimization.
- Administer SQL Server environments including stored procedures, security roles, and access controls in support of enterprise data applications.

Environment: AWS Glue, PySpark, Apache Kafka, Apache Flink, Snowflake, Jenkins, Docker, SQL Server, Looker, Power BI, Python, Azure Data Factory, Databricks

Hindustan Computers Limited | India

Jun 2019 – Sep 2021

Role: Data Engineer

- Collaborate with business analysts and solution architects to translate functional requirements into scalable Azure-based data architectures.
- Design and implement Azure Data Factory pipelines and Databricks workflows to support both batch and near-real-time data processing use cases.
- Develop Spark Streaming pipelines to ingest and analyze time-sensitive datasets, enabling faster operational insights.
- Provision, tune, and optimize Databricks clusters to balance performance and cloud cost efficiency.
- Implement near-real-time processing using Azure Event Hub, Service Bus, and Stream Analytics to support enterprise integration workflows.
- Write and deploy U-SQL scripts for large-scale data transformations targeting Azure Data Lake Analytics and Azure SQL Data Warehouse.
- Integrate heterogeneous data sources including Informix, Sybase, SFTP, and relational systems into centralized analytics platforms.
- Support production data pipelines through monitoring, issue resolution, and continuous performance improvements.

Environment: Azure Data Factory v2, Azure Data Lake, Azure SQL DB, Azure Synapse, Databricks (PySpark, Spark Streaming), Azure Event Hub, Azure Stream Analytics, U-SQL, SQL Server, Power BI, Python, Linux, Erwin Data Modeler