

Kevin Gibbons

McDonough, GA | (678) 364-2228 | kevin.gibbons.dev@gmail.com

SUMMARY

Senior Data Engineer with 8+ years of hands-on experience architecting and implementing enterprise data solutions within the Microsoft ecosystem. Proven expertise in Azure Data Factory, Synapse Analytics, Databricks, and SQL Server, with deep knowledge of dimensional modeling and metadata-driven pipeline orchestration. Successfully led end-to-end data engineering projects from requirements gathering through deployment, delivering scalable data platforms that drive business outcomes and enable self-service analytics across organizations.

EDUCATION

University of Texas at Austin | Bachelor's Degree | Sep 2012 - May 2016

SKILLS

- Microsoft Data Platform: Azure Data Factory, Synapse Analytics, Azure SQL Database, Azure Databricks, Fabric, ADLS Gen2, Power Platform, SQL Server Integration Services
- Data Engineering Architecture: Dimensional modeling, Star schema design, ETL/ELT pipelines, Data lakehouse architecture, Metadata-driven orchestration, Data quality frameworks, CDC patterns
- Programming SQL: Advanced SQL, Python, PySpark, T-SQL, Stored procedures, Complex query optimization, JSON transformation, API integration patterns
- Cloud Platforms Services: Microsoft Azure, AWS EMR, Google Cloud Dataproc, OAuth 2 authentication, REST API development, Cloud storage optimization, Infrastructure as code
- Data Processing Analytics: Apache Spark, Kafka Streams, Real-time analytics, Statistical modeling, Machine learning pipelines, Data validation, Performance tuning
- DevOps Collaboration: Azure DevOps, CI/CD pipelines, Terraform, Docker, Technical documentation, Workshop facilitation, Mentorship, Project leadership

CERTIFICATIONS

- Google Professional Data Engineer
- AWS Data Analytics – Specialty
- Databricks Data Engineer Associate

EXPERIENCE

GDIT

Sep 2022 - Present

Senior Data Engineer Consultant

Remote

- Architect and implement enterprise-scale data engineering solutions using Azure Data Factory, Databricks, and Synapse Analytics to modernize legacy data platforms and enable real-time analytics capabilities
- Design and optimize complex dimensional models and star schemas in Snowflake and Azure SQL Database, improving query performance by 60% for downstream reporting teams

- Build metadata-driven ETL/ELT pipelines using Python and advanced SQL that automate data quality validation, transformation logic, and orchestration workflows across multiple data sources
- Develop secure API integration patterns using OAuth 2 authentication to ingest data from external sources, standardizing JSON structures into normalized SQL formats
- Lead technical workshops and design sessions with client stakeholders to define data architecture strategies, roadmaps, and governance frameworks aligned with business objectives
- Create automated data quality monitoring and notification systems using Azure Functions and Power Platform, ensuring data reliability and proactive issue resolution
- Implement comprehensive data lake architectures on Azure Data Lake Storage with proper partitioning, lifecycle policies, and security controls for scalable data management
- Mentor junior data engineers and establish best practices for code reviews, documentation standards, and deployment processes using Azure DevOps CI/CD pipelines
- Collaborate with business intelligence developers and analysts to deliver self-service analytics capabilities through Power BI and Azure Synapse serverless SQL pools
- Drive adoption of modern data lakehouse patterns combining Delta Lake and Azure Synapse, reducing data processing costs by 40% while improving data freshness and availability

Nordstrom

Jul 2019 – May 2022

Senior Data Engineer

Dallas, TX

- Designed and delivered scalable data engineering solutions using Azure technologies and Spark processing to handle high-velocity retail data streams and support analytics initiatives
- Built robust ELT workflows in Airflow orchestrating complex transformation logic for customer behavior analysis, enabling data-driven merchandising decisions across multiple business units
- Implemented secure data lake architecture on AWS with encrypted S3 storage, CDC ingestion via DMS, and automated metadata cataloging through Glue for unified data access
- Developed SFTP-to-cloud ingestion pipelines using AWS Transfer Family and Terraform infrastructure-as-code, processing partner datasets exceeding 100GB daily volumes
- Optimized Spark transformations and applied partitioning strategies that reduced compute costs by 35% while maintaining sub-hour processing SLAs for critical business reports
- Enabled self-service analytics capabilities by implementing Athena and Presto query engines with Metabase visualization layer for business analyst teams
- Established comprehensive data governance practices including lineage documentation, access controls, and compliance frameworks supporting audit requirements
- Automated deployment processes using GitLab CI/CD for seamless pipeline updates and infrastructure changes with zero-downtime releases
- Collaborated with cross-functional product teams to define data modeling standards and transformation specifications that improved consistency across analytics platforms

RTI International

Oct 2016 – Jun 2019

Data Analytics Engineer

Austin, TX

- Developed ETL pipelines using Azure Data Factory and SQL Server Integration Services to process operational datasets exceeding 100GB, supporting analytical reporting requirements
- Built data transformation workflows using Hive SQL and Python on HDFS clusters, implementing validation routines that enhanced data quality for compliance and audit teams

- Created automated Oracle data ingestion processes using stored procedures and T-SQL, reducing manual coordination efforts and improving data freshness for critical business processes
- Implemented statistical forecasting models using Python libraries (ARIMA, regression analysis) to support training program optimization and resource planning initiatives
- Designed Redis-based caching mechanisms and lookup services that improved data retrieval performance for 3D modeling applications and enhanced user experience
- Delivered prototype machine learning pipelines on Google Cloud Platform using Vertex AI, demonstrating cloud-based model deployment patterns for predictive analytics use cases
- Established data lineage documentation and business rule specifications supporting GDPR and ISO 27001 compliance assessments and risk mitigation efforts
- Built Power BI dashboards and operational metrics visualizations enabling stakeholders to monitor data quality anomalies and system performance in real-time