

SHASHANK BASAVABATTULA

DATA ENGINEER

USA | +1 (945)-289-6460 | shashank445520@gmail.com | [LinkedIn](#)

SUMMARY

Result-oriented Data Engineer with 6 years of experience designing and optimizing large-scale data pipelines, integrating real-time and batch data solutions, and enabling scalable analytics infrastructure. Proven expertise in cloud platforms (AWS, Azure), big data technologies (Spark, Kafka), and modern data warehousing (Snowflake, Redshift). Strong collaborator with cross-functional teams, driving data-driven strategies across domains like finance, healthcare, and e-commerce

TECHNICAL SKILLS

- Cloud Platforms:** AWS (S3, Glue, Redshift, Lambda), Azure (Data Factory, Data Lake, Databricks)
- Big Data Technologies:** Apache Spark, Kafka, Hive, Hadoop
- Data Warehousing:** Snowflake, Amazon Redshift, Google BigQuery
- ETL/ELT Tools:** Apache Airflow, AWS Glue, dbt, Talend
- Programming Languages:** Python, SQL, PySpark
- CI/CD & DevOps:** GitLab CI/CD, Jenkins, Terraform
- Other Tools:** Delta Lake, Docker, Tableau, Power BI
- Compliance:** HIPAA, GDPR, SOC 2

WORK EXPERIENCE

Wells Fargo, USA

Sr. Data Engineer

Jan 2025 – Present

- Streamlined ETL and ELT workflows leveraging Apache Airflow, dbt, and AWS Glue, resulting in a 30% decrease in data processing time across customer and transaction datasets.
- Re-engineered dbt models to improve modularity and efficiency, reducing execution time by 33% and enhancing the reliability of end-to-end financial reporting pipelines.
- Tuned AWS Glue jobs to optimize performance, delivering a 25% faster turnaround for real-time processing of financial and transactional data.
- Scaled high-volume data ingestion using Apache Kafka and PySpark, increasing throughput by 40% and enabling near real-time analytics for account activity and transaction monitoring.
- Architected real-time data pipelines to power low-latency analytics on financial transactions and market trends, enhancing risk assessment and operational agility.
- Maintained and optimized cloud-based data warehouses (Snowflake, Redshift, and BigQuery), improving query response times and achieving a 25% reduction in storage expenditure.
- Empowered finance and compliance teams with improved access to actionable insights on asset management, profitability, and regulatory metrics through enhanced reporting workflows.
- Partnered with data scientists to streamline deployment of ML models for predictive financial analytics, accelerating delivery of AI-driven insights by 35%.
- Strengthened fraud detection systems, credit risk assessments, and market trend analyses to support real-time financial decision-making and strategic planning.
- Automated infrastructure and data workflow tasks using Python, reducing manual effort and increasing reliability across mission-critical financial pipelines.

CVS Health, USA

Data Engineer

Aug 2024 – Dec 2024

- Deployed Delta Lake over Azure Data Lake and AWS S3 to support version-controlled data storage, reducing redundancy and cutting storage costs by 30%.
- Enhanced data availability and optimized query performance for petabyte-scale healthcare datasets, supporting faster decision-making in clinical environments.
- Maintained strict adherence to HIPAA and healthcare data privacy standards, ensuring secure handling of sensitive patient information throughout the data lifecycle.
- Automated end-to-end CI/CD workflows for data pipelines using GitLab, Jenkins, and Terraform, cutting deployment times in half and improving reliability.
- Accelerated the delivery of real-time analytics to healthcare providers, minimizing delays in patient care decisions and improving operational responsiveness.
- Applied advanced data modeling strategies in Snowflake and BigQuery, reducing query execution time by 35% and boosting system efficiency for clinical reporting.
- Streamlined electronic health records (EHR) processing and clinical data retrieval, enabling clinicians to access critical insights more rapidly.
- Engineered scalable ETL pipelines using Azure Data Factory, Databricks, and PySpark, capable of processing over 5TB of healthcare data daily.
- Improved pipeline throughput by 40%, enabling real-time analytics integration with Azure Data Lake for faster, data-driven medical interventions.
- Facilitated seamless collaboration between data engineering, clinical, and data science teams to ensure alignment of

infrastructure with hospital workflows and patient outcomes.

- Supported strategic healthcare decisions by delivering timely and actionable insights through efficient data pipeline management and robust cloud architecture.

MiraiCoders Technology, India

Jan 2020 – Jun 2023

Jr. Data Engineer

- Spearheaded SQL and PL/SQL performance optimization efforts, resulting in a 40% increase in processing speed for complex transactional workloads in large-scale enterprise systems.
- Partnered with product development teams to implement GraphQL APIs, enabling more agile and efficient data querying and improving response times by 35%.
- Designed and implemented real-time streaming architectures with Oracle GoldenGate and Apache Kafka, capable of handling over 2 million events per minute while reducing latency by 30%. This significantly enhanced real-time analytics and data availability for operational intelligence.
- Automated end-to-end data ingestion from diverse sources—including IoT sensors, relational databases, and third-party APIs—into Oracle Cloud Infrastructure (OCI) Object Storage, reducing manual intervention by 60%. This automation accelerated integration cycles and ensured consistent, high-quality data delivery for analytics and reporting.

EDUCATION

Master in Computer Science

Auburn University at Montgomery, USA

Bachelor of Electrical & Electronics

Bharath institute of engineering and Technology, India