

SUMMARY

Data Engineer with 4+ years of experience implementing **Python-, SQL-, and Java-based data solutions** across **AWS, Azure, and GCP** environments. Demonstrated ability to improve **data throughput, query efficiency, and pipeline reliability** through API-driven ingestion, optimized database design, and automated validation. Strong background in transforming raw application and research data into **analytics-ready datasets** supporting enterprise reporting and decision-making.

WORK EXPERIENCE

Unum, USA | Data Engineer

Aug 2024 – Present

- Engineered **Python- and SQL-based ETL pipelines** loading application data into **MySQL and MongoDB**, processing **1M+ records daily**.
- Orchestrated scheduled workflows using **Apache Airflow DAGs**, reducing manual intervention and pipeline failures by **30%**.
- Transformed large datasets using **Apache Spark (PySpark)**, shortening batch processing time by **40%**.
- Developed **REST and GraphQL services with Spring Boot (Java)** to enable near-real-time ingestion for reporting systems.
- Optimized **database queries and Spark jobs** through indexing, partitioning, and execution plan tuning, cutting response times by **35%**.
- Implemented **data quality validation** using SQL constraints and Python reconciliation checks, lowering production defects by **40%**.
- Delivered curated datasets powering **Tableau and Power BI dashboards** used by **50+ stakeholders**.
- Monitored pipeline health using **AWS CloudWatch and structured logging**, maintaining **99.5% availability**.

Georgia State University, USA | Research Data Analyst

Jan 2023 -May 2024

- Built **Python- and SQL-driven ETL workflows** consolidating academic, survey, and behavioral datasets totaling **500K+ records**.
- Applied **Apache Spark (PySpark)** for scalable transformations and aggregations on large research datasets.
- Scheduled recurring dataset refreshes using **Apache Airflow**, reducing manual data preparation effort by **35%**.
- Prepared **analytics-ready tables** for **Tableau and Power BI**, improving query performance by **20%**.
- Performed **SQL-based reconciliation and audit checks**, increasing dataset accuracy by **25%**.
- Supported research outcomes through **statistical analysis** including trend and cohort comparisons.
- Authored technical documentation covering **data lineage, metric definitions, and transformation logic**.

Adons Softech, INDIA | Data Engineer

Aug 2019 – July 2022

- Implemented **ETL pipelines using Python, SQL, and Java** to support operational and application analytics.
- Integrated **Apache Kafka** for event-driven ingestion, capturing transactional updates with low latency.
- Processed batch and streaming data using **Apache Spark**, improving data availability for reporting systems.
- Designed **PostgreSQL and MongoDB schemas** optimized for high-volume ingestion and scalability.
- Automated extraction workflows via **RPA tools**, reducing manual reporting effort by **50%**.
- Applied **JWT and OAuth security frameworks** to protect internal data services.
- Improved pipeline resilience through **retry logic, exception handling, and alerting**, achieving **98%+ success rates**.
- Supported containerized deployments using **Docker and Kubernetes**.

SKILLS

Data Engineering & Pipeline: Python- and SQL-based ETL pipelines, Apache Airflow DAG orchestration, Apache Spark (PySpark) transformations, Apache Kafka event-driven ingestion, data validation using SQL constraints and Python reconciliation logic

Programming & Querying: Python (Pandas, NumPy), SQL (joins, window functions, query optimization), Java (Spring Boot), data structures, stored procedures

Data Processing: Batch processing with PySpark, near-real-time streaming via Kafka, scheduled workflows using Airflow, error handling and retry mechanisms, data reconciliation

Data Warehousing & Storage: Dimensional modeling (fact and dimension tables), PostgreSQL (partitioning, indexing), MySQL (query tuning), Oracle, MongoDB, NoSQL databases

Cloud & Big Data Platforms: AWS S3, AWS EC2, AWS Glue, AWS Lambda, AWS CloudWatch, Azure Data Factory, Azure Data Lake, Google Cloud Platform

APIs & Integration REST APIs using Spring Boot, GraphQL data services, event-driven architecture for system integration

Monitoring & Reliability: Pipeline monitoring with AWS CloudWatch, application logging frameworks, SQL- and Python-based data quality checks, performance tuning

DevOps & Delivery: CI/CD pipelines using Jenkins and Git, containerization with Docker and Kubernetes, infrastructure provisioning using Terraform

Collaboration & Process: Stakeholder communication, technical documentation, Agile methodology, SDLC

EDUCATION

Master in Computer Science | Georgia State University, USA

May 2024

Bachelor's in Information Technology | BVRIT, India

June 2021