

CARL SCOTT GOLDING

Senior Data Platform & MLOps Engineer | Big Data & Machine Learning Pipeline Specialist

(713) 979-4215 • carl.glding88a4@gmail.com • www.linkedin.com/in/carl-golding-282ab03a7 • Jacksonville, FL

Summary

Data Platform and Cloud Data Engineer with 13+ years of experience designing, building, and operating enterprise data platforms that enable analytics, business intelligence, real-time insights, and production machine learning. Proven at translating business goals—such as customer personalization, financial accuracy, healthcare outcomes, and operational efficiency—into scalable, secure, and compliant data solutions. Deep expertise across cloud data warehouses and lakehouse architectures, ELT/ETL frameworks, streaming pipelines, and analytics engineering, with strong hands-on skills in Python, SQL, Spark, and modern orchestration tools. Known for modernizing legacy systems, standardizing data quality and governance, and implementing CI/CD, DataOps, and MLOps practices that accelerate delivery while improving reliability. A collaborative partner to business, analytics, and data science teams, consistently delivering trusted, analytics-ready data that drives informed decision-making and measurable business value.

Experience

Datalere

Denver, CO

Data Platform Ops Engineer

01/2023 - 01/2026

- Architected and operated enterprise-scale cloud data platforms using DataOps, DevOps, and MLOps best practices, supporting analytics, BI, real-time, and production ML workloads.
- Automated infrastructure provisioning using Terraform and containerized platforms with Docker and Kubernetes, ensuring secure, compliant, and highly available data environments.
- Integrated Fivetran-managed data pipelines with dbt and dbt Fusion to standardize downstream transformations, enforce data quality checks, and support analytics engineering workflows.
- Deployed and managed self-hosted Airbyte using Docker and Kubernetes, enabling customization, scalability, and tighter control over data movement.
- Designed and maintained CI/CD pipelines (Jenkins, GitHub Actions, Azure Pipelines, GitLab) for data engineering, analytics, and ML workflows, including Python, PySpark, SQL, and dbt deployments.
- Administered and optimized the Databricks Lakehouse platform (Apache Spark, Delta Lake), integrating batch and streaming pipelines with Snowflake, BigQuery, Redshift, and Athena.
- Implemented modern ELT and analytics engineering frameworks, including dbt and dbt Fusion, to standardize transformations, enforce data quality, and accelerate analytics delivery.
- Built and supported real-time and event-driven data pipelines using Apache Kafka, enabling low-latency ingestion and streaming analytics across distributed systems.
- Led SQL Server analytics and BI modernization, developing complex T-SQL queries, stored procedures, and data models, and migrating IBM Cognos objects to SSRS reports requiring advanced transformations and performance tuning.
- Enabled end-to-end MLOps pipelines, supporting feature engineering, model training, deployment, monitoring, and observability with Datadog, while serving as a senior escalation point for platform, streaming, and ML production issues.

Epsilon US

Machine Learning Data Engineer

Irving, TX

01/2017 - 12/2022

- Collaborated with business and product teams to gather requirements, translate marketing and AI use cases into technical solutions, and implement privacy-aware data engineering practices supporting customer 360, personalization, and analytics initiatives.
- Designed and maintained end-to-end data pipelines and feature stores for machine learning workflows, enabling AI-driven personalization, predictive modeling, and privacy-centric customer identity resolution across paid, owned, and earned marketing channels.
- Built and automated ETL/ELT frameworks using SQL Server, SSIS, T-SQL, PostgreSQL, BML Express, and SQL Server Agent, ingesting, transforming, and validating large-scale structured and semi-structured datasets for model training, batch inference, and analytics.
- Developed and maintained data warehouses and data marts across Marketing, Sales, Product, HR, and Customer Engagement domains, implementing dimensional modeling, slowly changing dimensions, and partitioning to support scalable analytics and ML feature engineering.
- Partnered with data scientists and ML engineers to operationalize models, ensuring clean, reproducible datasets, efficient feature pipelines, and reliable ML scoring outputs for audience segmentation, churn prediction, attribution modeling, and personalization campaigns.
- Implemented data quality, validation, and monitoring frameworks using SQL Server and PostgreSQL, detecting anomalies and ensuring feature accuracy, which directly improved the reliability of ML-driven marketing decisions.
- Built mission-critical ML scoring pipelines and data streams feeding outputs into marketing automation, email, loyalty, digital media, and personalization platforms, enabling real-time insights and actionable recommendations.
- Developed interactive dashboards and analytics reports using Power BI and SSRS, and migrated legacy Cognos reports, empowering marketing and executive stakeholders to explore customer behavior, campaign performance, and model outputs.
- Led reusable ETL framework development and database CI/CD processes using BML Express and RoundhouseE, standardizing deployment of schema changes, version control, and multi-environment consistency for data pipelines.
- Optimized large-scale SQL Server and PostgreSQL workloads through indexing, query tuning, incremental loads, and partitioning, improving ETL performance, data refresh times, and ML model training efficiency.
- Collaborated with business, product, and analytics teams to translate marketing and AI objectives into technical solutions, designing privacy-aware data architectures and ML pipelines that drove customer engagement, personalization, and actionable marketing insights.

Empower

Cloud Data Engineer

Greenwood Village, CO

09/2012 - 12/2016

- Architected and delivered enterprise-scale data engineering solutions supporting Empower's retirement, recordkeeping, and wealth management platforms, serving 19M+ participants and 88K+ organizations with high data accuracy and reliability.
- Designed and maintained robust ETL and ELT pipelines to ingest, cleanse, and transform large volumes of financial, transactional, and customer data from relational databases (Oracle, SQL Server, PostgreSQL) and external data providers.
- Played a key role in early cloud data platform modernization, leveraging AWS, Azure, and Google Cloud services to transition legacy on-prem systems toward scalable, cloud-first analytics architectures.
- Implemented cloud data warehouses and analytical layers using Snowflake and Google BigQuery, enabling faster insights for actuarial, compliance, and business intelligence teams.
- Developed optimized ELT workflows that pushed transformations closer to the data warehouse layer, improving query performance, reducing data latency, and lowering infrastructure costs.
- Built and optimized dimensional and analytical data models (fact and dimension tables) to support retirement plan analytics, contribution tracking, investment performance, and regulatory reporting.
- Applied strong business logic and data validation rules to ensure the accuracy of financial metrics such as balances, vesting, contributions, distributions, and employer match calculations.
- Collaborated with product managers, compliance teams, and business stakeholders to translate complex business requirements into trusted, analytics-ready datasets used for executive dashboards and customer-facing insights.
- Implemented data governance, security, and compliance best practices, including role-based access control (RBAC), encryption at rest and in transit, and audit logging to meet financial services and regulatory standards.
- Automated data workflows using Python and SQL, reducing manual processing, minimizing operational risk, and improving overall data pipeline reliability.

Experience

WelbeHealth

San Jose, CA

Pipeline-Centric Data Engineer

09/2010 - 08/2012

- Designed, developed, and maintained scalable ETL data pipelines supporting healthcare operations, clinical data, and care-coordination workflows.
- Built batch and near-real-time data ingestion pipelines to integrate data from EHR systems, claims feeds, care management platforms, pharmacy, and third-party healthcare vendors.
- Implemented data integration and transformation processes to consolidate clinical, operational, and financial datasets into centralized data warehouses and early data lake architectures.
- Ensured data quality, validation, and reconciliation across pipelines handling sensitive PHI / HIPAA-regulated healthcare data.
- Optimized pipeline performance, reliability, and data freshness through monitoring, logging, alerting, and failure recovery mechanisms.
- Collaborated with clinicians, operations, analytics, and care-coordination teams to translate healthcare business requirements into production-ready data pipelines.
- Supported healthcare analytics, reporting, and population health insights by delivering clean, well-modeled datasets for downstream consumers.
- Implemented data governance and metadata documentation to improve lineage, traceability, and pipeline maintainability.
- Assisted in standardizing ETL frameworks and data modeling practices to improve pipeline scalability and reusability.
- Contributed to compliance-aligned data handling practices supporting HIPAA, healthcare reporting, and audit requirements.

Education

University of North Florida

Jacksonville, United States

Bachelor's degree of Computer Science

2006 - 2010

Skills

Data Processing • Data Management • Cloud Platform • Big Data • Data Modeling • Data Warehousing • ETL/ELT • Data Pipeline • Data Streaming • Data Analytics • Data Lineage • Data Integration

Python • SQL • Scala • Java • PySpark • AWS • T-SQL • S3 • Redshift • Glue • Lambda • Kinesis • GCP • BigQuery • Dataflow • Pub/Sub • Azure • Synapse • Data Factory • Snowflake • PostgreSQL • Cassandra • MongoDB • Airflow • dbt • Fivetran • Informatica • Dataflow • Apache Spark • Hadoop • Dataproc • EMR • Spark Structured Streaming • Step Functions • Cloud Composer • Luigi • Terraform • CloudFormation • Docker • Kubernetes • Jenkins • GitHub Actions • Oracle • SQL Server • MySQL • DB2 • Sybase • DynamoDB • Amazon Athena

Hive • Databricks • Delta Lake • Talend • DataDog • CDP • Cognos • Teradata • SSIS • SSRS • Batch Scoring • RoundhouseE • Kafka • DataStage • Pentaho • Apache Nifi • T-SQL • PL/SQL • Shell • R • Flink • Star Schema • OLAP • OLTP • Data Marts • OLAP cubes • Git • Ansible • Tableau • Power BI • Qlik • Looker • MicroStrategy • Cognos • PCI DSS • SOX • GDPR • HIPAA • Data Encryption • RBAC • Pig • HBase • HiveQL