

ZANE SIDIQUE

Senior Data Engineer | DataPipelineArchitect | BigDataand Cloud Analytics

LosAngeles, CA • (323) 553-3802 • zaane.sidd@gmail.com • [linkedin.com/in/zane-sidiqu](https://www.linkedin.com/in/zane-sidiqu)

Senior Data Engineer with 8+ years of experiencedesigningand optimizing scalable data pipelines, implementing robust ETL and ELT processes, and managing large-scale cloud data infrastructures acrossAWS, Azure, and GCP. Expert in real-time streaming architectures, data warehouse design, and big data frameworks. Proven track records of improving data and reliability, reducing processing latency, and enabling advanced analytics and machine learning applications in compliance-driven environments.

TECHNICAL SKILLS

Data Warehousing and Lakes: Snowflake, Redshift, BigQuery, Delta Lake, Apache Iceberg, AWS S3, Azure Data Lake, dbt
ETL and ELT Pipelines: Apache Airflow, AWS Glue, Azure Data Factory, Talend, Alteryx Designer, Apache NiFi, Dataflow
Big Data and Streaming: Apache Spark (PySpark), Kafka, Hadoop, Hive, Apache Beam, Flink, Kudu, Impala, Kinesis
Programming and SQL: Python, PySpark, Scala, SQL (complex queries, query optimization), Bash, Node.js, Go (Golang)
Modeling and Governance: Dimensional modeling, schema design, Quality frameworks, data cataloging, lineage tracking
Machine Learning Libraries: TensorFlow, PyTorch, Scikit-learn, Keras, Pandas, DASK, SVM, Autoencoders, XGBoost
CI/CD and DevOps: Docker, Kubernetes, Git, GitHub, GitLab, Jenkins, CI/CD pipelines for data workflows, Helm, ArgoCD
Data Visualization: Power BI, Tableau, Google Data Studio, AWS QuickSight, Looker, Seaborn, Plotly, Apache Superset
Cloud Platforms: AWS(Glue, Redshift, Athena, EMR, S3, Lambda, SageMaker, EC2), GCP(BigQuery, Dataflow, CloudStorage, Azure(Synapse, Databricks, Data Factory)

WORK EXPERIENCE

Senior Data Engineer | [Deloitte](#) March2021–Present

- Architected real-time datapipelines using Apache Airflow, Spark (PySpark), and Kafka processing millions of daily events and improving system performance by 30%.
- Built cloud data warehousing on Snowflake and BigQuery using Alteryx and dbt, implementing automated ETL and ELT workflows that cut processing times by 40%.
- Designed dimensional data models and schema structures supporting predictive analytics, ML workloads, and reporting requirements for multiple cross-functional teams.
- Led migration of on-premise data systems to AWS and GCP, improving scalability, fault tolerance, and overall data availability for the entire organization.
- Established data quality frameworks and governance standards, implementing automated validation checks and lineage tracking across all pipelines.
- Optimized Spark jobs and SQL performance through partition pruning, caching, and cluster right-sizing, reducing compute costs across all production workloads by 25%.
- Built Tableau and Power BI dashboards enabling business stakeholders to self-serve real-time operational and analytical insights without engineering support.
- Collaborated with data science and ML teams to build feature engineering pipelines and automated model retraining workflows on AWS SageMaker and Google Vertex AI.
- Mentored junior data engineers on pipeline design, PySpark optimization, and cloud best practices, consistently improving Overall team delivery and productivity.
- Designed and managed Kafka-based streaming ingestion, handling high-throughput event streams from multiple upstream sources producers into a centralized data lake.

WORK EXPERIENCE (CONTINUED)

Data Engineer | [Verifone](#) February 2019 – March 2021

- Developed scalable ingestion pipelines from structured and unstructured sources using Apache NiFi, integrating with Cloudera, HDFS, and Impala for enterprise data platform consolidation.
- Leveraged Azure Data Factory and Databricks to automate complex ETL and ELT flows across data domains, reducing pipeline latency by 30% and eliminating manual processing.
- Built low-latency storage layers using Apache Kudu and Ozone for real-time analytics workloads serving multiple downstream business intelligence and ML consumers.
- Performed large-scale data processing in Amazon Redshift and S3, delivering high-performance data marts and optimized SQL queries, cutting report generation time by 45%.
- Implemented CI/CD pipelines for data workflow deployments using GitLab and Jenkins, enabling automated testing, code review gates, and version-controlled releases.
- Designed and enforced data governance policies including schema versioning, role-based access controls, and automated data quality monitoring using alerting frameworks.
- Migrated batch processing workloads to Apache Spark on EMR, improving throughput by 3x and enabling parallel processing of terabyte-scaled datasets for analytics teams.

Software Engineer | [Wells Fargo](#) April 2017 – January 2019

- Delivered API-based data segmentation and process automation solutions, increasing operational efficiency by 15% across payment data workflows serving millions of daily transactions.
- Built and maintained Node.js microservices for real-time data ingestion, integrating third-party REST and GraphQL APIs to enhance platform interoperability across distributed systems.
- Implemented data pipeline monitoring and alerting, increasing automated test coverage to 95% using Unit and Jest test frameworks to ensure application stability.
- Provisioned and managed AWS infrastructure (EC2, S3, Lambda) supporting data ingestion, transformation, and storage for high-availability payment processing systems.
- Collaborated with data science teams to design feature engineering pipelines, producing clean and structured datasets for machine learning model training and validation.
- Optimized relational database schemas and query performance across PostgreSQL and MySQL, reducing average query response time by 35% for critical reporting workflows.

CORE COMPETENCIES

- Real-time and batch data pipeline architecture across multi-cloud environments (AWS, Azure, GCP)
- Advanced SQL query optimization, complex joins, window functions, CTEs, and database performance tuning
- Data lake architecture design using Delta Lake, Apache Iceberg, and cloud-native object storage
- End-to-end machine learning pipeline support, including feature stores, data versioning, and model serving
- Data governance, data quality frameworks, metadata management, and regulatory compliance (HIPAA, GDPR, SOC 2)
- Stream processing using Apache Kafka, Flink, Kinesis, and Apache Beam for high-throughput event-driven systems
- Cross-functional stakeholder collaboration translating business requirements into scalable data solutions

EDUCATION

Bachelor of Science in Computer Science | Jiwaji University

ACHIEVEMENTS

- Technical Excellence Award — Deloitte
- Rising Star of the Year — Wells Fargo
- Certified Data Engineer — AWS, GCP