Saritha Reddy Kamasani

(940) 843-1998 | sarithakamasani98@gmail.com

PROFESSIONAL SUMMARY:

Data Engineer with 5+ years of experience designing and building scalable data platforms across healthcare, finance, and cloud-native environments. Expertise in developing batch and real-time data pipelines, modern lakehouse architectures, and analytics-ready data models using Databricks, Apache Spark, Airflow, and dbt. Skilled in multi-cloud ecosystems including AWS, Azure, and GCP, with strong experience in data governance, performance optimization, and building reliable data solutions that power enterprise analytics, BI, and machine learning initiatives.

SKILLS:

- **Programming Languages:** Python, R, C++, C#, SQL, Scala, SAS, YAML Java, JavaScript, Matlab, HTML5
- **Database Management Systems (DBMS):** MySQL, PostgreSQL, NoSQL, Oracle, SQL Server, T-SQL, MongoDB, RDS, Cassandra, Elasticsearch. OLTP, RDBMSs, DeltaStreaam
- **Data Warehousing:** Amazon Redshift, Google BigQuery, Snowflake, Apache Hive, Teradata
- **ETL Tools:** Apache Spark, Apache Airflow, Talend, Informatica, Apache NiFi, Data Build Tool (DBT), Dataproc, DataFlow, Fivetran, Stitch, Rivery, Airbyte
- **Big Data Technologies:** Apache Hadoop, Apache Kafka, Apache HBase, Apache Flink, Apache Storm, Apache Iceberg, Apache Druid, Databricks, EMR, Kinesis, Cloudera
- **Cloud Platforms:** AWS, Azure (Microsoft Fabric, Azure Synapse Analytics, ADLS Gen2), GCP, Amazon S3, Azure Data Lake Storage, GCS
- **Version Control Systems:** Git, GitHub
- **Data Visualization:** Tableau, Power BI, Crystal Reports OBIEE, Qlik sense, Alteryx, Cognos
- **Machine Learning/AI:** TensorFlow, PyTorch, TAMR, LLM
- **Operating System:** Linux/Unix, Windows, macOS
- **Containerization and Orchestration**: Docker, Kubernetes
- **Tools & IDE:** Git, IntelliJ, Visual Studio Code, Jupyter Notebook, PyCharm, ER Studio, JIRA, Confluence
- **Hadoop Ecosystem:** HDFS, YARN, MapReduce
- **Monitoring and Logging:** Prometheus, Grafana, ELK stack

EXPERIENCE

Data Engineer

Abbott – Chicago, IL | 07/2025 - Present

- Designed and developed scalable ETL/ELT pipelines using Azure Data Factory and Databricks to process large healthcare and life-sciences datasets.
- Built and optimized PySpark and SQL-based transformations to support batch and near real-time data ingestion.
- Implemented data modeling solutions (star and snowflake schemas) to support enterprise analytics and reporting.
- Managed and optimized data storage in Azure Data Lake (ADLS Gen2) ensuring performance, scalability, and cost efficiency.
- Designed and implemented scalable data solutions using Microsoft Fabric, leveraging Lakehouse architecture, Dataflows, and integrated Power BI semantic models to support enterprise analytics and self-service reporting.
- Developed and optimized data pipelines and dedicated SQL pools in Azure Synapse Analytics, enabling high-performance data transformations, large-scale analytics, and seamless integration with Azure Data Lake.
- Enforced data quality, validation, and reconciliation checks to meet regulatory and audit requirements.
- Applied data governance and access controls using Unity Catalog and Azure Purview to ensure compliance with HIPAA, GxP, and SOX standards.
- Automated pipeline deployments using CI/CD practices with Git and Azure DevOps.
- Collaborated with cross-functional teams (Analytics, QA, Manufacturing, and IT) to translate business requirements into technical data solutions.
- Optimized datasets and semantic layers to support Power BI and Qlik dashboards, improving reporting performance and usability.
- Monitored, troubleshot, and tuned data pipelines to ensure high availability and reliability in production environments.

Data Engineer

Goldman Sachs - New York, NY | 01/2024 – 06/2025

- Engineered scalable data pipelines and distributed processing frameworks using Python, Scala, and Apache Spark to process large enterprise datasets.
- Built event-driven ETL pipelines using AWS Lambda and Python to automate data migration from DynamoDB to Amazon Redshift.
- Developed ingestion frameworks using AWS Glue to process CSV, JSON, and Parquet data from Amazon S3 into Snowflake, Redshift, and DynamoDB.
- Designed and optimized Snowflake data warehouse architecture, including staging, ODS, and dimensional layers using Kimball star schema modeling.
- Implemented scalable dbt and Snowpark transformation pipelines to support analytics and improve pipeline efficiency.

- Orchestrated complex ETL workflows using Apache Airflow DAGs deployed on Kubernetes for reliable pipeline scheduling and monitoring.
- Developed high-performance PL/SQL procedures, packages, functions, and triggers supporting complex business logic and reporting workloads.
- Optimized query performance for large analytical datasets through indexing strategies, aggregations, and materialized views.
- Designed and implemented RESTful and gRPC APIs using Flask and Django for secure microservice communication.
- Built and maintained a centralized log analytics and monitoring platform using Elasticsearch, Logstash, Kibana, Kinesis, and CloudWatch.
- Developed and managed Fivetran ingestion pipelines, configuring connectors and schema mappings for Snowflake and Redshift data integration.
- Integrated large-scale customer behavior and marketing datasets to enable segmentation, attribution, and advanced analytics.
- Prepared structured and unstructured datasets for AI/ML and LLM workloads, including data cleansing, anonymization, and metadata enrichment.
- Automated operational tasks using Bash scripting and UNIX utilities and implemented Infrastructure as Code using Terraform to manage AWS resources.
- Delivered interactive dashboards and analytical insights using Tableau, QlikView, and Splunk to support data-driven decision making.

Data Engineer
IBM – India | 12/2020 – 07/2022

- Designed and implemented enterprise data warehouse and BI solutions using AWS services including S3, Redshift, Lambda, API Gateway, DynamoDB, and AWS Glue.
- Built scalable ETL pipelines using Apache Spark, PySpark, and Scala (RDDs, DataFrames, Spark SQL) for large-scale data transformation and aggregation.
- Developed automated data ingestion frameworks integrating data from APIs, Amazon S3, Teradata, and Snowflake using Python and Scala.
- Engineered batch and real-time data pipelines to process multi-channel customer engagement and marketing analytics datasets.
- Implemented CI/CD pipelines to integrate data engineering workflows with DevOps practices and improve deployment automation.
- Designed and maintained Apache Airflow DAGs to orchestrate complex ETL workflows and scheduled data pipelines.
- Built serverless data processing pipelines using AWS Lambda integrated with API Gateway and DynamoDB for event-driven architectures.
- Processed structured, semi-structured, and unstructured datasets while implementing data profiling and quality validation using Python and SQL.
- Developed Snowflake schemas and ingestion pipelines using Snowpipe and Matillion to process batch and streaming data from AWS S3 data lakes.
- Leveraged Spark Streaming, Hadoop (HDFS, Hive), and GCP services including Dataproc, GCS, Cloud Functions, and BigQuery for distributed data processing and analytics.

Big Data Developer
General Elelectric - India | 06/2018 – 05/2020

- Experience in working multiple projects and teams involved in analytics and cloud platforms.
- Experience in building scalable data pipelines in Azure cloud platform using different tools.
- Developed multiple optimized PySpark applications using Azure Databricks.
- Developed data pipelines using Azure Data Factory that process cosmos activity.
- Delivered real-time marketing and operational dashboards in Power BI, providing actionable insights into customer engagement and business KPIs.
- Developed ETL solutions using SSIS, Azure Data Factory and Azure Data Bricks.
- Expert in working continuous integration and deployment using Jenkins.
- Developed real time ingestion data pipelines from Event Hub into different tools.
- Experience in building ETL solutions using Hive and Spark with Python and Scala.
- Expert in working on optimizing applications built using tools like Spark and Hive.
- Developed job automations from different clusters using Airflow Scheduler.
- Worked on Talend integration with on prem cluster and Azure Cloud Sql for data migrations.
- Developed code coverage and test cases integrations using sonar and Mockito.
- Developed lambda functions on AWS to automate the process of daily manual adhocs.
- Deployed Stream sets an application using a Docker container on the application platform.
- Worked on streaming tools like Stream set to stream the live data into HDFS.
- Worked on automation tools like Oozie to automate the Hive and spark workflows.

EDUCATION

- **Master of Science in Computer Science**
  University of North Texas